


# COS 226

## Chapter 2 Mutual Exclusion

# Acknowledgement



- Some of the slides are taken from the companion slides for “The Art of Multiprocessor Programming” by Maurice Herlihy & Nir Shavit




# Why is Concurrent Programming so Hard?

- Try preparing a seven-course banquet
  - By yourself
  - With one friend
  - With twenty-seven friends ...
- Before we can talk about programs
  - Need a language
  - Describing time and concurrency


# Events

- An *event*  $a_0$  of thread A is
  - Instantaneous
  - No simultaneous events (break ties)



# Threads

- A *thread* A is (formally) a sequence  $a_0, a_1, \dots$  of events
  - “Trace” model
  - Notation:  $a_0 \rightarrow a_1$  indicates order



# Concurrency

- Thread A




- Thread B




# Interleavings

- Events of two or more threads
  - Interleaved
  - Not necessarily independent




# Intervals

- An *interval*  $A_0 = (a_0, a_1)$  is
  - Time between events  $a_0$  and  $a_1$






# Intervals may Overlap




# Intervals may be Disjoint




# Precedence

Interval  $A_0$  **precedes** interval  $B_0$




# Precedence




- Notation:  $A_0 \rightarrow B_0$
- Formally,
  - End event of  $A_0$  before start event of  $B_0$
  - Also called “happens before” or “precedes”

# Precedence Ordering



- Remark:  $A_0 \rightarrow B_0$  is just like saying
  - 1066 AD  $\rightarrow$  1492 AD,
  - Middle Ages  $\rightarrow$  Renaissance,

# Precedence Ordering



- Never true that  $A \rightarrow A$
- If  $A \rightarrow B$  then not true that  $B \rightarrow A$
- If  $A \rightarrow B$  &  $B \rightarrow C$  then  $A \rightarrow C$
- Funny thing:  $A \rightarrow B$  &  $B \rightarrow A$  might both be false!

# Repeated Events

```
while (mumble) {  
    a0; a1;  
}
```

*k*-th occurrence  
of event  $a_0$

$a_0^k$

$A_0^k$

*k*-th occurrence of  
interval  $A_0 = (a_0, a_1)$



# Implementing a Counter

```
public class Counter {  
    private long value;  
  
    public long getAndIncrement() {  
        value++;  
    }  
}
```





# Implementing a Counter

```
public class Counter {  
    private long value;  
  
    public long getAndIncrement() {  
        temp  = value;  
        value = temp + 1;  
        return value;  
    }  
}
```



# Critical Section

- Block of code that can be executed by only one thread at a time
- Needs Mutual Exclusion
- Standard way to approach mutual exclusion is through locks

# Locks (Mutual Exclusion)

```
public interface Lock {  
    public void lock();  
    public void unlock();  
}
```

# Locks (Mutual Exclusion)

```
public interface Lock {
```

```
    public void lock();
```

**acquire lock**

```
    public void unlock();  
}
```

# Locks (Mutual Exclusion)

```
public interface Lock {
```

```
    public void lock();
```

acquire lock

```
    public void unlock();
```

release lock

```
}
```

# Using Locks

```
public class Counter {  
    private long value;  
    private Lock lock;  
    public long getAndIncrement() {  
        lock.lock();  
        try {  
            int temp = value;  
            value = value + 1;  
        } finally {  
            lock.unlock();  
        }  
        return temp;  
    }  
}
```

# Using Locks

```
public class Counter {  
    private long value;  
    private Lock lock;  
    public long getAndIncrement() {  
        lock.lock();  
        try {  
            int temp = value;  
            value = value + 1;  
        } finally {  
            lock.unlock();  
        }  
        return temp;  
    }  
}
```

**lock.lock();** → **acquire Lock**

# Using Locks

```
public class Counter {  
    private long value;  
    private Lock lock;  
    public long getAndIncrement() {  
        lock.lock();  
        try {  
            int temp = value;  
            value = value + 1;  
        } finally {  
            lock.unlock();  
        }  
        return temp;  
    }  
}
```

Release lock  
(no matter what)



# Using Locks

```
public class Counter {  
    private long value;  
    private Lock lock;  
    public long getAndIncrement() {  
        lock.lock();  
        try {  
            int temp = value;  
            value = value + 1;  
        } finally {  
            lock.unlock();  
        }  
        return temp;  
    }  
}
```

**Critical  
section**



# Properties of a good Lock algorithm

- Mutual Exclusion
- Deadlock-free
- Starvation-free



# Mutual Exclusion

- Threads do not access critical section at same time



# Deadlock-free

- If some thread attempts to acquire the lock, some thread will succeed in acquiring the lock



# Deadlock-free

- If **some** thread attempts to acquire the lock, **some** thread will succeed in acquiring the lock
  - System as a whole makes progress
  - Even if individuals starve
  - At least one thread is completing



# Starvation-free

- Every thread that attempts to acquire the lock will eventually succeed



# Starvation-free

- **Every** thread that attempts to acquire the lock will eventually succeed
  - If a thread calls lock() it will eventually acquire the lock
  - Individual threads make progress



# Locks

- Let's start with lock solutions for 2 concurrent threads...



# Two-Thread Conventions

```
class ... implements Lock {  
    ...  
    // thread-local index, 0 or 1  
    public void lock() {  
        int i = ThreadID.get();  
        int j = 1 - i;  
        ...  
    }  
}
```

# Two-Thread Conventions

```
class ... implements Lock {  
    ...  
    // thread-local index, 0 or 1  
    public void lock() {  
        int i = ThreadID.get();  
        int j = 1 - i;  
        ...  
    }  
}
```

Henceforth: *i* is current thread, *j* is other thread



# LockOne

## ■ Basic idea:

- Thread indicates interest in acquiring lock
- Checks to see if other thread is currently in critical section
  - If true, waits until other thread finishes
  - If not, enters critical section



# LockOne

```
class LockOne implements Lock {  
    private boolean[] flag = new boolean[2];  
  
    public void lock() {  
        flag[i] = true;  
        while (flag[j]) {}  
    }  
  
    public void unlock() {  
        flag[i] = false;  
    }  
}
```

# LockOne

```
class LockOne implements Lock {  
    private boolean[] flag =  
                                new boolean[2];  
    public void lock() {  
        flag[i] = true;  
        while (flag[j]) {}  
    }  
  
    public void unlock() {  
        flag[i] = false;  
    }  
}
```

**Set my flag**

# LockOne

```
class LockOne implements Lock {  
    private boolean[] flag =  
        new boolean[2];  
    public void lock() {  
        flag[i] = true;  
        while (flag[j]) {}  
    }  
    public void unlock() {  
        flag[i] = false;  
    }  
}
```

Set my flag

Wait for other  
flag to go false

# LockOne

```
class LockOne implements Lock {  
    private boolean[] flag =  
        new boolean[2];  
    public void lock() {  
        flag[i] = true;  
        while (flag[j]) {}  
    }  
    public void unlock() {  
        flag[i] = false;  
    }  
}
```

When release lock  
set my flag again

# Deadlock Freedom?

- Concurrent execution:

<pre>flag[i] = true;    flag[j] = true; while (flag[j]){} while (flag[i]){} </pre>
--

- If each thread sets its flag to true and waits for the other, they will wait forever
- No deadlock freedom





# LockOne Summary

- LockOne offers mutual exclusion
- When accessed sequentially, LockOne works fine
- However with concurrent threads, LockOne is not Deadlock-free



# LockTwo

## ■ Basic idea:

- When attempting to acquire lock, offer to be the victim that has to defer to other thread
- While current thread is the victim, wait until other thread becomes the victim
- When current thread no longer the victim, enter the critical section



# LockTwo

```
public class LockTwo implements Lock {  
    private int victim;  
    public void lock() {  
        victim = i;  
        while (victim == i) {};  
    }  
  
    public void unlock() {}  
}
```

# LockTwo

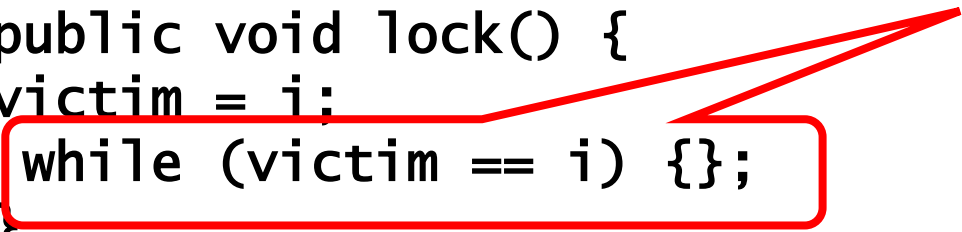
```
public class LockTwo implements Lock {  
    private int victim;  
    public void lock() {  
        victim = i;  
        while (victim == i) {};  
    }  
  
    public void unlock() {}  
}
```

Let other go  
first

# LockTwo

```
public class LockTwo implements Lock {  
    private int victim;  
    public void lock() {  
        victim = i;  
        while (victim == i) {};  
    }  
  
    public void unlock() {}  
}
```

**Wait for permission**



# LockTwo

```
public class Lock2 implements Lock {  
    private int victim;  
    public void lock() {  
        victim = i;  
        while (victim == i) {};  
    }  
    public void unlock() {}  
}
```

Nothing to do



# LockTwo Claims

- Satisfies mutual exclusion

- If thread **i** in CS
- Then **victim == j**
- Cannot be both 0 and 1

```
public void LockTwo() {  
    victim = i;  
    while (victim == i) {}  
}
```



# LockTwo Summary

- LockTwo offers Mutual Exclusion
- Works fine with concurrent threads
- However results in Deadlock with sequential threads
- LockOne and LockTwo thus complement each other





# Peterson Lock

- Combine LockOne and LockTwo
  - Enable successful sequential access provided by LockOne
  - Enables successful concurrent access provided by LockTwo



# Peterson Lock

## ■ Basic idea:

- ☐ Current thread indicates interest in acquiring the lock
- ☐ Current thread offers to be the victim
- ☐ If no interest from other thread and no longer the victim, then continue to critical section

# Peterson's Algorithm

```
public void lock() {  
    flag[i] = true;  
    victim = i;  
    while (flag[j] && victim == i) {};  
}  
public void unlock() {  
    flag[i] = false;  
}
```

# Peterson's Algorithm

Announce I'm  
interested

```
public void lock() {  
    flag[i] = true;  
    victim = i;  
    while (flag[j] && victim == i) {};  
}  
public void unlock() {  
    flag[i] = false;  
}
```

# Peterson's Algorithm

```
public void lock() {  
    flag[i] = true;  
    victim = i;  
    while (flag[j] && victim == i) {};  
}  
public void unlock() {  
    flag[i] = false;  
}
```

Announce I'm  
interested

Defer to other

# Peterson's Algorithm

```
public void lock() {  
    flag[i] = true;  
    victim = i;  
    while (flag[j] && victim == i) {};  
}  
public void unlock() {  
    flag[i] = false;  
}
```

Announce I'm  
interested

Defer to other

Wait while other  
interested & I'm  
the victim

# Peterson's Algorithm

```
public void lock() {  
    flag[i] = true;  
    victim = i;  
    while (flag[j] && victim == i) {};  
}  
public void unlock() {  
    flag[i] = false;  
}
```

Announce I'm interested

Defer to other

Wait while other interested & I'm the victim

No longer interested

# Mutual Exclusion

```
public void lock() {  
    flag[i] = true;  
    victim = i;  
    while (flag[j] && victim == i) {};
```

- If thread 0 in critical section,
  - flag[0] = true,
  - victim = 1
- If thread 1 in critical section,
  - flag[1] = true,
  - victim = 0

Cannot both be true



# Deadlock Free

```
public void lock() {  
    ...  
    while (flag[j] && victim == i) {};  
}
```

- Thread blocked
  - only at **while** loop
  - Only if other's flag is true
  - only if it is the **victim**
- Solo: other's flag is false
- Both: one or the other must not be the victim

# Starvation Free

- Thread *i* would be blocked only if *j* repeatedly re-enters so that

`flag[j] == true` and  
`victim == i`

- When *j* re-enters
  - it sets `victim` to *j*.
  - So *i* gets in
- Thus: Starvation free

```
public void lock() {  
    flag[i] = true;  
    victim  = i;  
    while (flag[j] && victim == i) {};  
}  
  
public void unlock() {  
    flag[i] = false;  
}
```



# Locks

- Moving on to solutions for  $n$  concurrent threads




# Filter Lock

- Peterson lock adapted to work with  $n$  threads instead of just 2
- Thread has to traverse  $n-1$  waiting rooms in order to acquire the lock

# The Filter Algorithm for $n$ Threads

There are  $n-1$  “waiting rooms” called levels

- At each level
  - At least one enters level
  - At least one blocked if many try



- Only one thread makes it through

# Filter

```
class Filter implements Lock {  
    int[] level; // level[i] for thread i  
    int[] victim; // victim[L] for level L  
  
    public Filter(int n) {  
        level = new int[n];  
        victim = new int[n];  
        for (int i = 1; i < n; i++) {  
            level[i] = 0;  
        }  
        ...  
    }  
}
```

# Filter

```
class Filter implements Lock {  
    ...  
  
    public void lock(){  
        for (int L = 1; L < n; L++) {  
            level[i] = L;  
            victim[L] = i;  
            while (( $\exists k \neq i$ ) level[k] >= L) &&  
                victim[L] == i );  
        }  
    }  
    public void unlock() {  
        level[i] = 0;  
    }  
}
```

# Filter

```
class Filter implements Lock {
```

```
...
```

```
public void lock() {
```

```
    for (int L = 1; L < n; L++) {
```

```
        level[i] = L;
```

```
        victim[L] = i;
```

```
        while (( $\exists k \neq i$ ) level[k] >= L) &&  
            victim[L] == i);
```

```
    }
```

```
public void release(int i) {
```

```
    level[i] = 0;
```

```
}
```

One level at a time



# Filter

```
class Filter implements Lock {
```

```
...
```

```
public void lock() {
```

```
    for (int L = 1; L < n; L++) {
```

```
        level[i] = L;
```

```
        victim[L] = 1;
```

```
        while (( $\exists k \neq i$ ) level[k] >= L) &&  
            victim[L] == 1); // busy wait
```

```
    }
```

```
public void release(int i) {
```

```
    level[i] = 0;
```

```
}
```

**Announce  
intention to  
enter level L**

# Filter

```
class Filter implements Lock {  
    int level[n];  
    int victim[n];  
    public void lock() {  
        for (int L = 1; L < n; L++) {  
            level[i] = L;  
            victim[L] = i;  
            while (( $\exists k \neq i$ ) level[k] >= L) &&  
                victim[L] == i);  
        }  
    }  
    public void release(int i) {  
        level[i] = 0;  
    }  
}
```

**Give priority to  
anyone but me**

# Filter

Wait as long as someone else is at  
same or higher level, and I'm  
designated victim

```
class Filter implements Lock {
    int level[n];
    int victim[n];
    public void lock() {
        for (int L = 1; L < n; L++) {
            level[i] = L;
            victim[L] = i;
            while (( $\exists k \neq i$ ) level[k] >= L) &&
                victim[L] == i);
        }
    }
    public void release(int i) {
        level[i] = 0;
    }
}
```

# Filter

```
class Filter implements Lock {  
    int level[n];  
    int victim[n];  
    public void lock() {  
        for (int L = 1; L < n; L++) {  
            level[i] = L;  
            victim[L] = i;  
            while (( $\exists k \neq i$ ) level[k] >= L) &&  
                victim[L] == i);  
        }  
    }  
    public void release(int i) {  
        level[i] = 0;  
    }  
}
```

**Thread enters level L when it completes the loop**



# No Starvation

- Filter Lock satisfies properties:
  - Just like Peterson Algorithm at any level
  - So no one starves
- But what about fairness?
  - Threads can be overtaken by others



# Waiting

- Starvation freedom guarantees that every thread that calls `lock()` eventually enters the critical section
- It however makes no guarantee about how long that can take



# Waiting

- Ideally if A calls lock() before B, then A should enter critical section before B
- However this does not currently work since we cannot determine which thread called lock() first
- Locks should thus be further defined



# Fairness

- Locks should be first-some-first served





# Filter Lock again

- Filter Lock satisfies properties:
  - No one starves
  - But very weak fairness
    - Can be overtaken **arbitrary** # of times
  - That's pretty lame...



# Bakery Algorithm

- Provides First-Come-First-Served
- How?
  - Take a “number”
  - Wait until lower numbers have been served
- Each thread takes a number when attempting to acquire the lock and waits until no thread with an earlier number is trying to acquire it

# Bakery Algorithm

```
class Bakery implements Lock {  
    boolean[] flag;  
    Label[] label;  
    public Bakery (int n) {  
        flag = new boolean[n];  
        label = new Label[n];  
        for (int i = 0; i < n; i++) {  
            flag[i] = false; label[i] = 0;  
        }  
    }  
    ...  
}
```



# Bakery Algorithm

- $\text{flag}[A]$  is a boolean flag indicating whether  $A$  wants to enter the critical section
- $\text{label}[A]$  is an integer that contains thread  $A$ 's “number” when entering the bakery

# Bakery Algorithm

```
class Bakery implements Lock {  
    ...  
    public void lock() {  
        flag[i] = true;  
        label[i] = max(label[0], ..., label[n-1])+1;  
        while ( $\exists k$  flag[k]  
                && (label[k] < label[i]));  
    }
```

# Bakery Algorithm

```
class Bakery implements Lock {  
    ...  
    public void lock() {  
        flag[i] = true;  
        label[i] = max(label[0], ..., label[n-1])+1;  
        while ( $\exists k$  flag[k]  
                && (label[k] < label[i]));  
    }  
}
```

I'm interested

# Bakery Algorithm

Take increasing  
label (read labels  
in some arbitrary  
order)

```
class Bakery implements Lock {  
    ...  
    public void lock() {  
        flag[i] = true;  
        label[i] = max(label[0], ..., label[n-1])+1;  
        while ( $\exists k$  flag[k]  
                && (label[k] < label[i]));  
    }  
}
```

Label is created as one greater than the maximum of the other thread's labels

# Bakery Algorithm

```
class Bakery implements Lock {  
    ...  
    public void lock() {  
        flag[i] = true;  
        label[i] = max(label[0], ..., label[n-1])+1;  
        while ( $\exists k$  flag[k]  
                && (label[k] < label[i]));  
    }
```

Someone is  
interested



# Bakery Algorithm

```
class Bakery implements Lock {  
    ...  
    public void lock() {  
        flag[i] = true;  
        label[i] = max(label[0], ..., label[n-1])+1;  
        while ( $\exists k$  flag[k]  
            && (label[k] < label[i]));  
    }
```

Someone is  
interested

And someone has  
a smaller number  
than me

THEORETICALLY



# Bakery Algorithm

- If two threads try to acquire the lock concurrently, they may read the same maximum number
- Threads thus have unique pairs consisting of number as well as thread ID



# Bakery Algorithm

$$(\text{label}[i], i) \ll (\text{label}[j], j)$$

If and only if

$\text{label}[i] < \text{label}[j]$  OR  $\text{label}[i] = \text{label}[j]$  and  $i < j$

# Bakery Algorithm

```
class Bakery implements Lock {  
    boolean flag[n];  
    int label[n];  
  
    public void lock() {  
        flag[i] = true;  
        label[i] = max(label[0], ..., label[n-1])+1;  
        while ( $\exists k$  flag[k]  
                && (label[k], k) << (label[i], i));  
    }  
}
```

Someone is  
interested

With lower (label,i)  
in lexicographic order



# Bakery Algorithm

- In other words:
- Thread A must wait if:
  - Another thread is interested AND the other thread's number is lower than thread A
    - OR
  - Another thread is interested AND the two threads have the same number but the other's threads ID is smaller than A

# Bakery Algorithm

```
class Bakery implements Lock {  
  
    ...  
  
    public void unlock() {  
        flag[i] = false;  
    }  
}
```

# Bakery Algorithm

```
class Bakery implements Lock {
```

```
    ...
```

```
    public void unlock() {
```

```
        flag[i] = false;
```

```
    }
```

```
}
```

**No longer  
interested**





# To analyse:

- Does the lock provide:

- Mutual exclusion?

- YES – two threads cannot be in the critical section at the same time since one of them will have an earlier label pair





# To analyse:

- Starvation freedom?

- YES – if a thread exists the critical section and immediately wants to reacquire the lock, he will first have to take a new, later number allowing the other waiting threads to gain access first



# To analyse:

- Deadlock freedom?

- YES – there is always one thread with the earliest label, ties are not possible because of labels consist of number and order in array



# To analyse:

- The Bakery algorithm also provides First-come-first-served
  - If A calls lock() before B, then A's number is smaller than B's number
  - So B is locked out while flag[A] is true



# Potential issue:

- With the current Bakery algorithm we are assuming that we have an infinite amount of numbers to use
- In practice this is not the case



# Bounded timestamps

- Labels in the Bakery lock grow without bounds
- In a long-lived system we may have to worry about overflow
- If a thread's label silently rolled over from a large number to zero, the first-come-first-served property no longer holds



# Bounded timestamps

- In the Bakery algorithm, the idea of labels can be replaced by timestamps
- Timestamps can ensure order among the contending threads
- We will thus need to ensure that if one thread takes a label after another, then the latter has the higher timestamp



# Bounded timestamps

- Timestamps need the ability to:
  - Scan – read the other thread's timestamps
  - Label – assign itself a larger timestamp



# Possible solution

- To construct a **Sequential** timestamping system
- Each thread perform scan-and-label completely one after the other
- Uses mutual exclusion





# Bakery algorithm

- The Bakery algorithm is elegant and fair
- However it is not considered practical
  - Why?
  - Principal drawback is the need to read  $n$  distinct location where  $n$  can be very large