

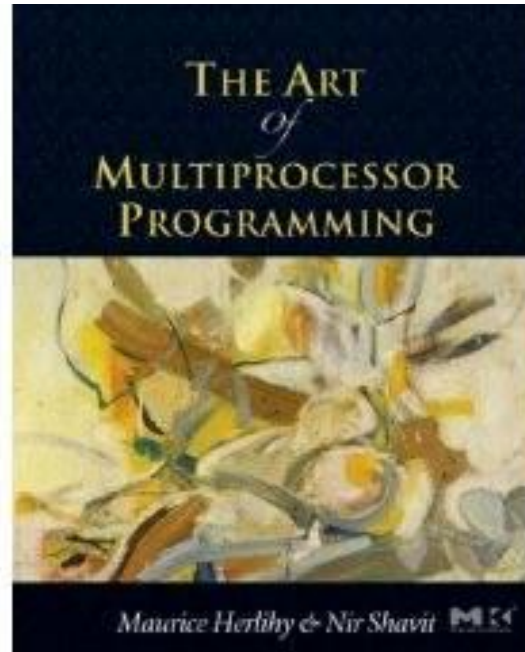


# COS 226

## Chapter 7

## Spin Locks and Contention

# Acknowledgement



- Some of the slides are taken from the companion slides for “The Art of Multiprocessor Programming” by Maurice Herlihy & Nir Shavit



# Concurrency issues

- Memory contention:
  - Not all processors can access the same memory at the same time and if they try they will have to queue
- Contention for communication medium:
  - If everyone wants to communicate at the same time, some of them will have to wait
- Communication latency:
  - If takes more time for a processor to communicate with memory or with another processor.



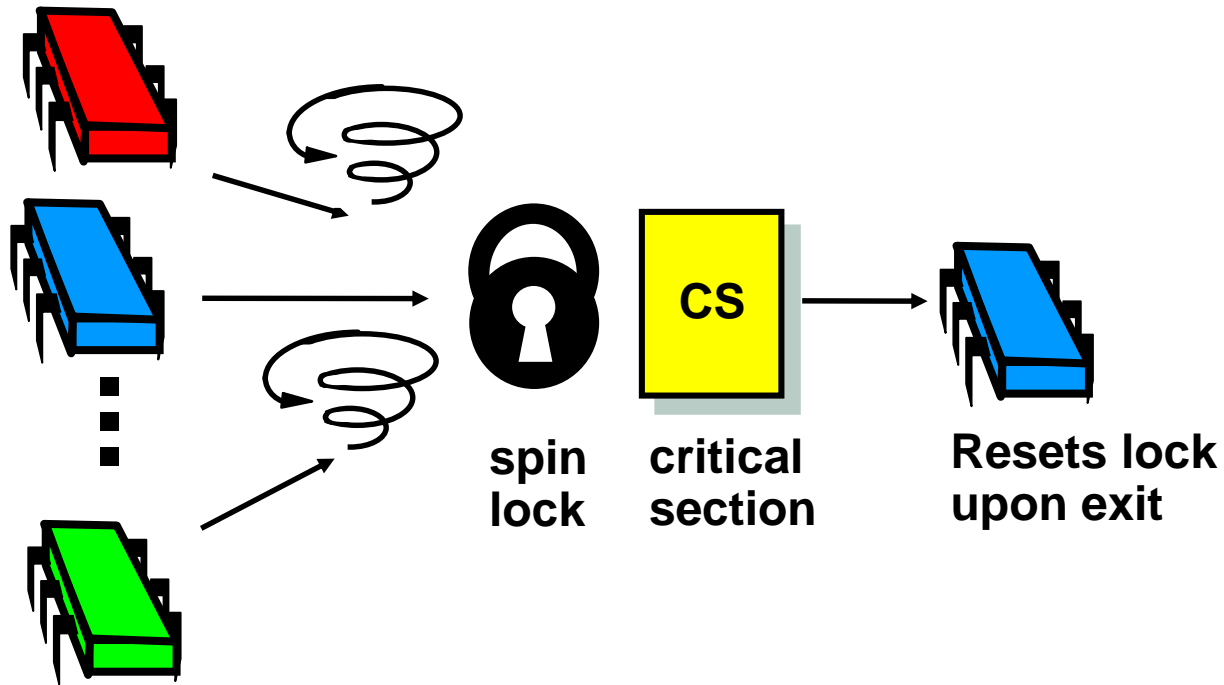
# New goals

- Think of performance, not just correctness and progress
- Understand the underlying architecture
- Understand how the architecture affects performance
- Start with Mutual Exclusion

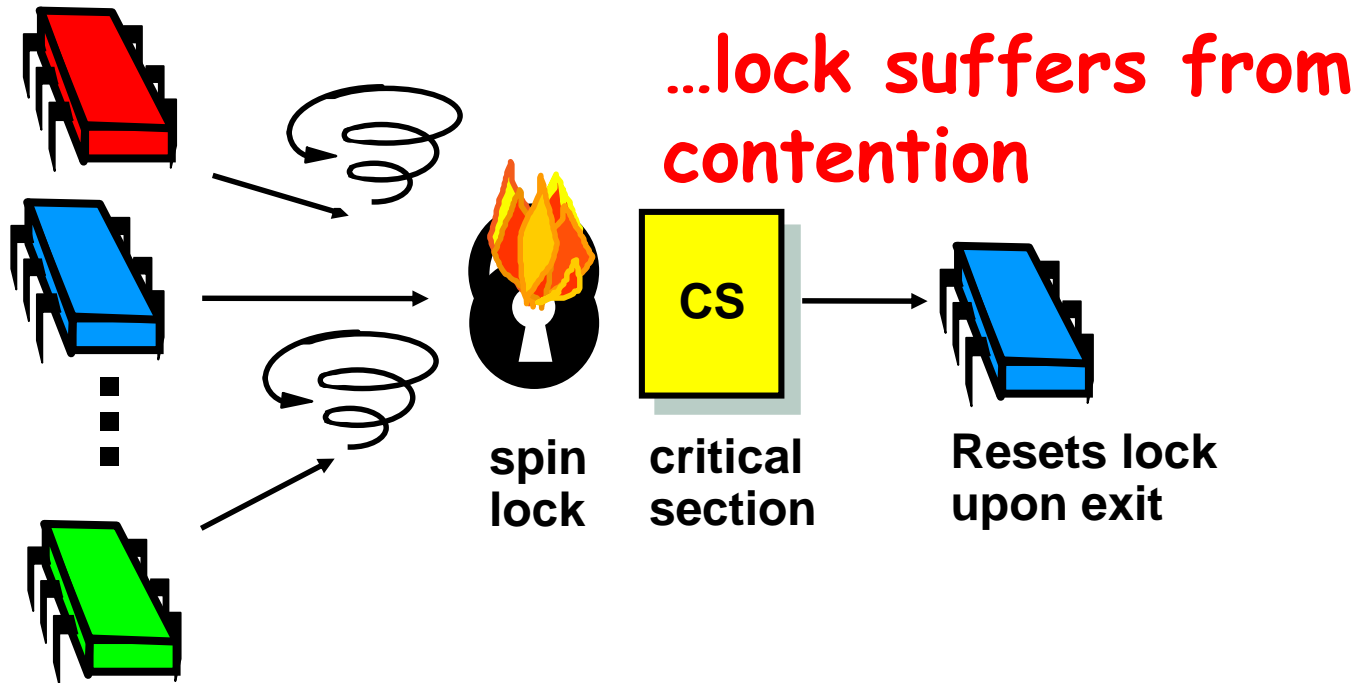
# What should you do if you can't get a lock?

- Keep trying
    - “spin” or “busy-wait”
    - Good if delays are short
  - Give up the processor
    - Suspend yourself and ask the schedule to create another thread on your processor
    - Good if delays are long
    - Always good on uniprocessor
- our focus

# Basic Spin-Lock



# Basic Spin-Lock





# Contention

- Contention:

- When multiple threads try to acquire a lock at the same time

- High contention:

- There are many such threads

- Low contention:

- The opposite





# Welcome to the real world

- Java Lock interface

- `java.util.concurrent.locks` package

```
Lock mutex = new LockImpl (...);  
...  
mutex.lock();  
try {  
    ...  
} finally {  
    mutex.unlock();  
}
```



# Why don't we just use the Filter or Bakery Locks?

- The principal drawback is the need to read and write  $n$  distinct locations where  $n$  is the number of concurrent threads
- This means that the locks require space linear in  $n$




# What about the Peterson lock?

```
class Peterson implements Lock {  
    private boolean[] flag = new boolean[2];  
    private int victim;  
    public void lock() {  
        int i = ThreadID.get();  
        int j = 1 - i;  
        flag[i] = true;  
        victim = i;  
        while (flag[j] && victim == i) {}  
    }  
}
```



# Peterson lock?

- It is not our logic that fails, but our assumptions about the real world
- We assumed that read and write operations are atomic
- Our proof relied on the assumption that any two memory accesses by the same thread, even to different variables, take place in program order



# Why does it not take place in program order?

- Modern multiprocessors do not guarantee program order
- Due to:
  - Compilers
    - reorder instructions to enhance performance
  - Multiprocessor hardware itself
    - writes to multiprocessor memory do not necessarily take effect when they are issued
    - writes to shared memory are buffered and written to memory only when needed

# What about the Peterson lock?

```
class Peterson implements Lock {  
    private boolean[] flag = new boolean[2];  
    private int victim;  
    public void lock() {  
        int i = ThreadID.get();  
        int j = 1 - i;  
        flag[i] = true;  
        victim = i;  
        while (flag[j] && victim == i) {};  
    }  
}
```

Important that these  
steps take place in  
program order



# How can one fix this?

- Memory barriers (or memory fences) can be used to force outstanding operations to take effect
- It is the programmer's responsibility to know when to insert a memory barrier
- However, memory barriers are expensive



# Memory barriers

- Synchronization instructions such as `getAndSet()` or `compareAndSet()` often include memory barriers
- As do reads and writes to **volatile** fields





# Review: Test-and-Set

- Boolean value
- Test-and-set (TAS)
  - Swap **true** with current value
  - Return value tells if prior value was **true** or **false**
- Can reset just by writing **false**
- TAS aka “getAndSet”



# Review: Test-and-Set

```
public class AtomicBoolean {  
    boolean value;  
  
    public synchronized boolean  
        getAndSet(boolean newValue) {  
        boolean prior = value;  
        value = newValue;  
        return prior;  
    }  
}
```

# Review: Test-and-Set

```
public class AtomicBoolean {
```

```
    boolean value;
```

```
    public synchronized boolean  
    getAndSet(boolean newValue) {
```

```
        boolean prior = value;
```

```
        value = newValue;
```

```
        return prior;
```

```
    }
```

```
}
```

Package

java.util.concurrent.atomic

# Review: Test-and-Set

```
public class AtomicBoolean {  
    boolean value;
```

```
    public synchronized boolean  
    getAndSet(boolean newValue) {  
        boolean prior = value;  
        value = newValue;  
        return prior;  
    }  
}
```

**Swap old and new  
values**



# Review: Test-and-Set

```
AtomicBoolean lock
= new AtomicBoolean(false)
...
boolean prior = lock.getAndSet(true)
```

# Review: Test-and-Set

```
AtomicBoolean lock  
= new AtomicBoolean(false)
```

```
boolean prior = lock.getAndSet(true)
```

Swapping in `true` is called  
"test-and-set" or TAS



# Test-and-Set Locks

- Locking
  - Lock is free: value is false
  - Lock is taken: value is true
- Acquire lock by calling TAS
  - If result is false, you win
  - If result is true, you lose
- Release lock by writing false



# Test-and-set Lock

```
class TASLock {  
    AtomicBoolean state =  
        new AtomicBoolean(false);  
  
    void lock() {  
        while (state.getAndSet(true)) {}  
    }  
  
    void unlock() {  
        state.set(false);  
    }  
}
```



# Test-and-set Lock

```
class TASLock {  
    AtomicBoolean state =  
        new AtomicBoolean(false);  
  
    void lock() {  
        while (state.getAndSet(true)) {}  
    }  
  
    void unlock() {  
        state.set(false);  
    }  
}
```

**Lock state is  
AtomicBoolean**

# Test-and-set Lock

```
class TASLock {  
    AtomicBoolean state =  
        new AtomicBoolean(false);
```

```
    void lock() {  
        while (state.getAndSet(true)) {}  
    }
```

```
    void unlock() {  
        state.set(false);  
    }  
}
```

**Keep trying until  
lock acquired**

# Test-and-set Lock

```
class TASLock {  
    AtomicBoolean state =  
        new AtomicBoolean(false);  
  
    void lock() {  
        while (state.getAndSet(true)) {}  
    }  
  
    void unlock() {  
        state.set(false);  
    }  
}
```

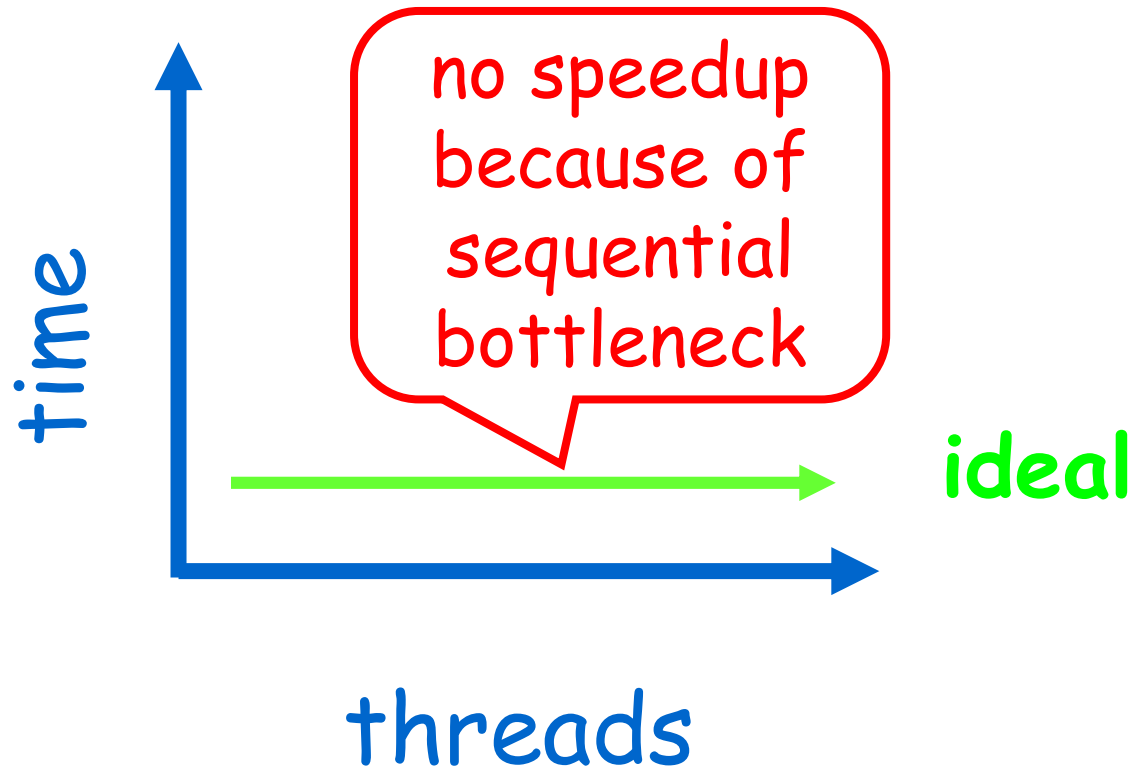
**Release lock by  
resetting state to  
false**



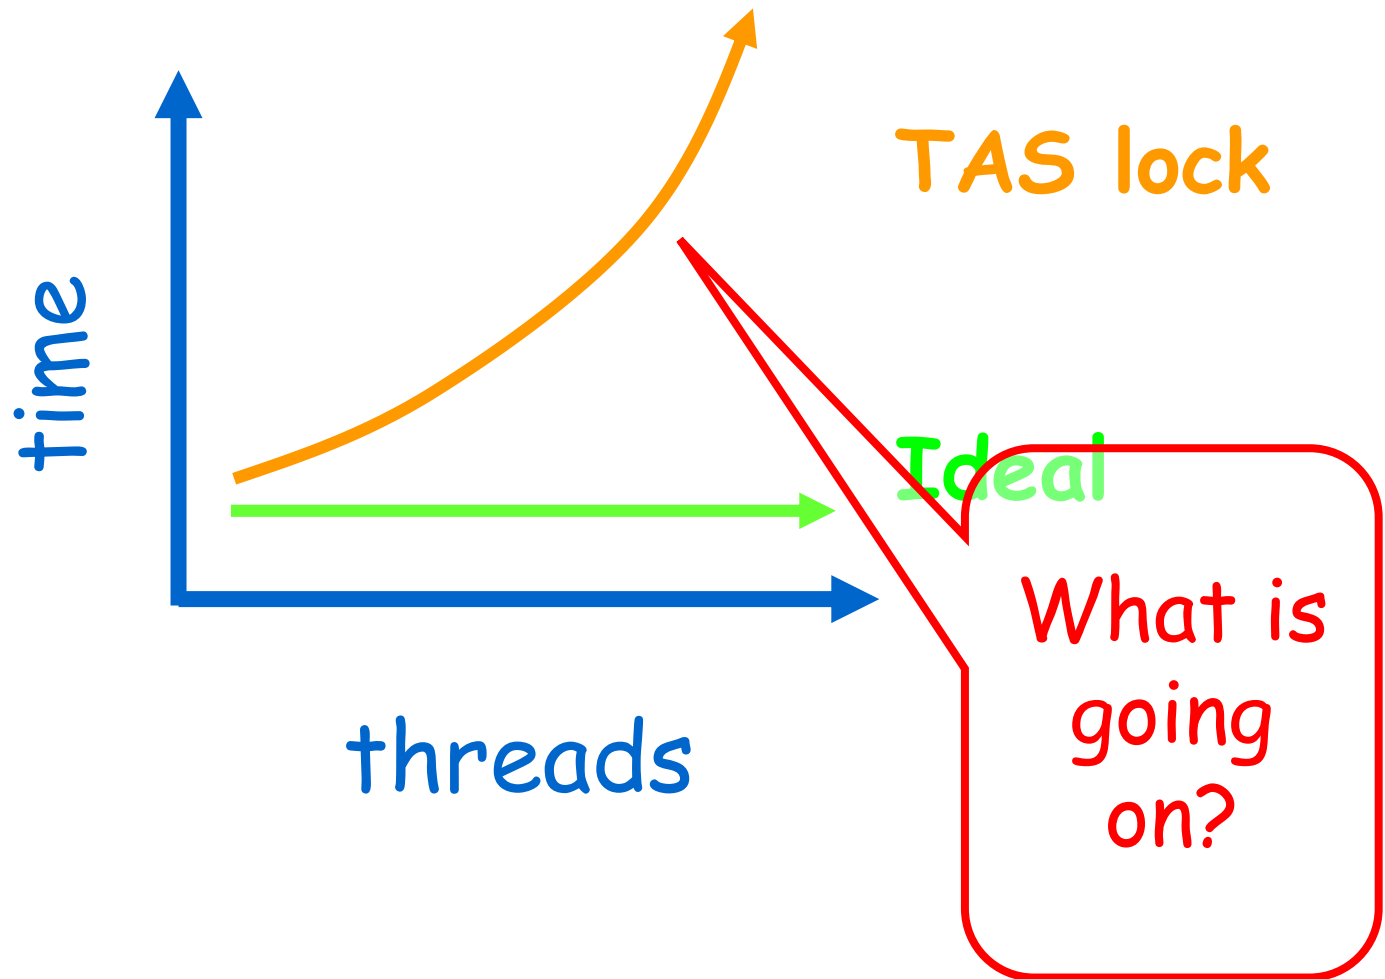
# Performance

- Experiment
  - $n$  threads
  - Increment shared counter 1 million times
- How long should it take?
- How long does it take?

# Graph



# Mystery #1

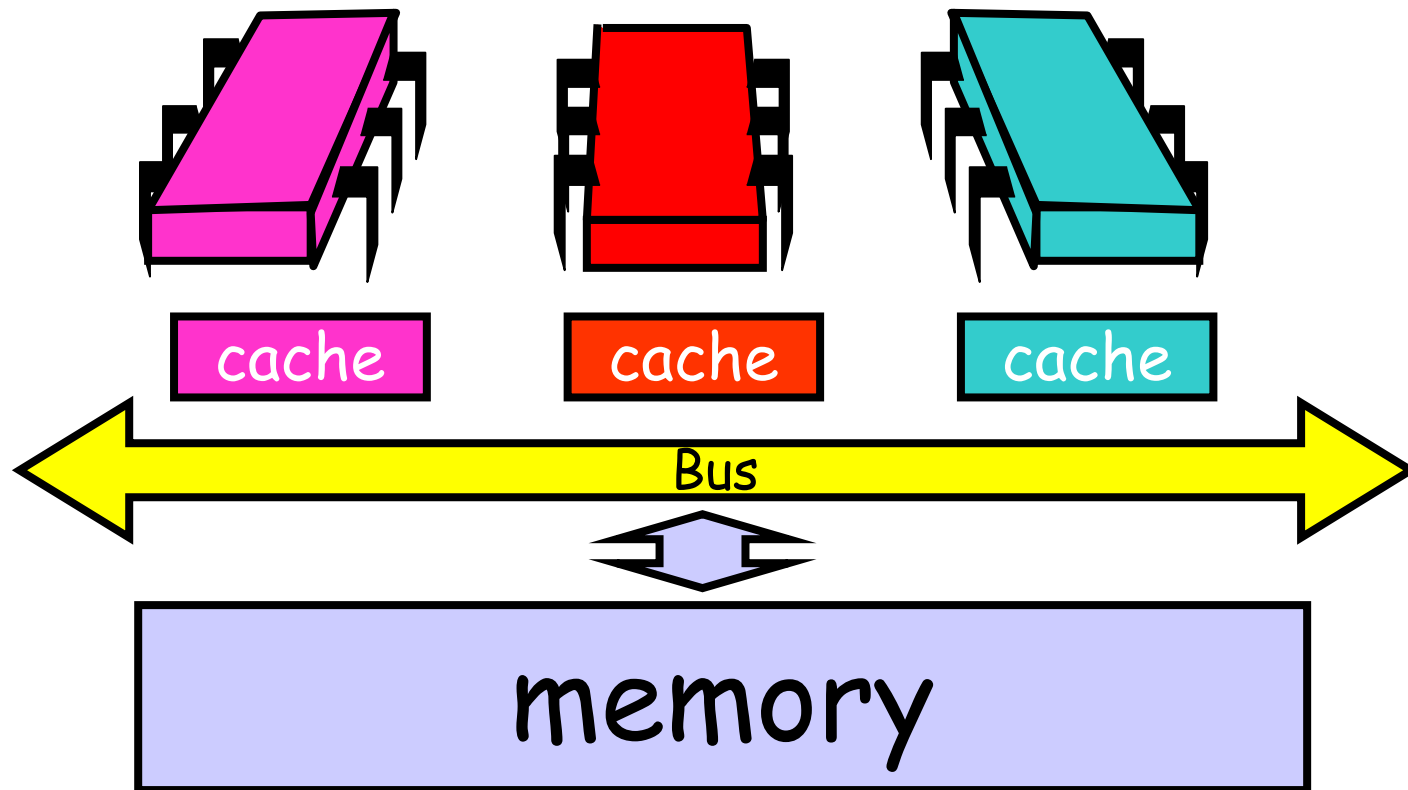




# Questions

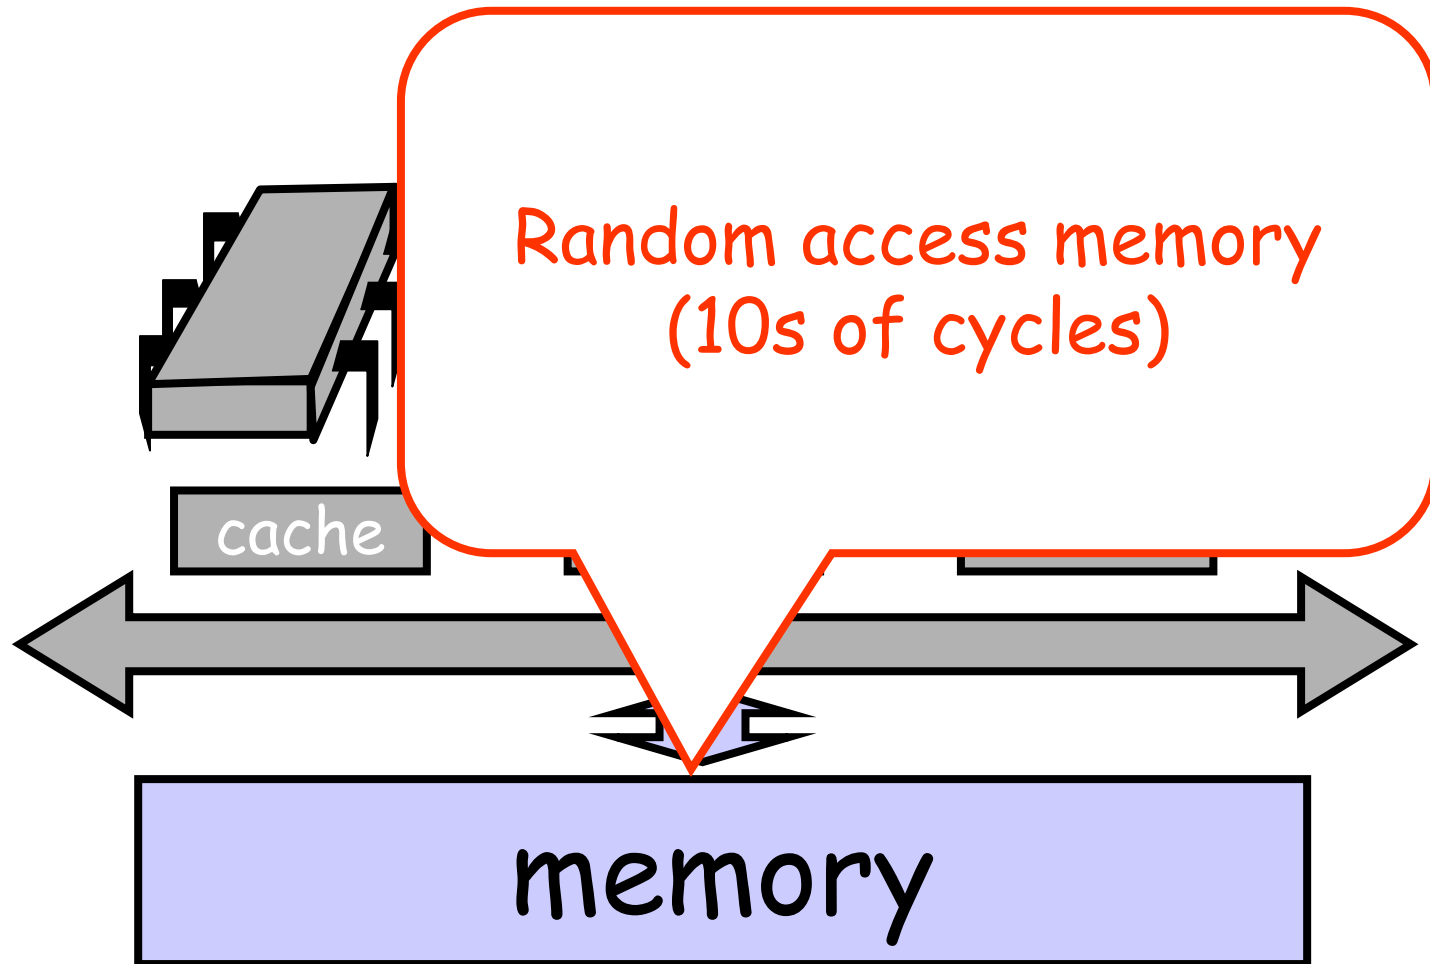
- Why is TAS so bad (so much worse than ideal)?

# Bus-Based Architectures






# Bus-Based Architectures



# Bus-Based Architectures

## Shared Bus

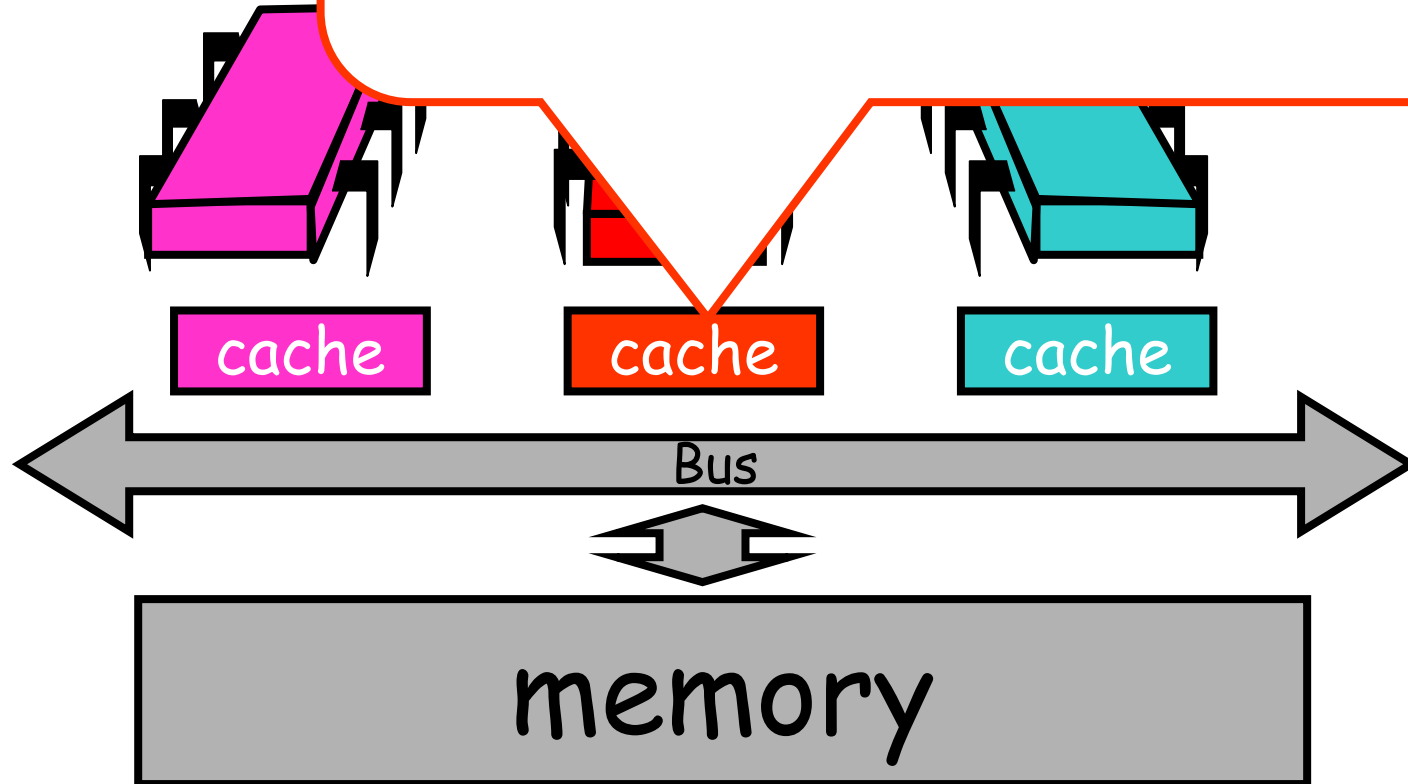
- Broadcast medium
- One broadcaster at a time
- Processors and memory all "snoop"



# Bus-Bas

## Per-Processor Caches

- Small
- Fast: 1 or 2 cycles
- Address & state information





# Jargon Watch

- Cache hit

- ☐ “I found what I wanted in my cache”
- ☐ Good Thing™



# Jargon Watch

## ■ Cache hit

- “I found what I wanted in my cache”
- Good Thing™

## ■ Cache miss

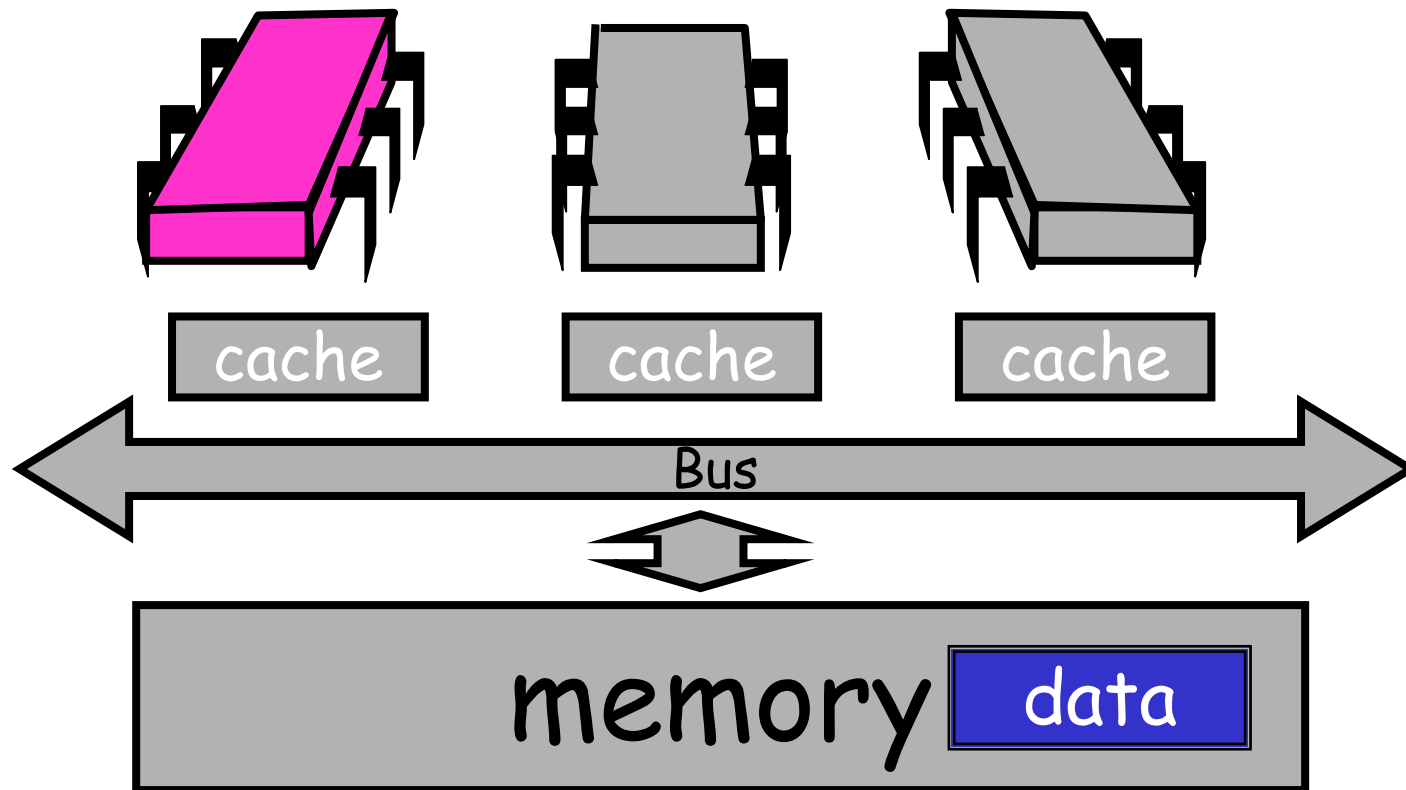
- “I had to shlep all the way to memory for that data”
- Bad Thing™




# Cave Canem

- This model is still a simplification
  - But not in any essential way
  - Illustrates basic principles
- Will discuss complexities later

# Processor Issues Load Request

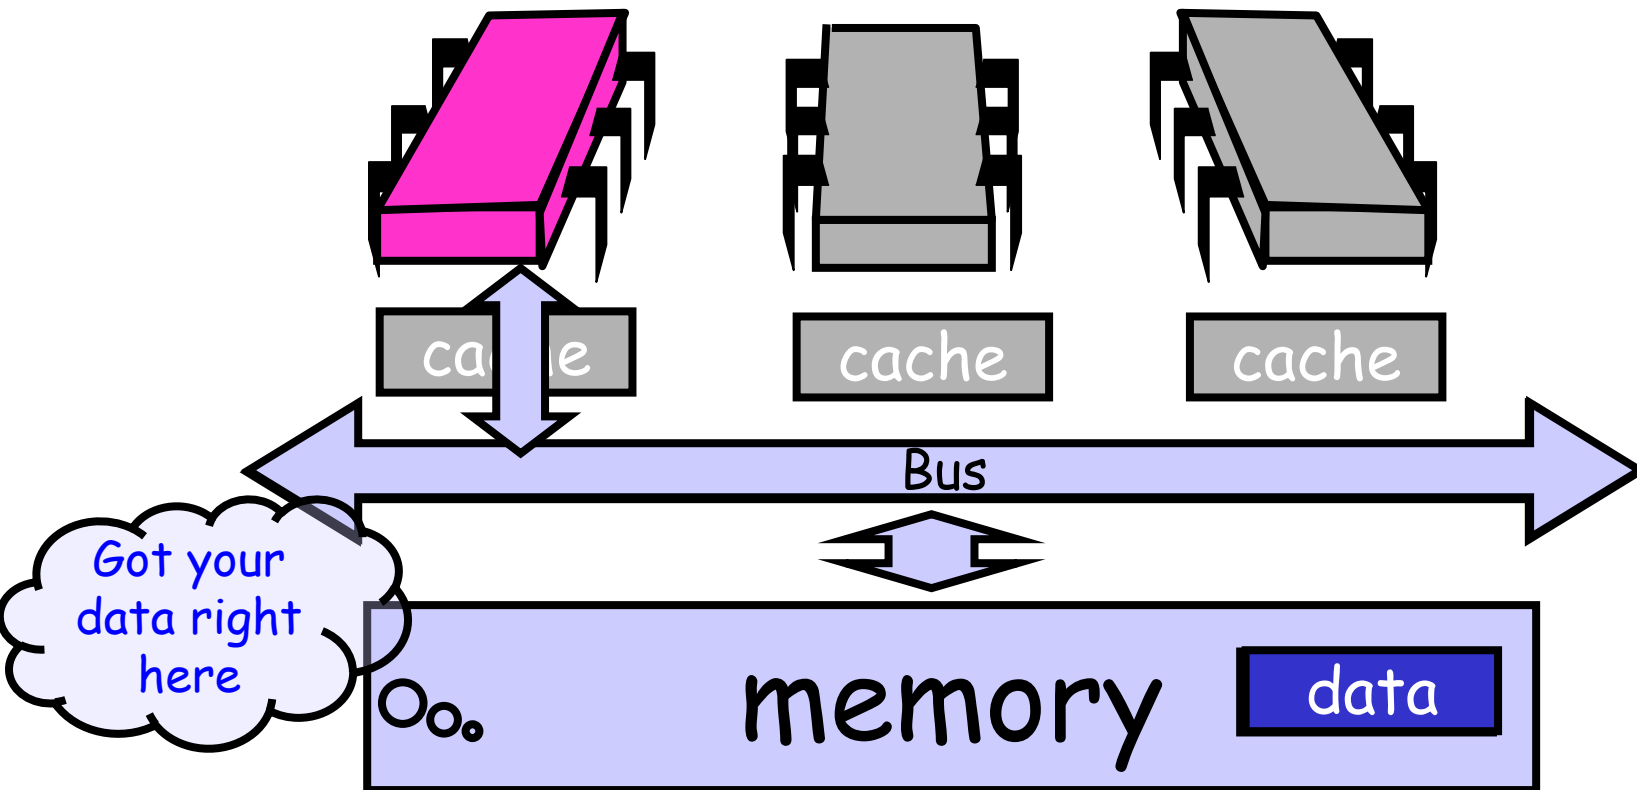


# Processor Issues Load Request

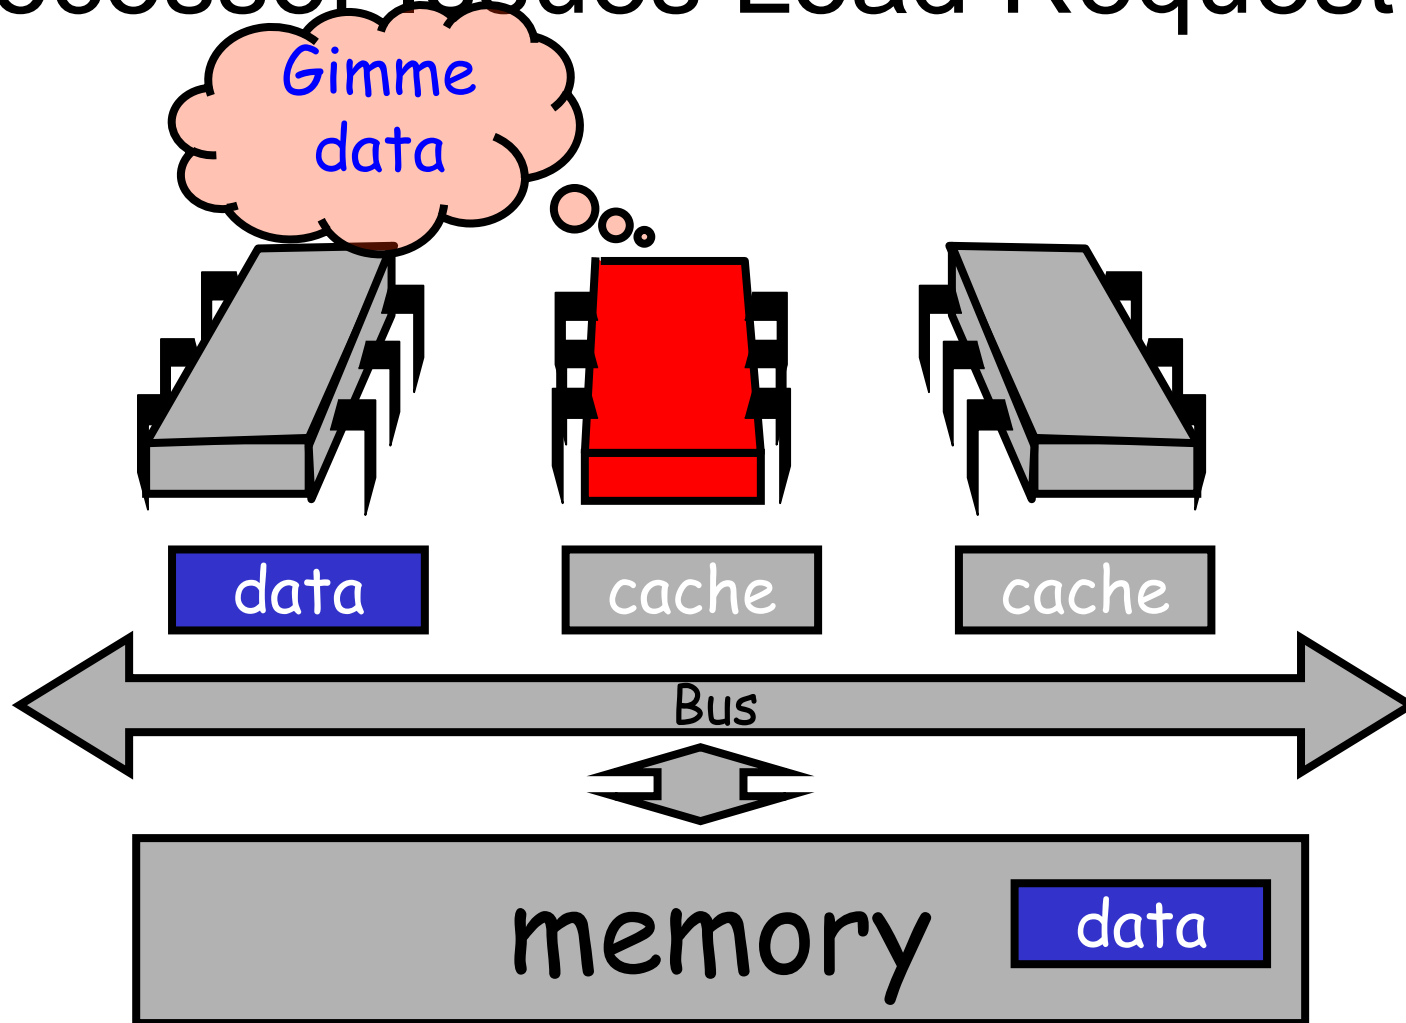




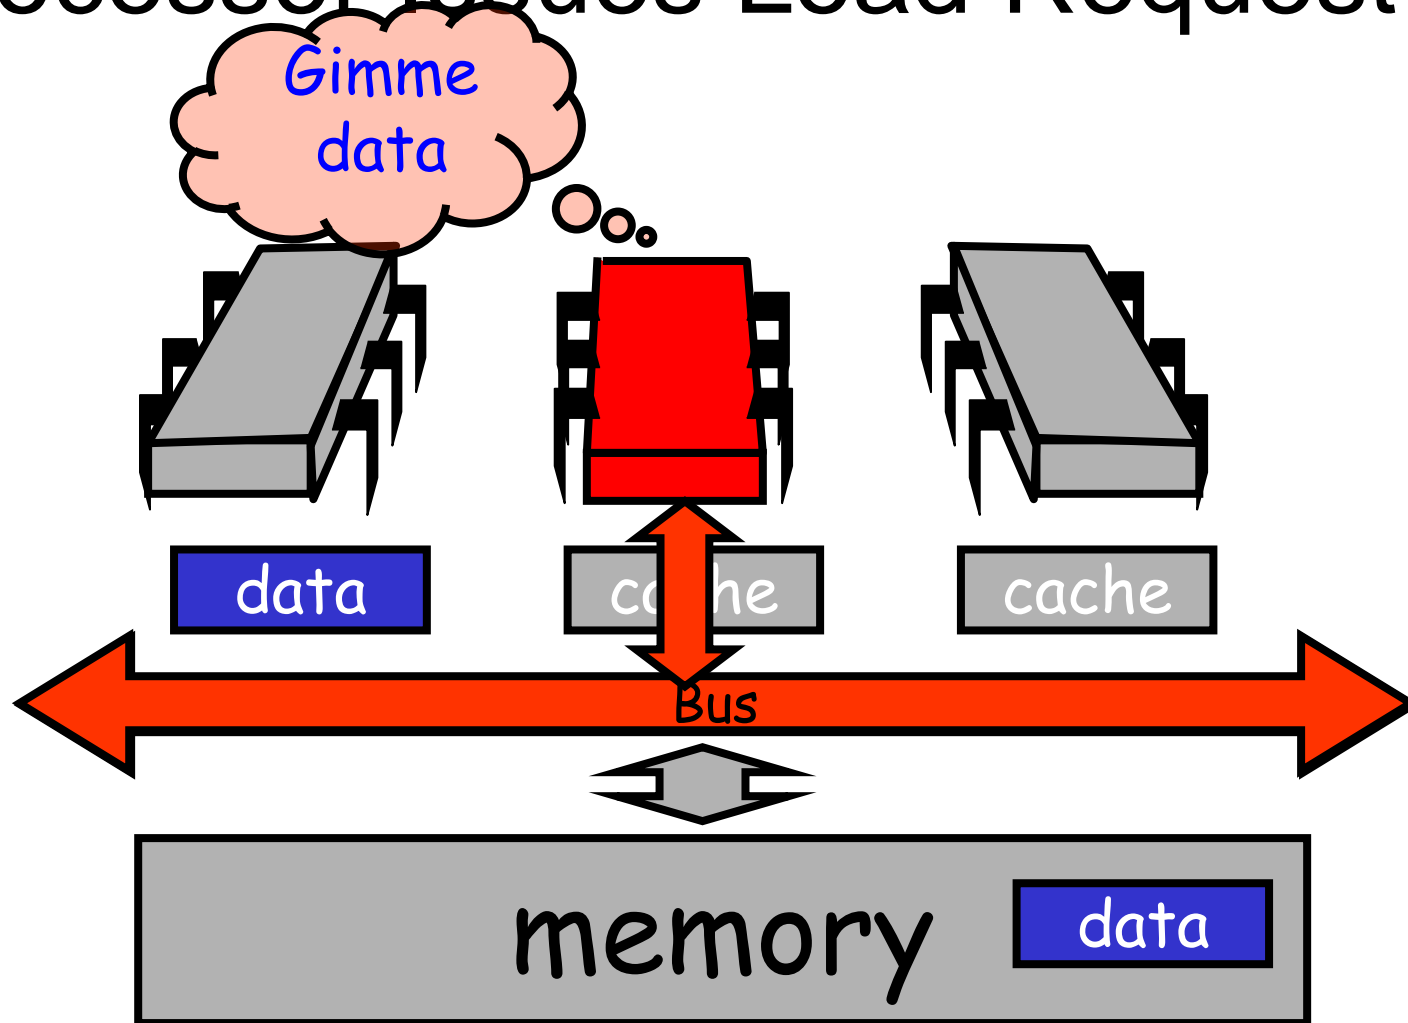
# Memory Responds



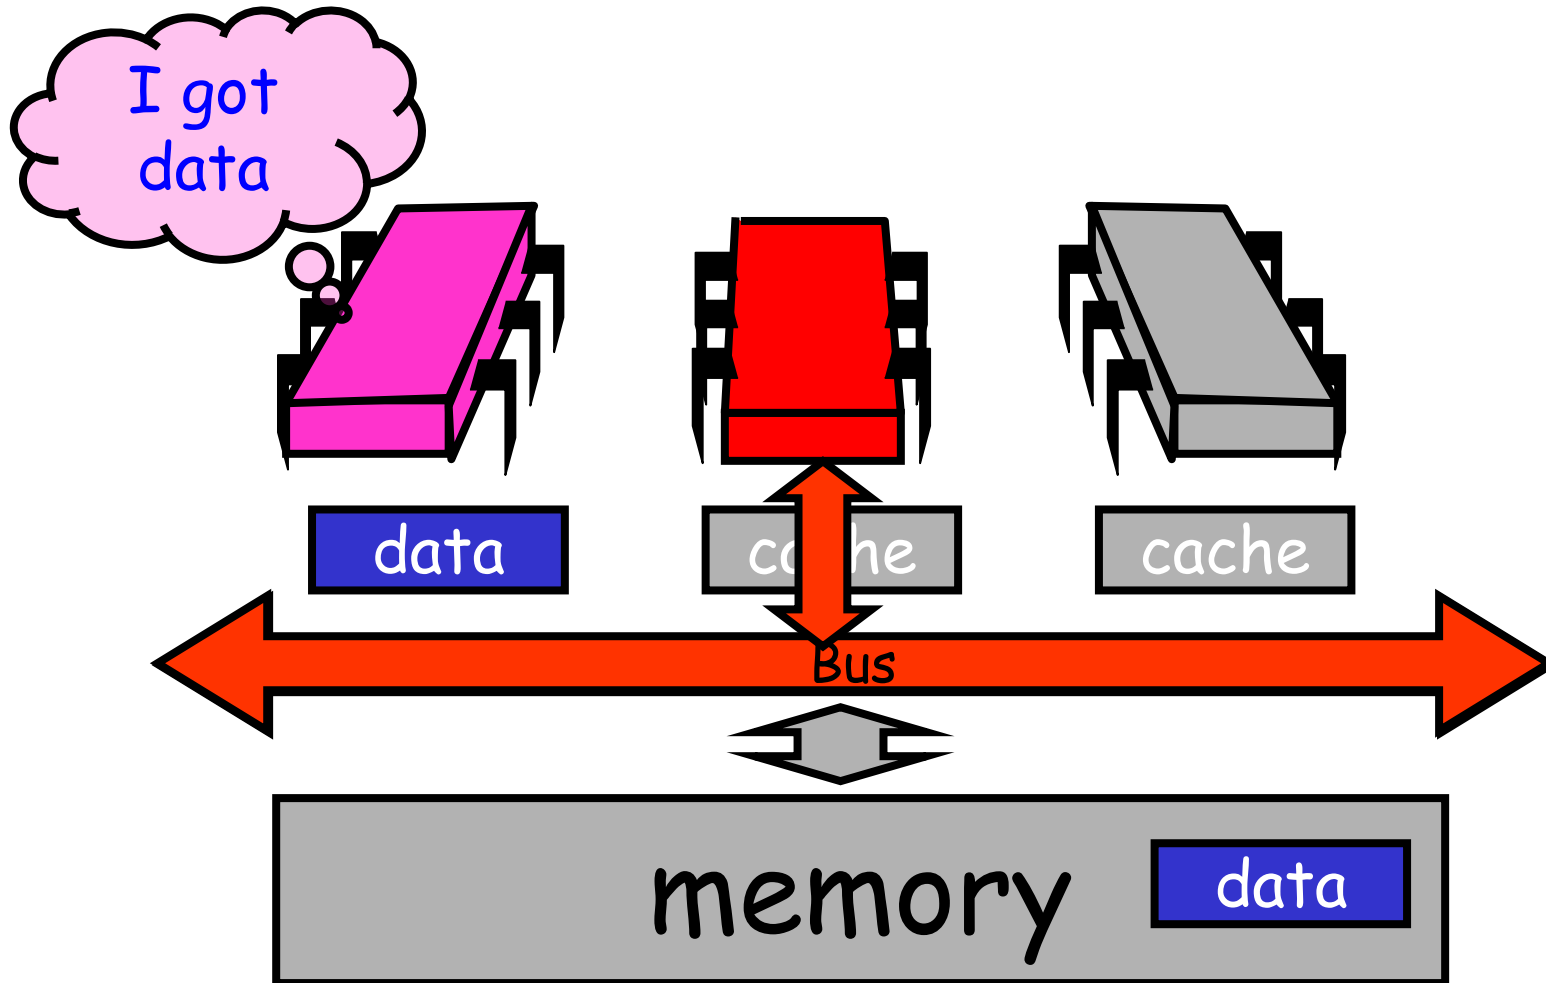
# Processor Issues Load Request



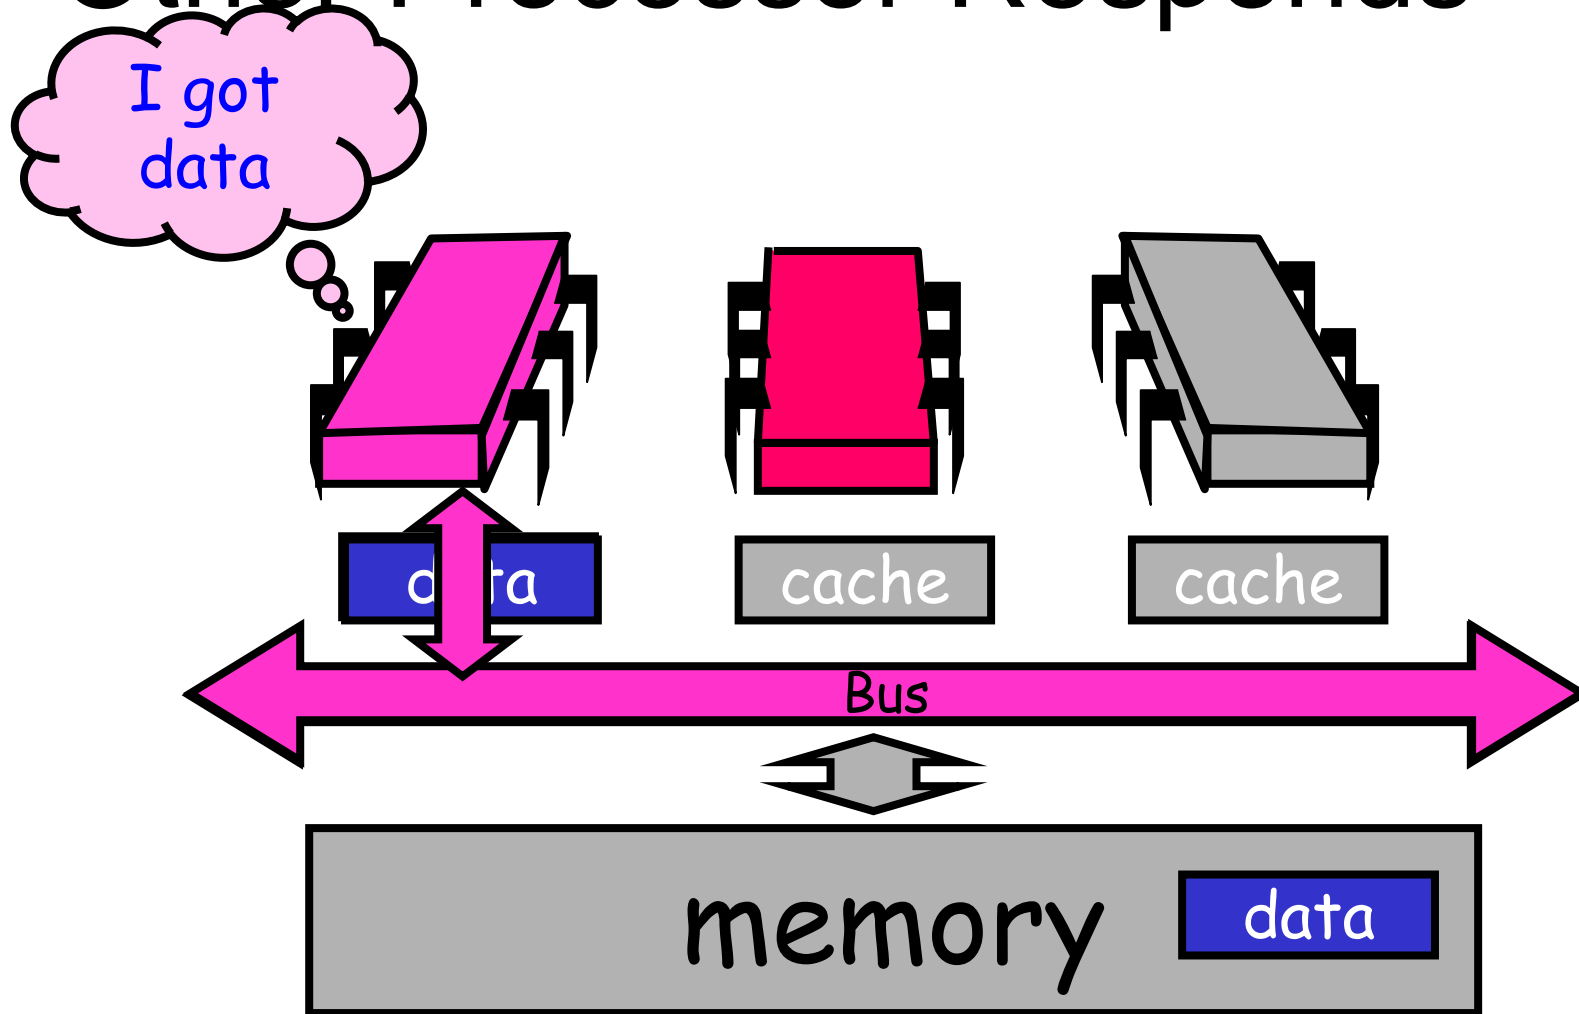
# Processor Issues Load Request



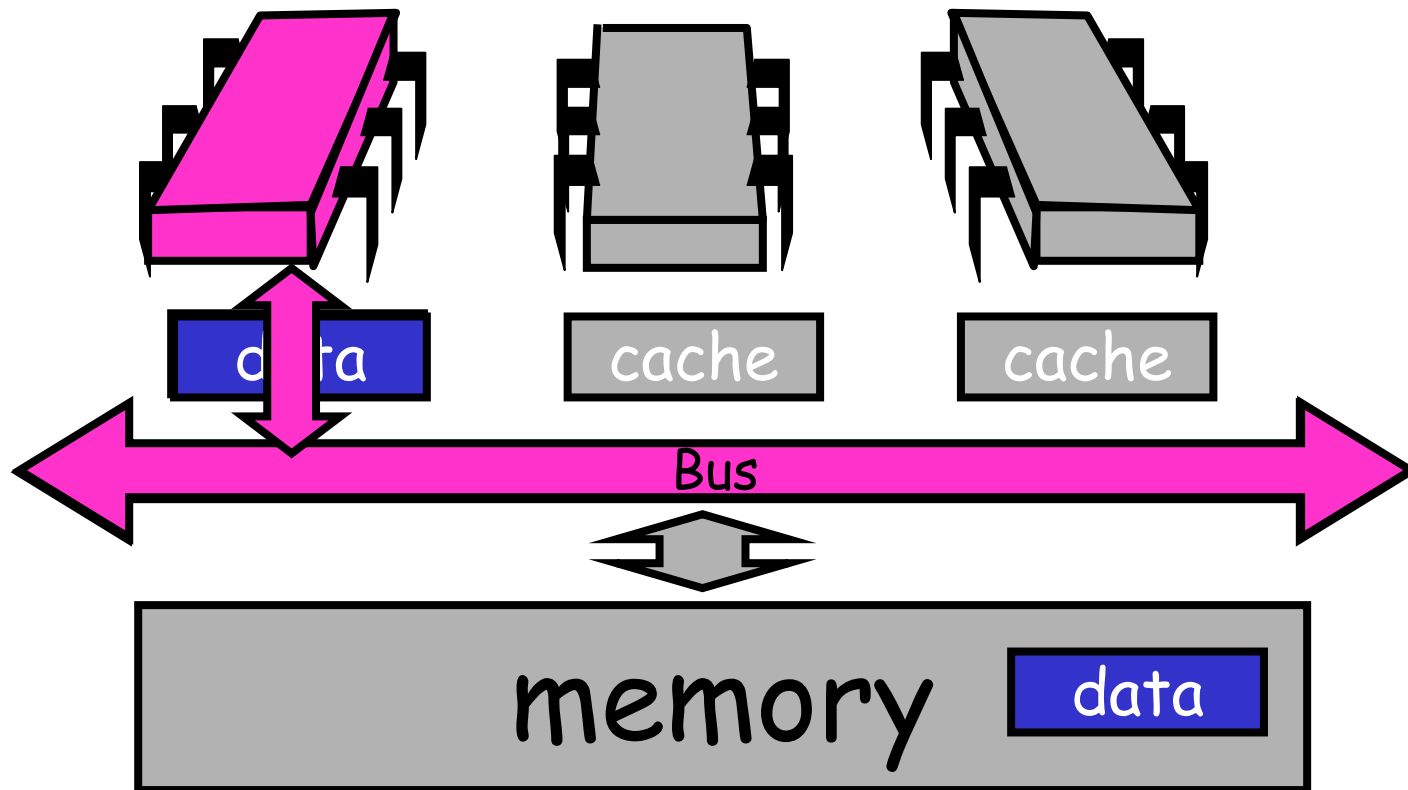
# Processor Issues Load Request



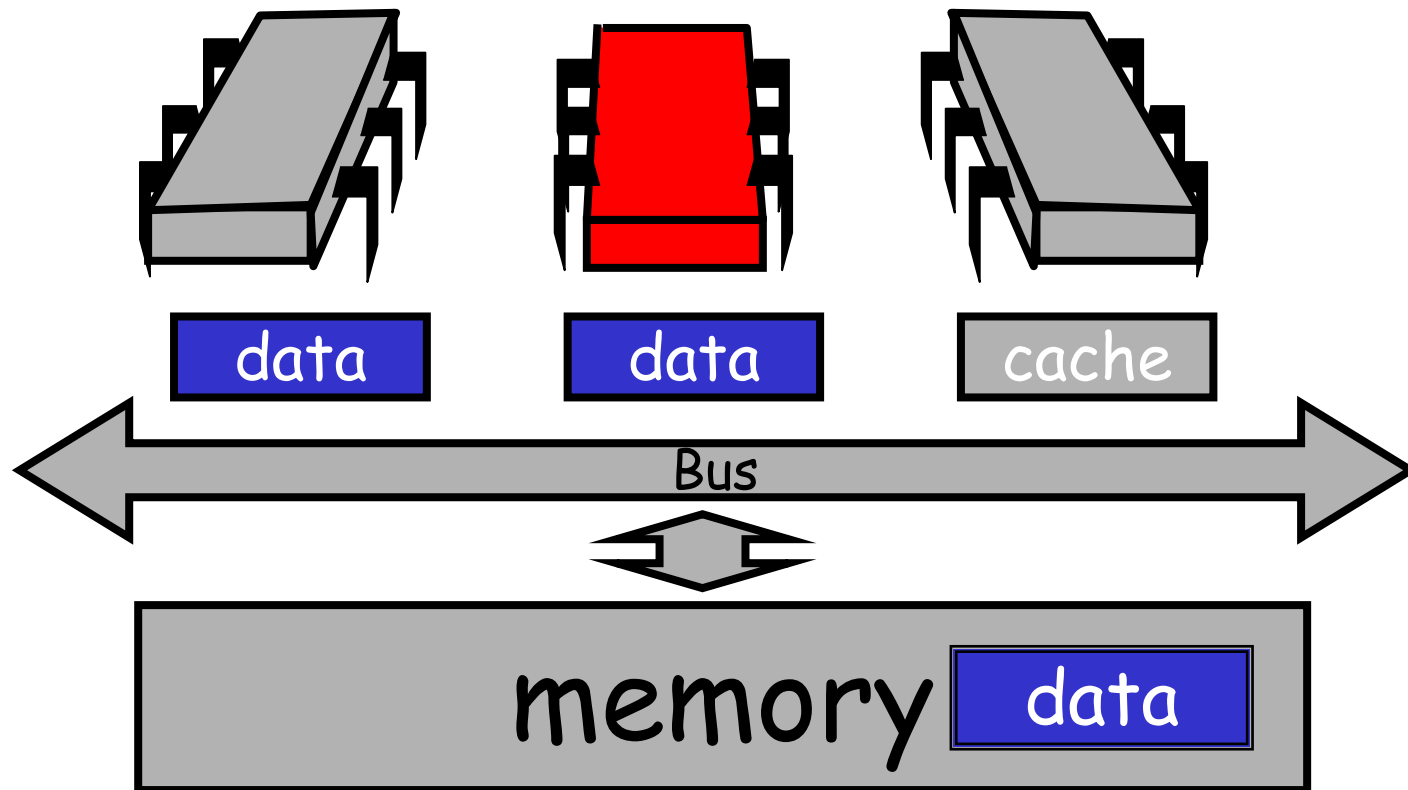
# Other Processor Responds



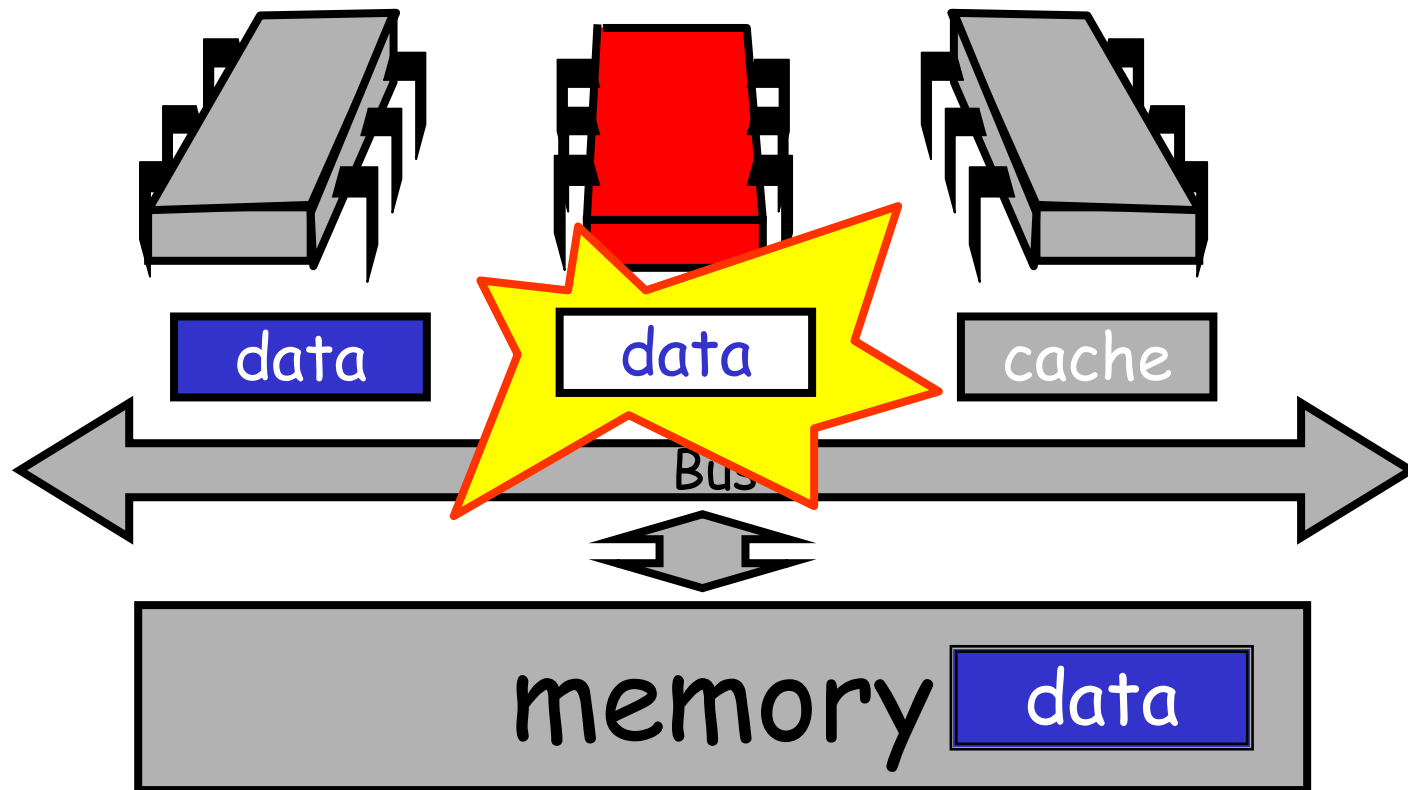
# Other Processor Responds



# Modify Cached Data

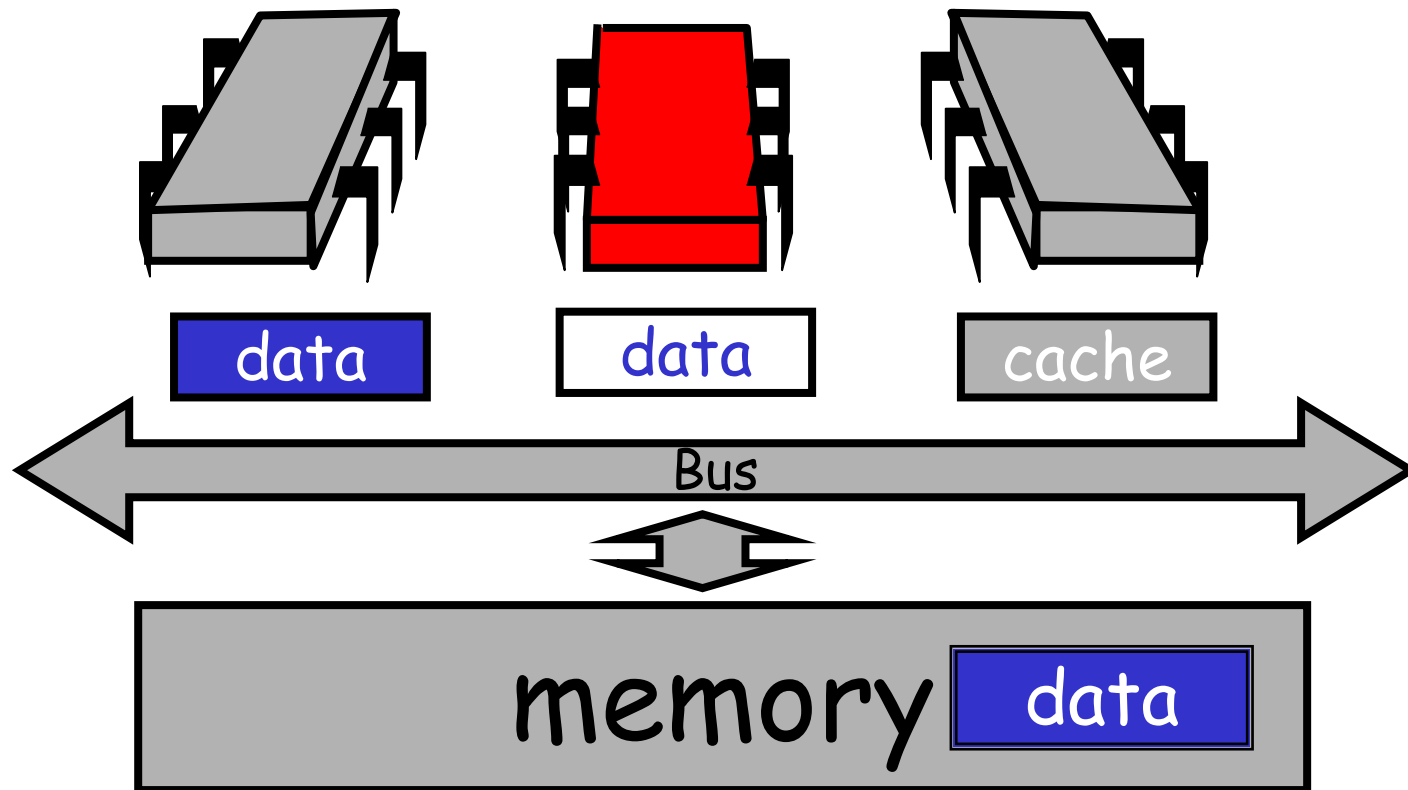


# Modify Cached Data






# Modify Cached Data



# Modify Cached Data





# Cache Coherence

- We have lots of copies of data
  - Original copy in memory
  - Cached copies at processors
- Some processor modifies its own copy
  - What do we do with the others?
  - How to avoid confusion?



# Write-Back Cache Coherence Protocol

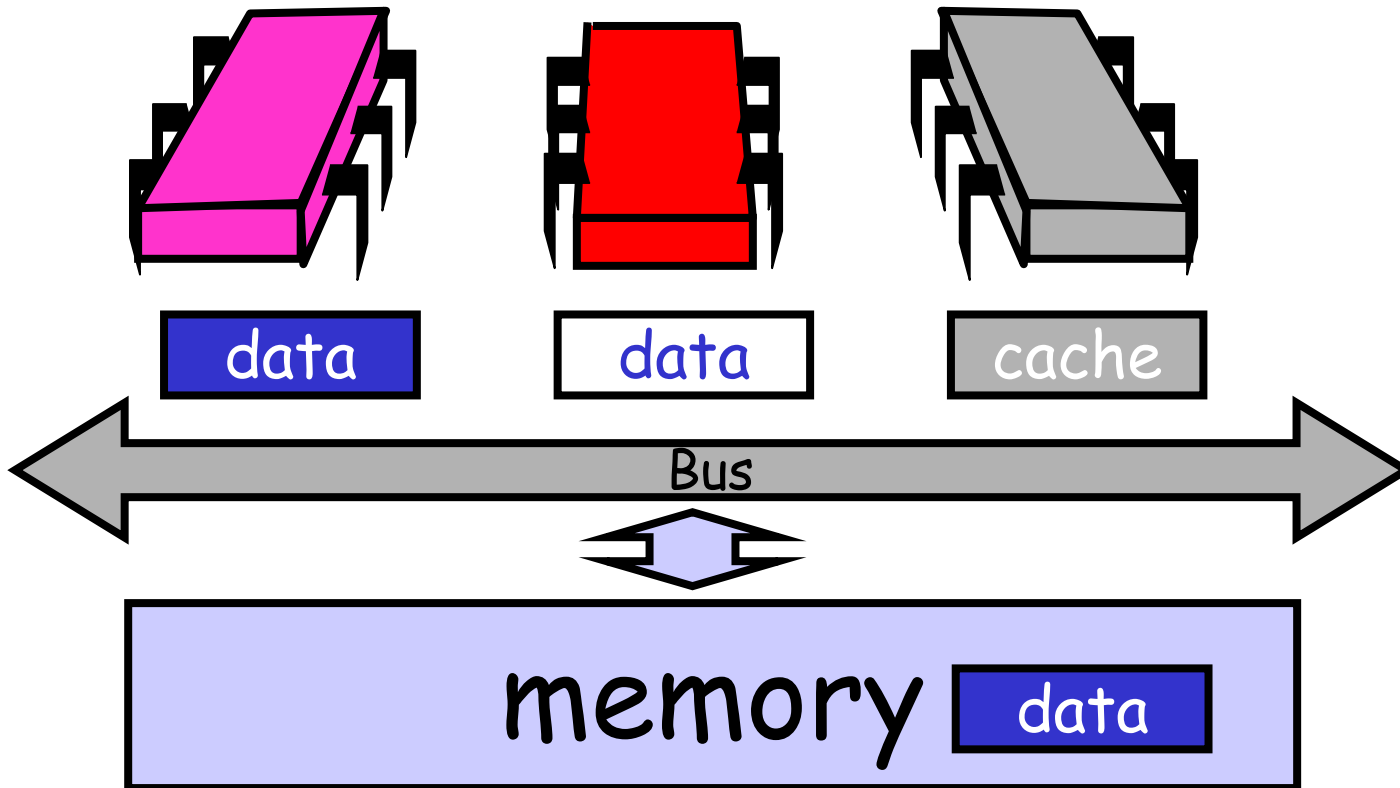
- Accumulate changes in cache
- Write back when needed
  - Need the cache for something else
  - Another processor wants it
- Write-back coherence protocol:
  - Invalidate other entries
  - Requires non-trivial protocol ...




# Write-Back Caches

- Cache entry has three states
  - Invalid: contains raw seething bits (meaningless)
  - Valid: I can read but I can't write because it may be cached elsewhere
  - Dirty: Data has been modified
    - Intercept other load requests
    - Write back to memory before using cache

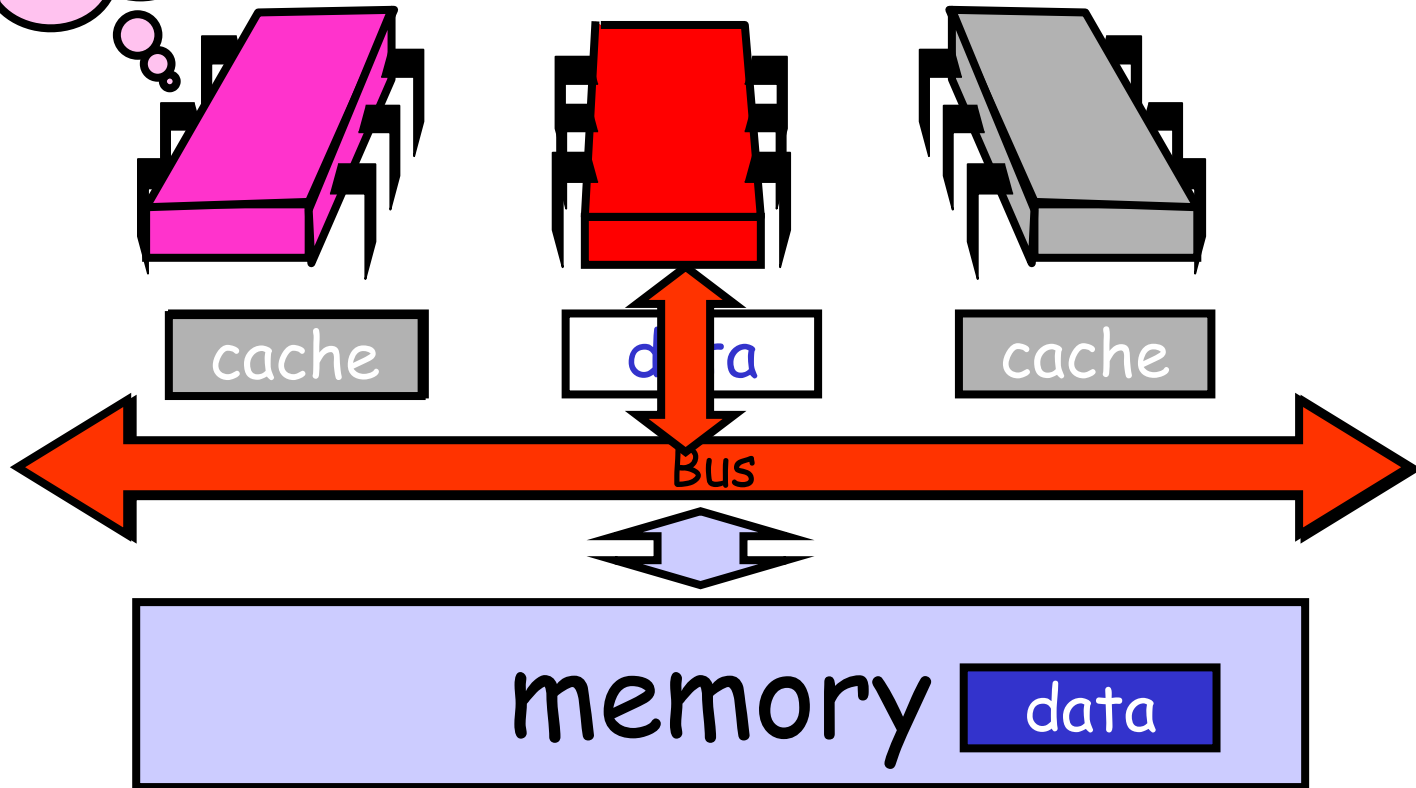
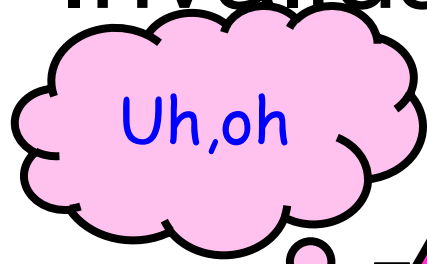
# Invalidate



# Invalidate



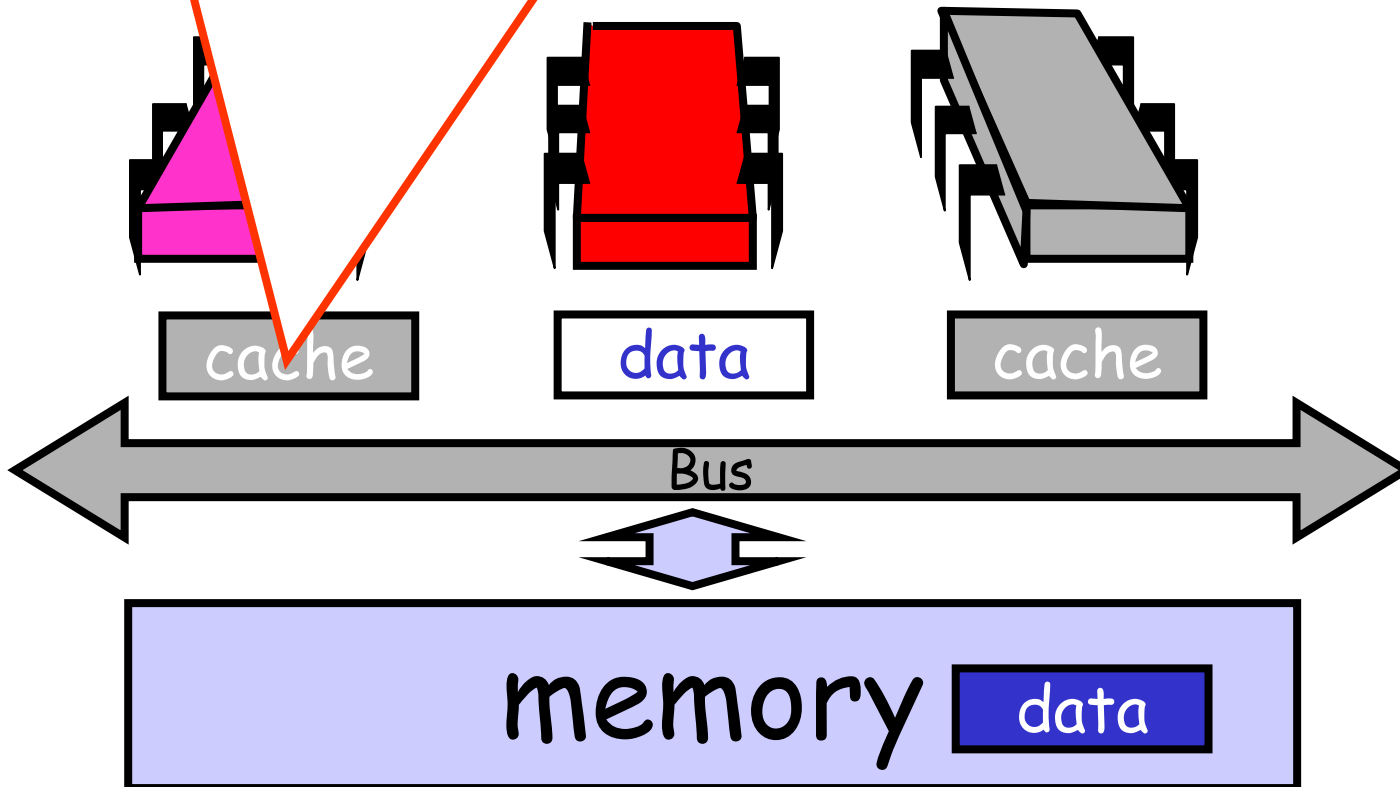
# Invalidate





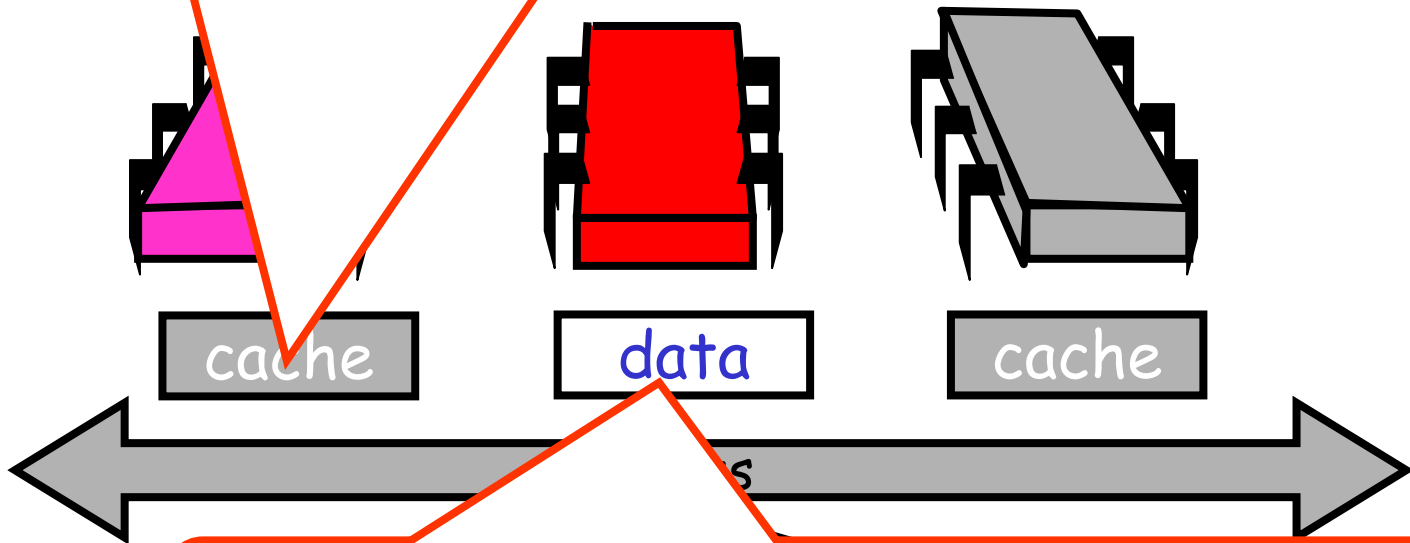
# Invalidate

Other caches lose read permission




# Invalidate

Other caches lose read permission

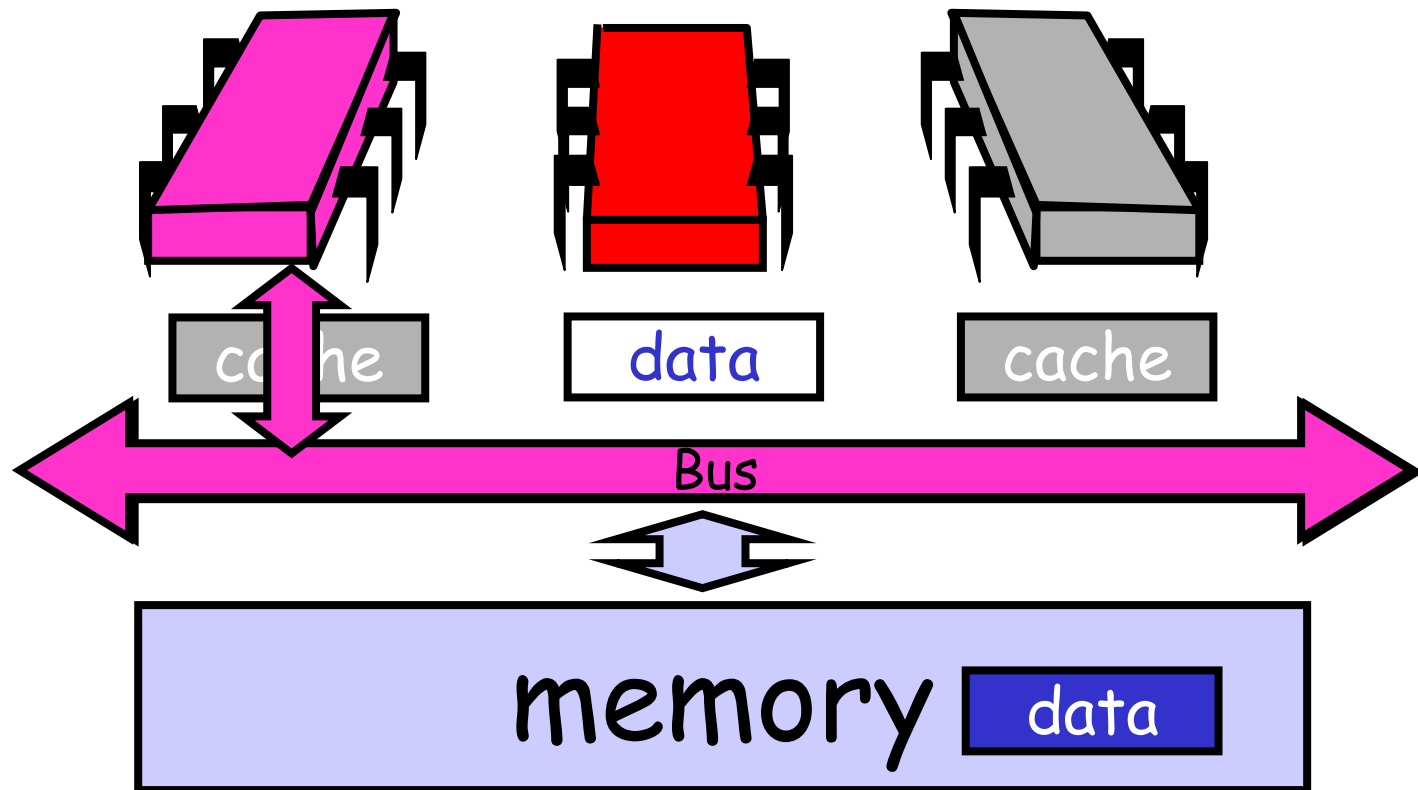


This cache acquires write permission

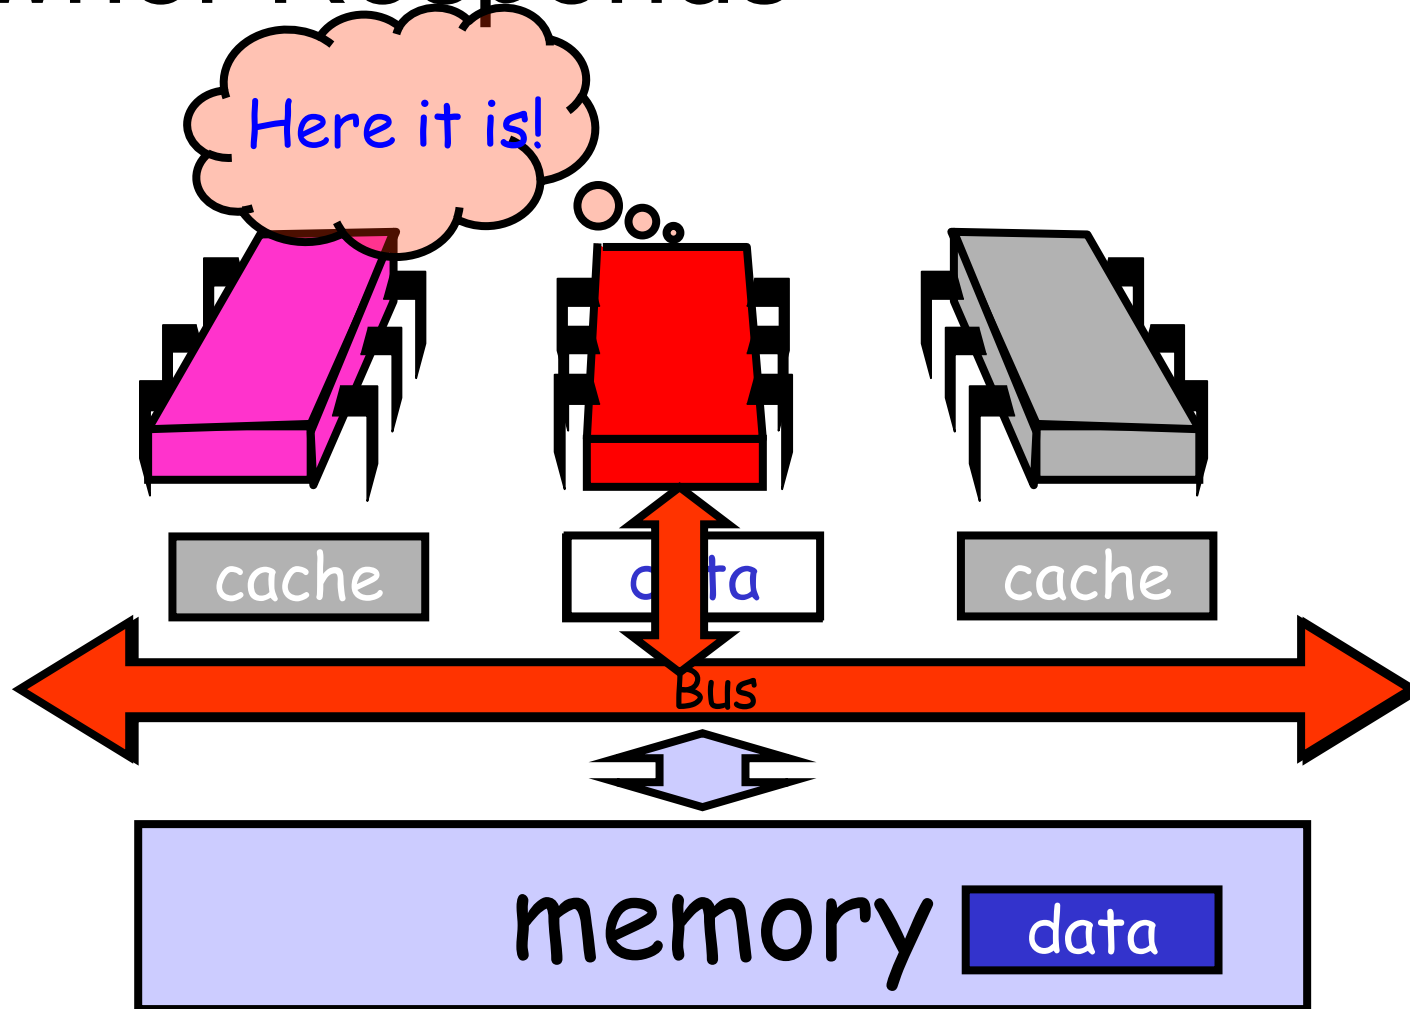
# Invalidate



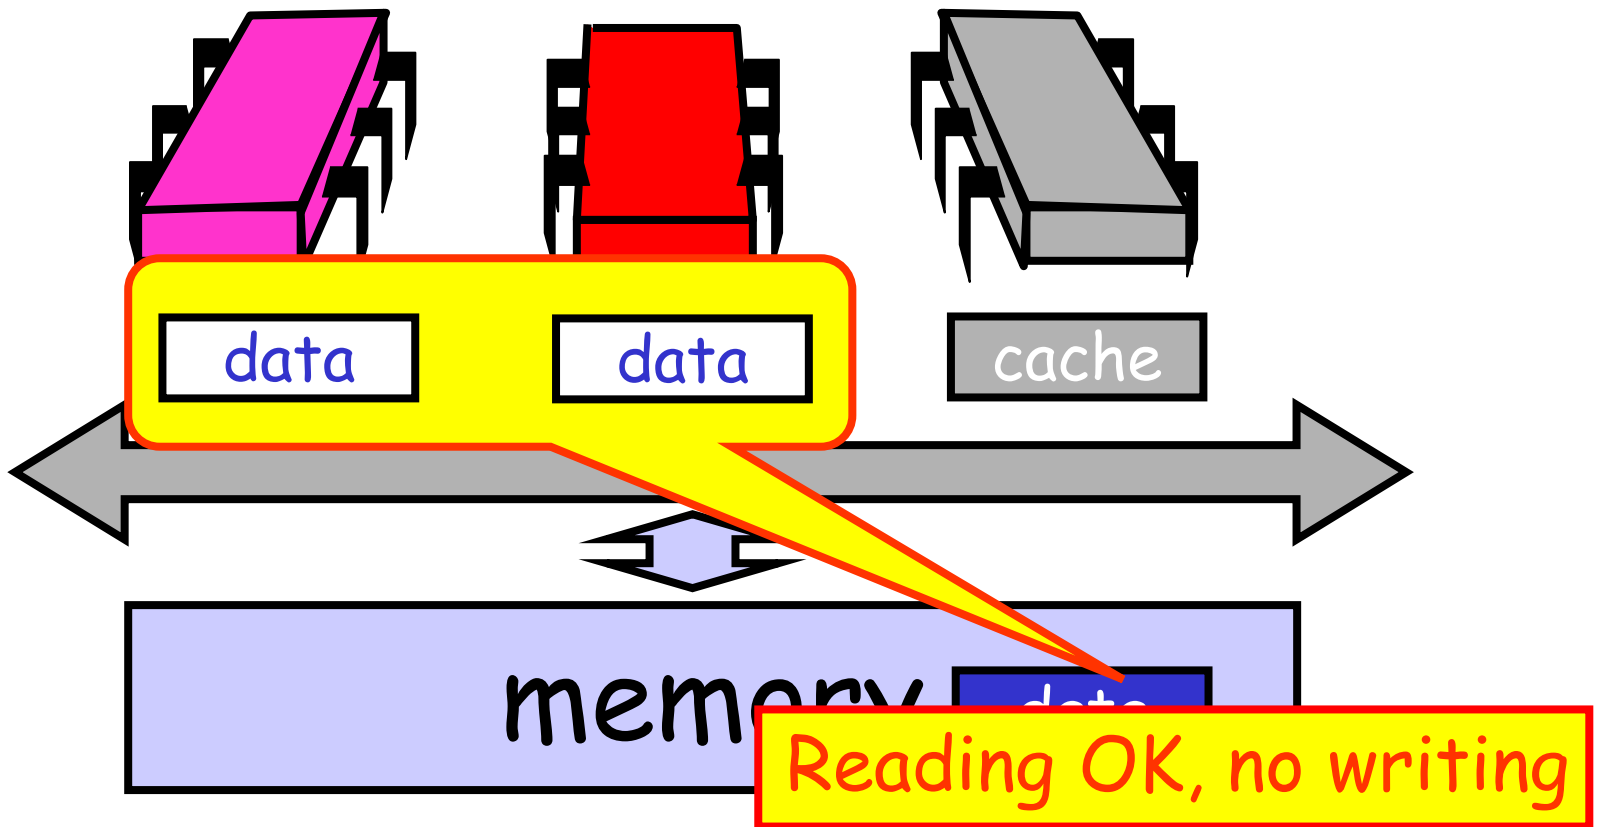
# Another Processor Asks for Data



# Owner Responds



# End of the Day ...



# Test-and-set Lock

```
class TASLock {  
    AtomicBoolean state =  
        new AtomicBoolean(false);  
  
    void lock() {  
        while (state.getAndSet(true)) {}  
    }  
  
    void unlock() {  
        state.set(false);  
    }  
}
```

Has to change to  
current value while  
spinning



# Back to TASLocks

- How does a TASLock perform on a write-back shared-bus architecture?
  - Because it uses the bus, each getAndSet() call delays all the other threads
    - Even those not waiting for the lock
  - The getAndSet() call forces the other processors to discard their own cached copies – resulting in a cache miss every time
  - They must then use the bus to fetch the new, but unchanged value





# TASLock

- When the thread wants to release the lock it may be delayed because the bus is being monopolized by the spinners



# Test-and-Test-and-Set Locks

## ■ Lurking stage

- ☐ Wait until lock “looks” free
- ☐ Spin while read returns true (lock taken)

## ■ Pouncing state

- ☐ As soon as lock “looks” available
- ☐ Read returns false (lock free)
- ☐ Call TAS to acquire lock
- ☐ If TAS loses, back to lurking

# Test-and-test-and-set Lock

```
class TTASLock {  
    AtomicBoolean state =  
        new AtomicBoolean(false);  
  
    void lock() {  
        while (true) {  
            while (state.get()) {}  
            if (!state.getAndSet(true))  
                return;  
        }  
    }  
}
```

# Test-and-test-and-set Lock

```
class TTASLock {  
    AtomicBoolean state =  
        new AtomicBoolean(false);  
  
    void lock() {  
        while (true) {  
            while (state.get()) {}  
            if (!state.getAndSet(true))  
                return;  
        }  
    }  
}
```

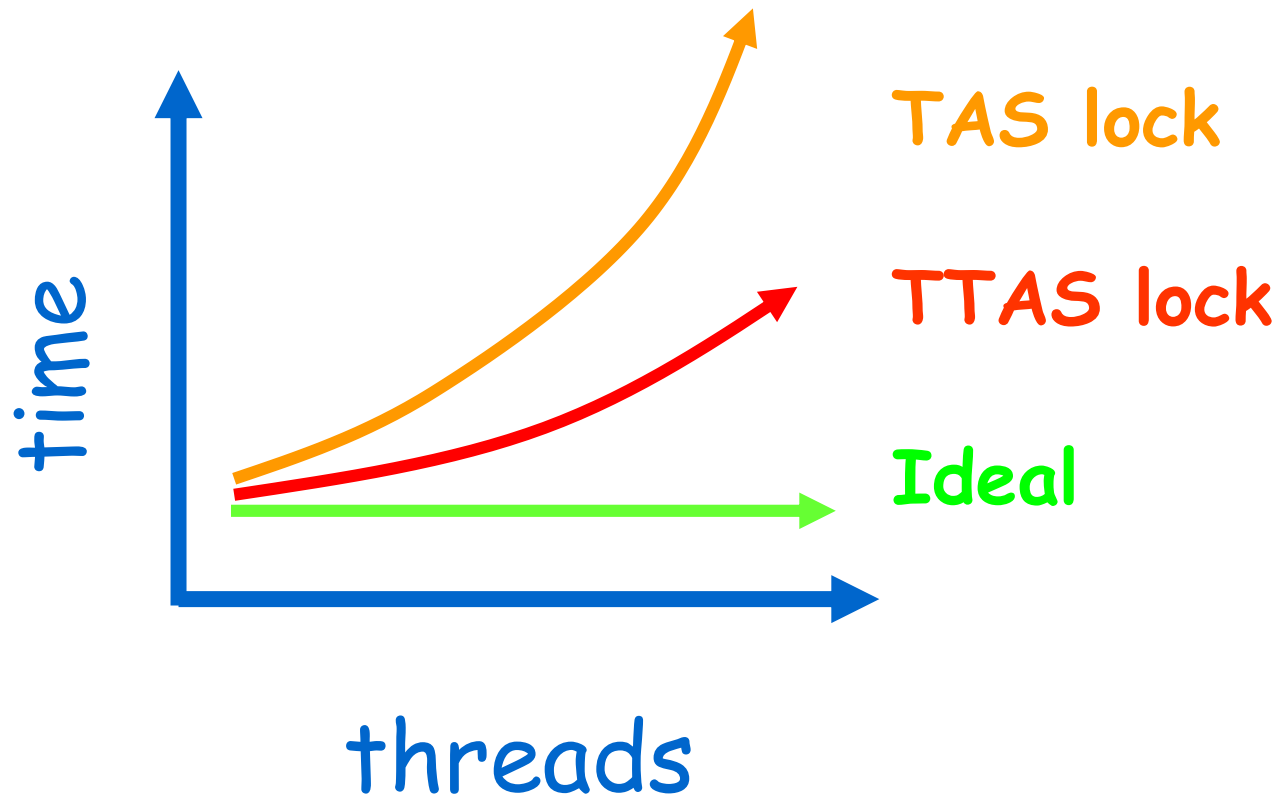
Wait until lock looks free

# Test-and-test-and-set Lock

```
class TTASLock {  
    AtomicBoolean state =  
        new AtomicBoolean(false);  
  
    void lock() {  
        while (true) {  
            while (state.get()) {}  
            if (!state.getAndSet(true))  
                return;  
        }  
    }  
}
```

Then try to  
acquire it

# Mystery #2





# What about the TTASLock?

- Suppose thread A acquires the lock.
- The first time thread B reads the lock it takes a cache miss and has to use the bus to fetch the new value
- As long as A holds the lock however, B repeatedly rereads the value – resulting in a cache hit every time
- B thus produces no extra traffic



# What about the TTASLock?

- However when A releases the lock:
  - A writes false to the lock variable
  - The spinner's cached copies are invalidated
  - Each one takes a cache miss
  - They all use the bus to read a new value
  - They all call getAndSet() to acquire the lock
  - The first one to acquire the lock invalidates the others who must then reread the value
  - Storm of traffic





# Local spinning

- Threads repeatedly reread cached values instead of repeatedly using the bus



# Exponential Backoff

- Recall that in the TTASLock, the thread first reads the lock and if it appears to be free it attempts to acquire the lock
- “If I see that the lock is free, but then another thread acquires it before I can, then there must be high contention for that lock”
- Better to back off and try again later



# For how long should a thread back off?

- Rule of thumb:

- The larger number of unsuccessful tries, the higher the contention, the longer the thread should back off.

- Approach:

- Whenever the thread sees the lock has become free, but fails to acquire it, it backs off before retrying



# What about lock-step?

- What happens if all the threads backs off the same amount of time?
- Instead the threads should back off for a random amount of time
- Each time the thread tries and fails to get the lock, it doubles the back-off time, up to a fixed maximum.



# Exponential Backoff Lock

```
public class Backoff implements Lock {  
    public void lock() {  
        int delay = MIN_DELAY;  
        while (true) {  
            while (state.get()) {}  
            if (!lock.getAndSet(true))  
                return;  
            sleep(random() % delay);  
            if (delay < MAX_DELAY)  
                delay = 2 * delay;  
        }  
    }  
}
```

# Exponential Backoff Lock

```
public class Backoff implements Lock {  
    public void lock() {  
        int delay = MIN_DELAY;  
        while (true) {  
            while (state.get()) {}  
            if (!lock.getAndSet(true))  
                return;  
            sleep(random() % delay);  
            if (delay < MAX_DELAY)  
                delay = 2 * delay;  
        }  
    }  
}
```

Fix minimum delay

# Exponential Backoff Lock

```
public class Backoff implements Lock {  
    public void lock() {  
        int delay = MIN_DELAY;  
        while (true) {  
            while (state.get()) {}  
            if (!lock.getAndSet(true))  
                return;  
            sleep(random() % delay);  
            if (delay < MAX_DELAY)  
                delay = 2 *  
            Wait until lock looks free  
        }  
    }  
}
```

# Exponential Backoff Lock

```
public class Backoff implements Lock {  
    public void lock() {  
        int delay = MIN_DELAY;  
        while (true) {  
            while (state.get()) {}  
            if (!lock.getAndSet(true))  
                return;  
            sleep(random() % delay);  
            if (delay < MAX_DELAY)  
                delay = 2 * delay;  
        }  
    }  
}
```

If we win, return



# Exponential Backoff Lock

```
public class Backoff implements Lock {  
    public void lock() {  
        int delay = MIN_DELAY;  
        while (true) {  
            while (state.get()) {}  
            if (!lock.getAndSet(true))  
                return;  
            sleep(random() % delay);  
            if (delay < MAX_DELAY)  
                delay = 2 * delay;  
        }  
    }  
}
```

Otherwise back off for random duration

# Exponential Backoff Lock

```
public class Backoff implements Lock {  
    public void lock() {  
        int delay = MIN_DELAY;  
        while (true) {  
            while (state.get()) {}  
            if (!lock.getAndSet(true))  
                return;  
            sleep(random() % delay);  
            if (delay < MAX_DELAY)  
                delay = 2 * delay;  
        }  
    }  
}
```

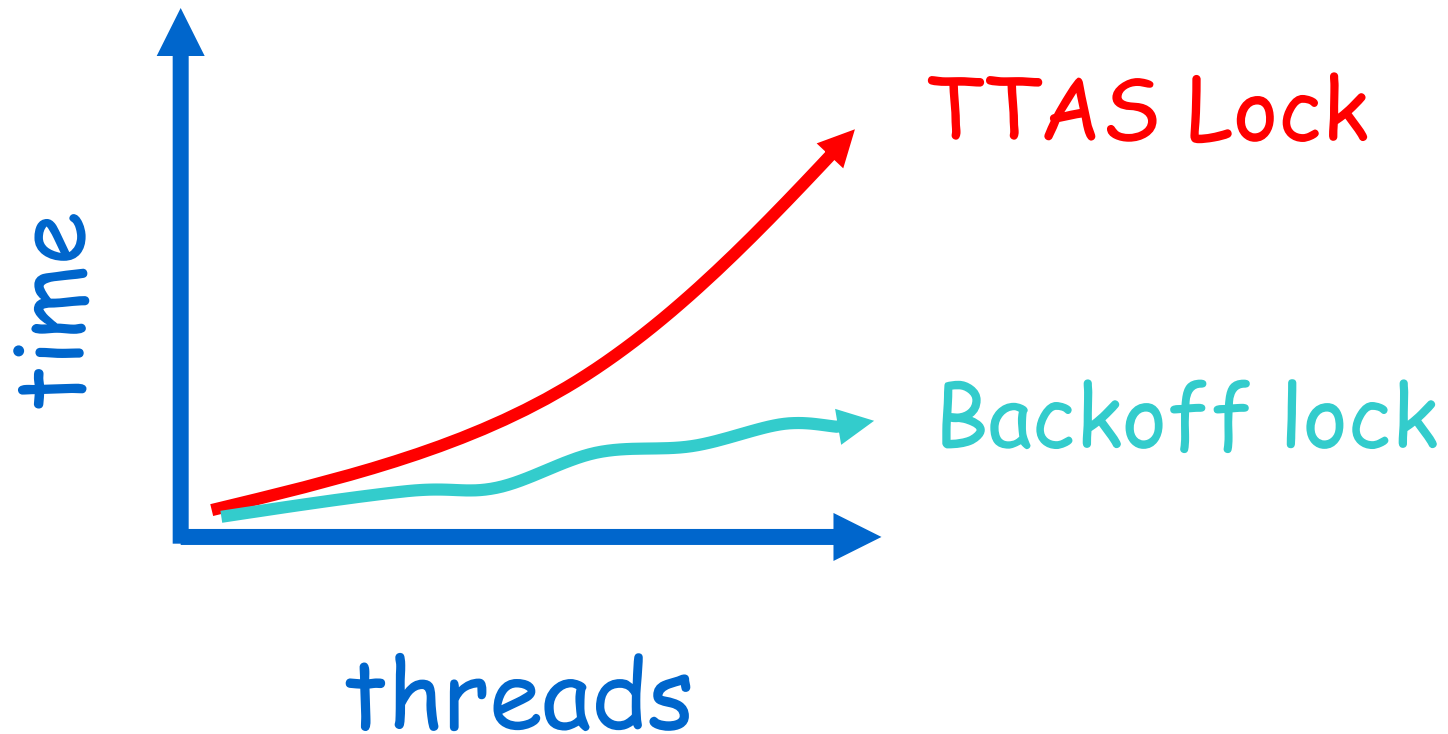
Double max delay,  
within reason



# Exponential Backoff Lock

- Important to note that a thread backs off only when it fails to acquire a lock that was just available, but is not available anymore
- Observing that a lock is held by another thread does not imply backoff

# Spin-Waiting Overhead





# Backoff: Other Issues

## ■ Good

- Easy to implement
- Beats TTAS lock

## ■ Bad

- Must choose parameters carefully
  - Sensitive to choice of minimum and maximum delays
- Sensitive to number of processors and their speed
  - Cannot have a general solution for all platforms and machines



# BackoffLock drawbacks

- Cache-coherence Traffic:
  - All threads spin on the same location
- Critical Section Underutilization:
  - Threads delay longer than necessary



# Idea

- Avoid useless invalidations
  - By keeping a queue of threads
- Each thread
  - Notifies next in line
  - Without bothering the others

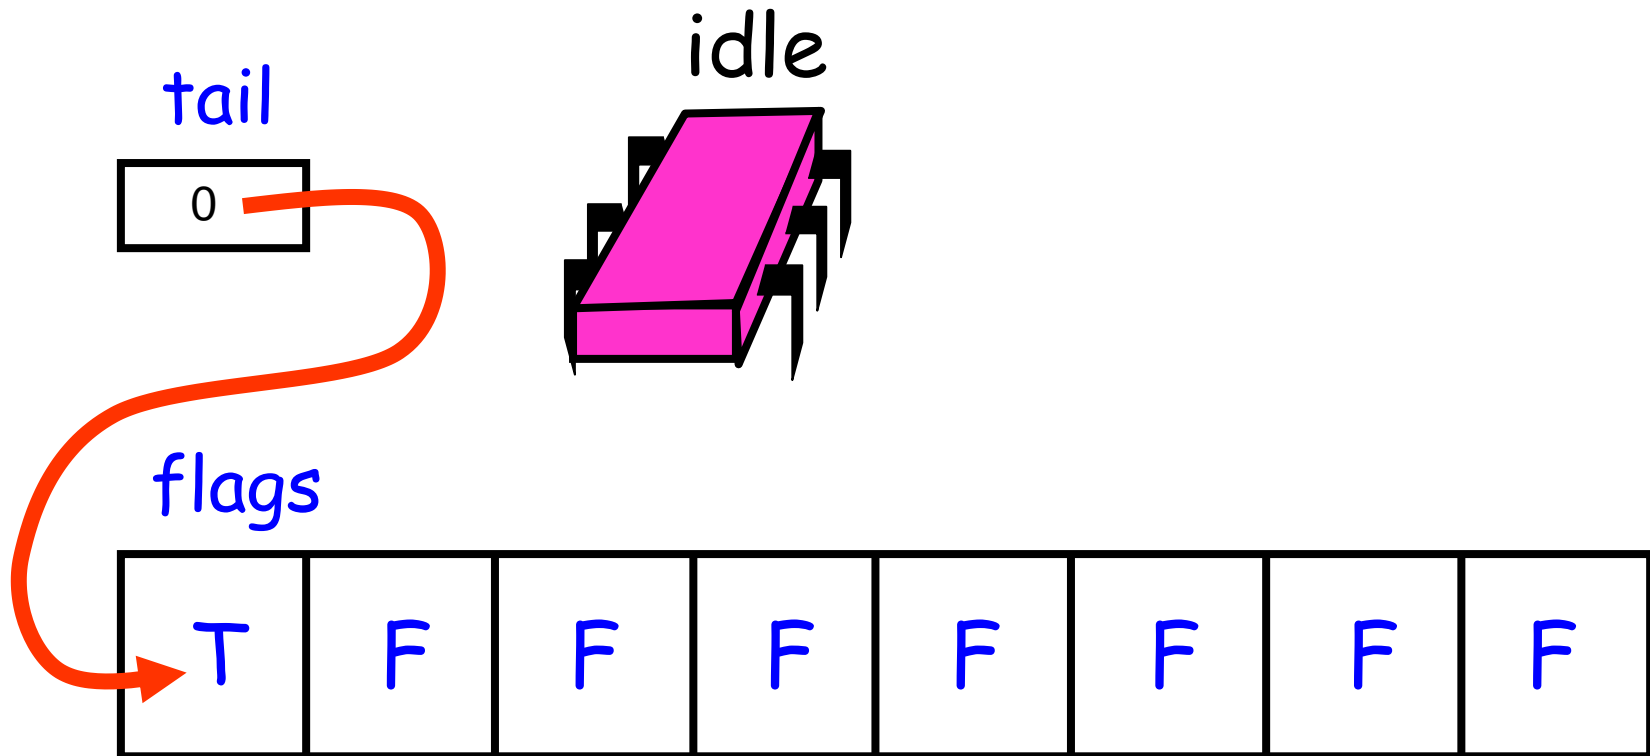


# Queue Locks

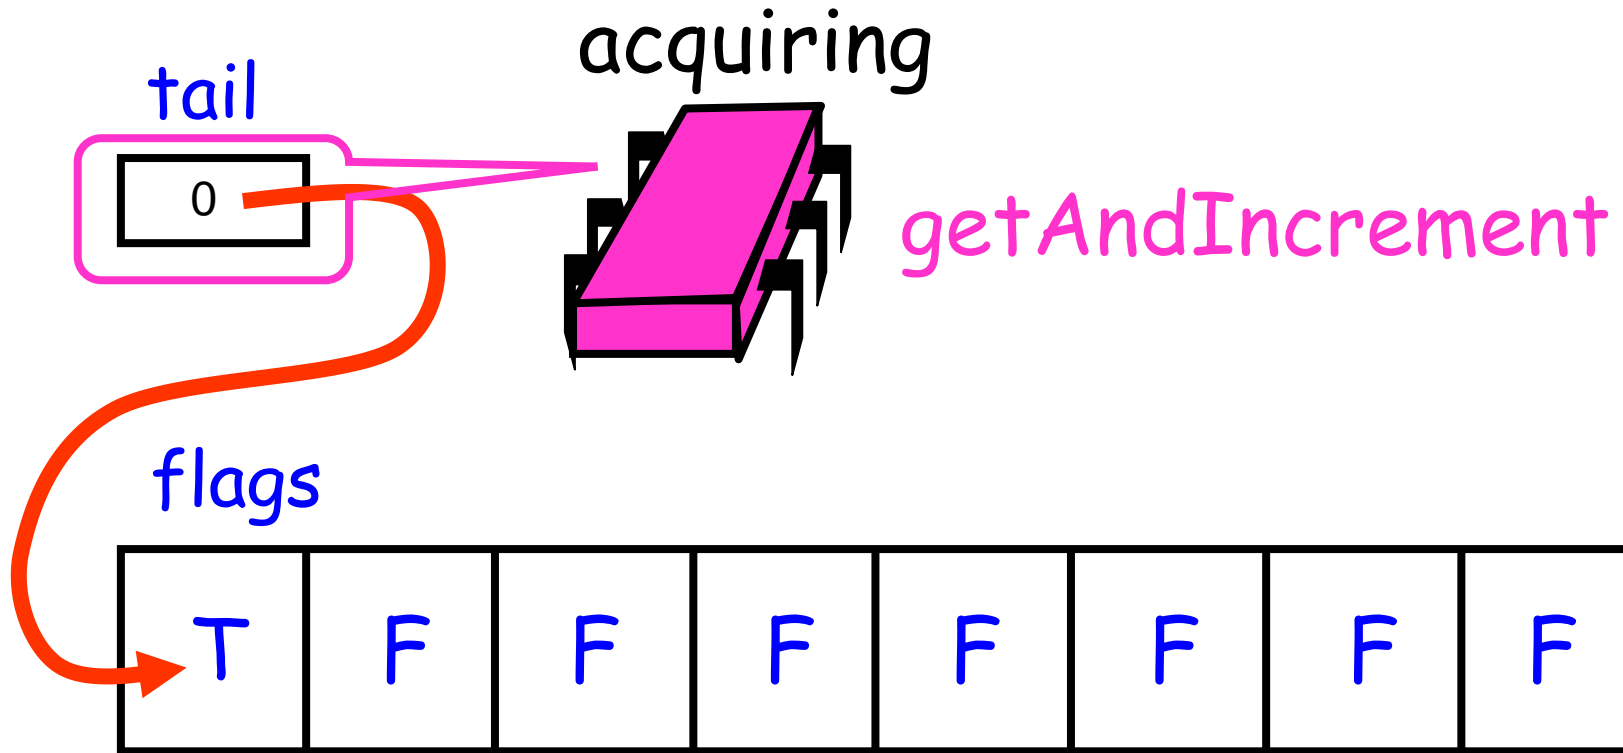
- Cache-coherence traffic is reduced since each thread spins on a different location
- No need to guess when to attempt to access lock – increase critical section utilization
- First-come-first-served fairness



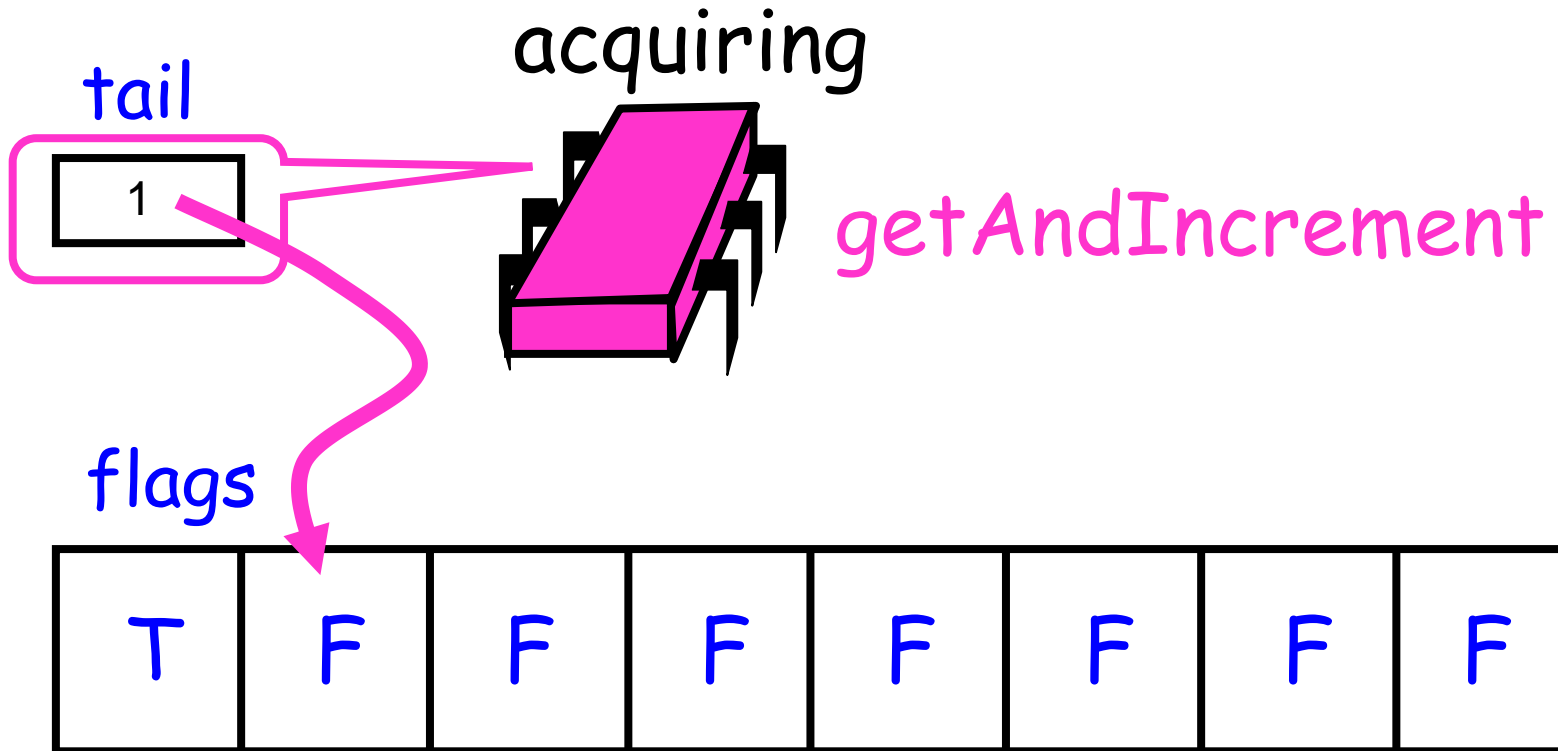
# Anderson Queue Lock



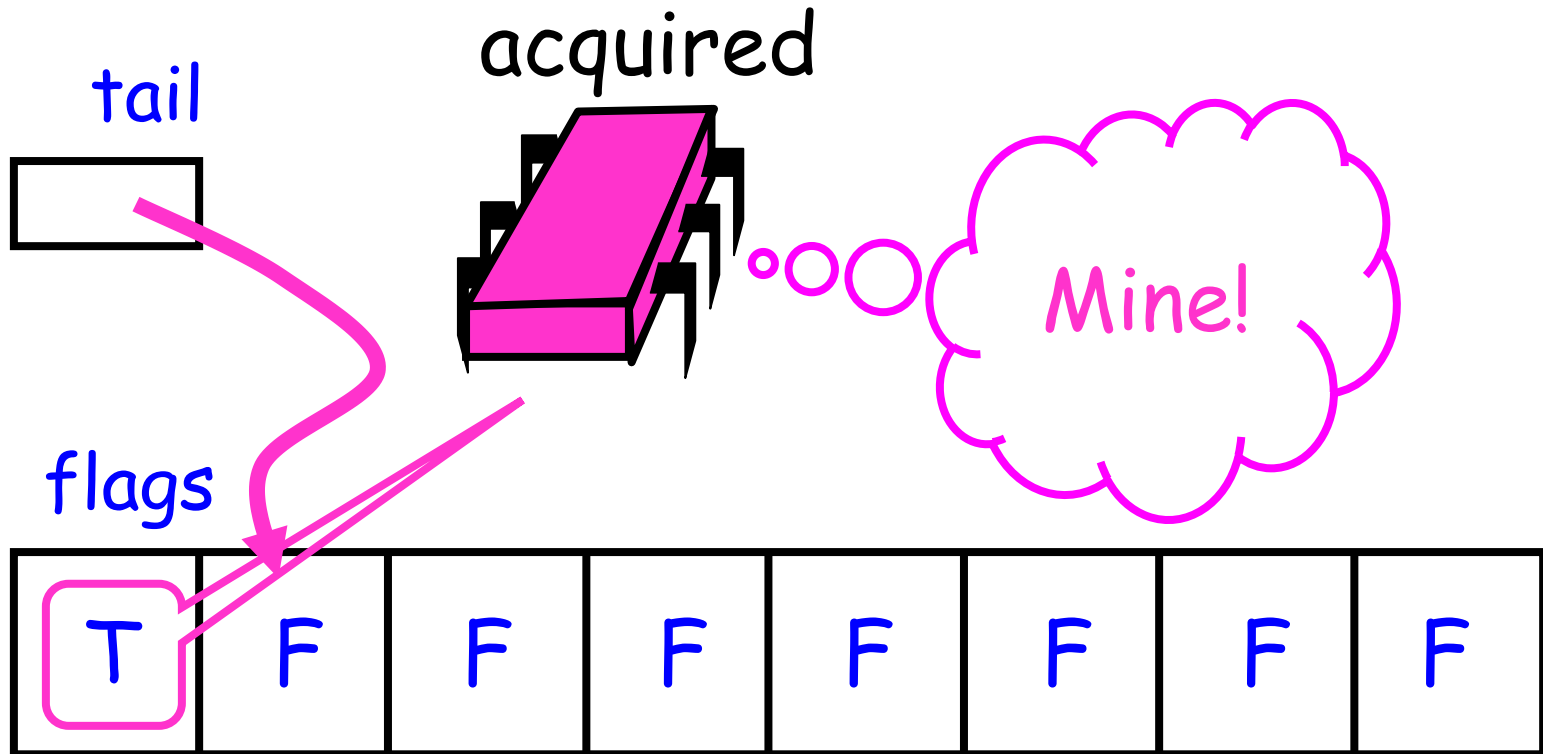
# Anderson Queue Lock



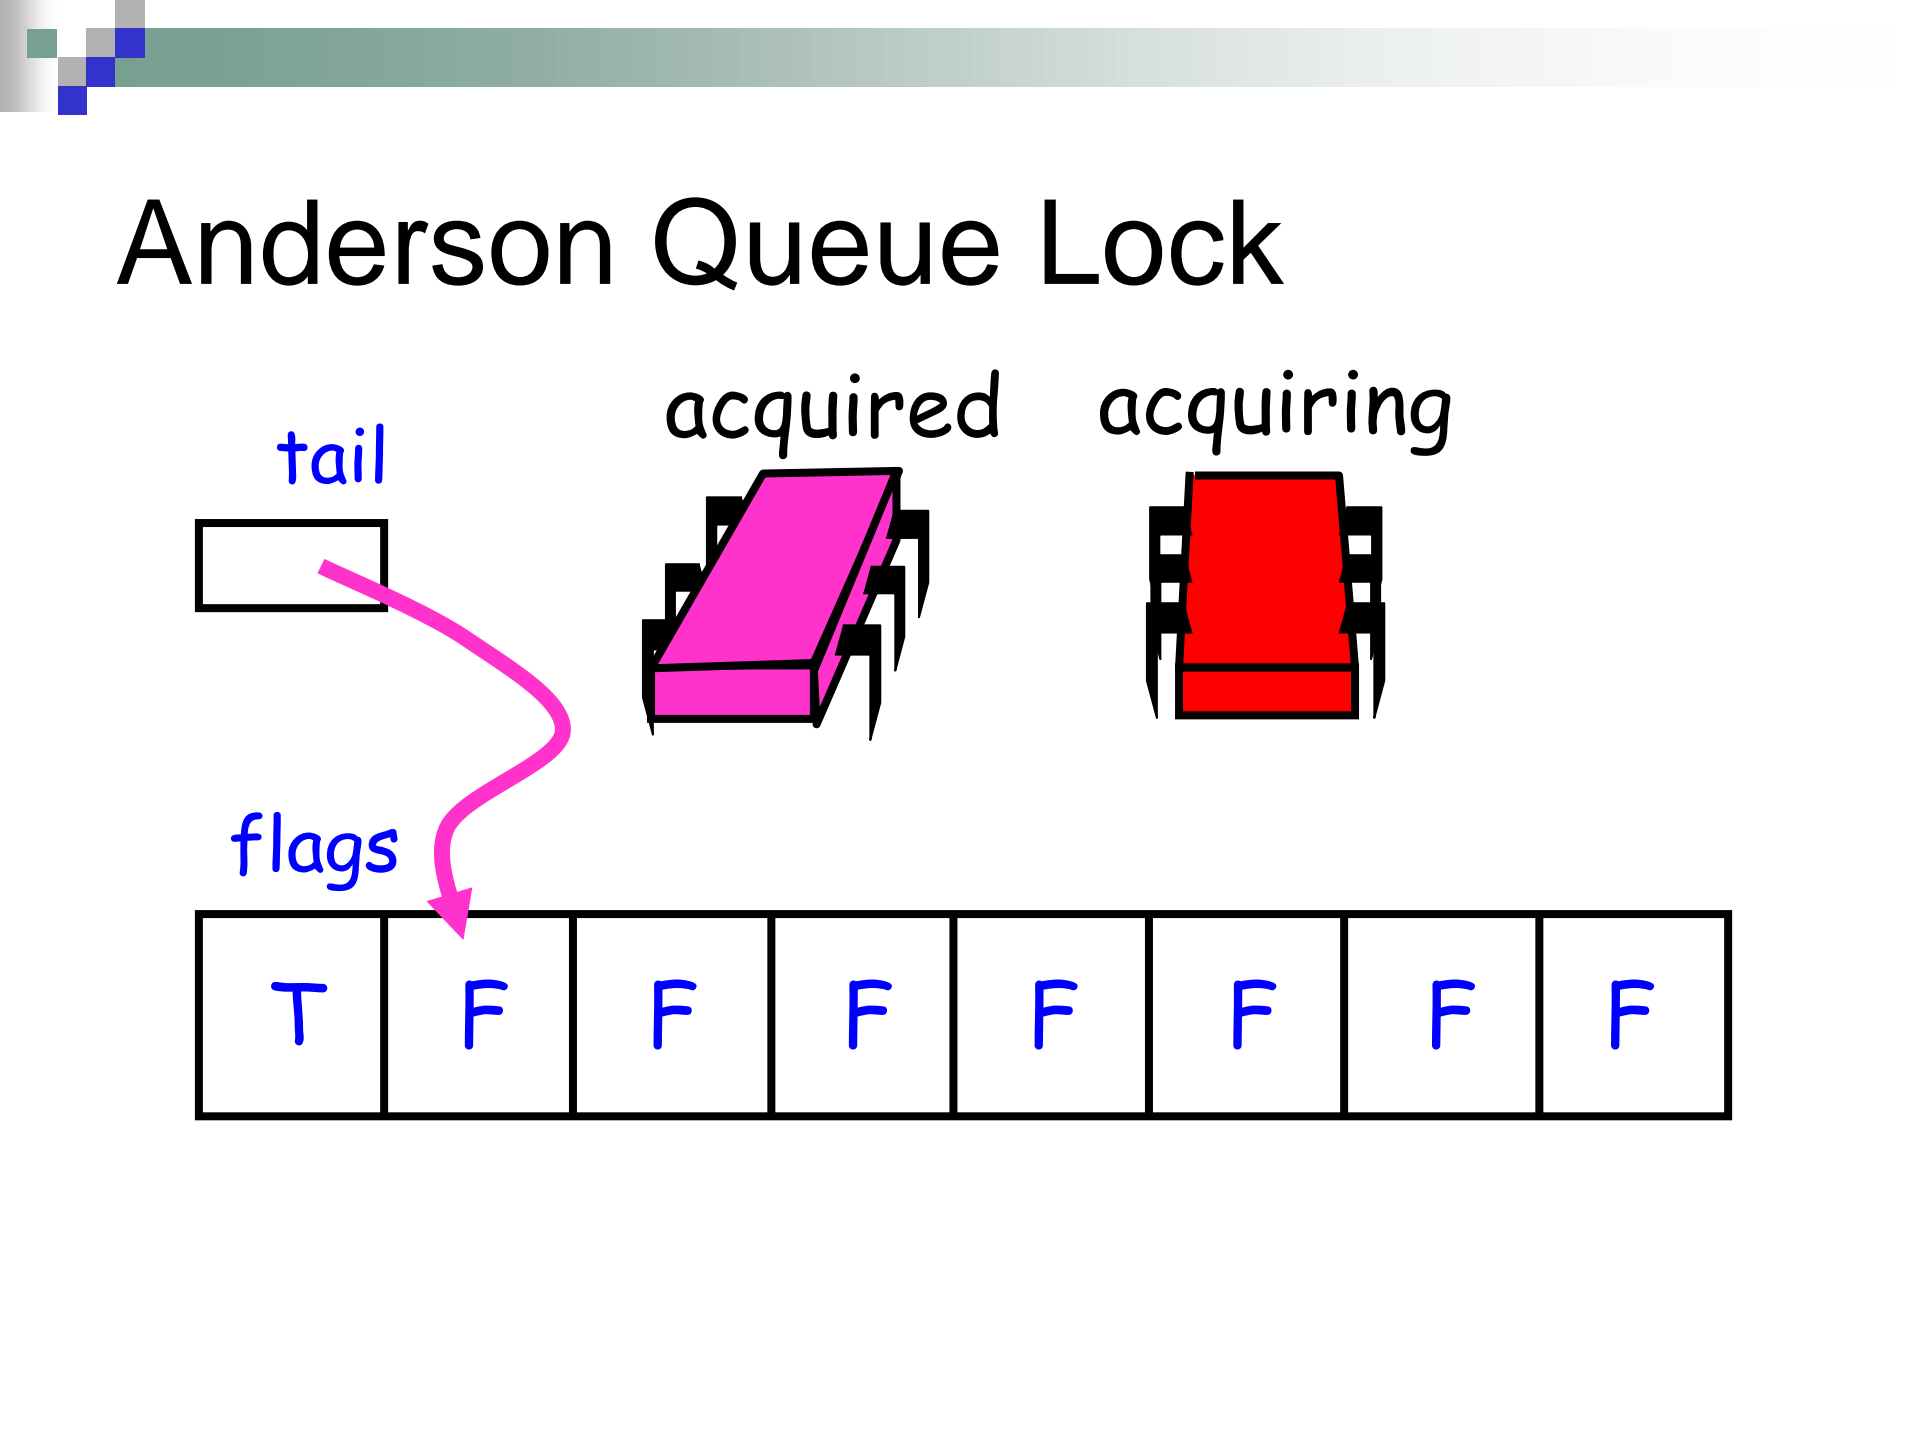
# Anderson Queue Lock



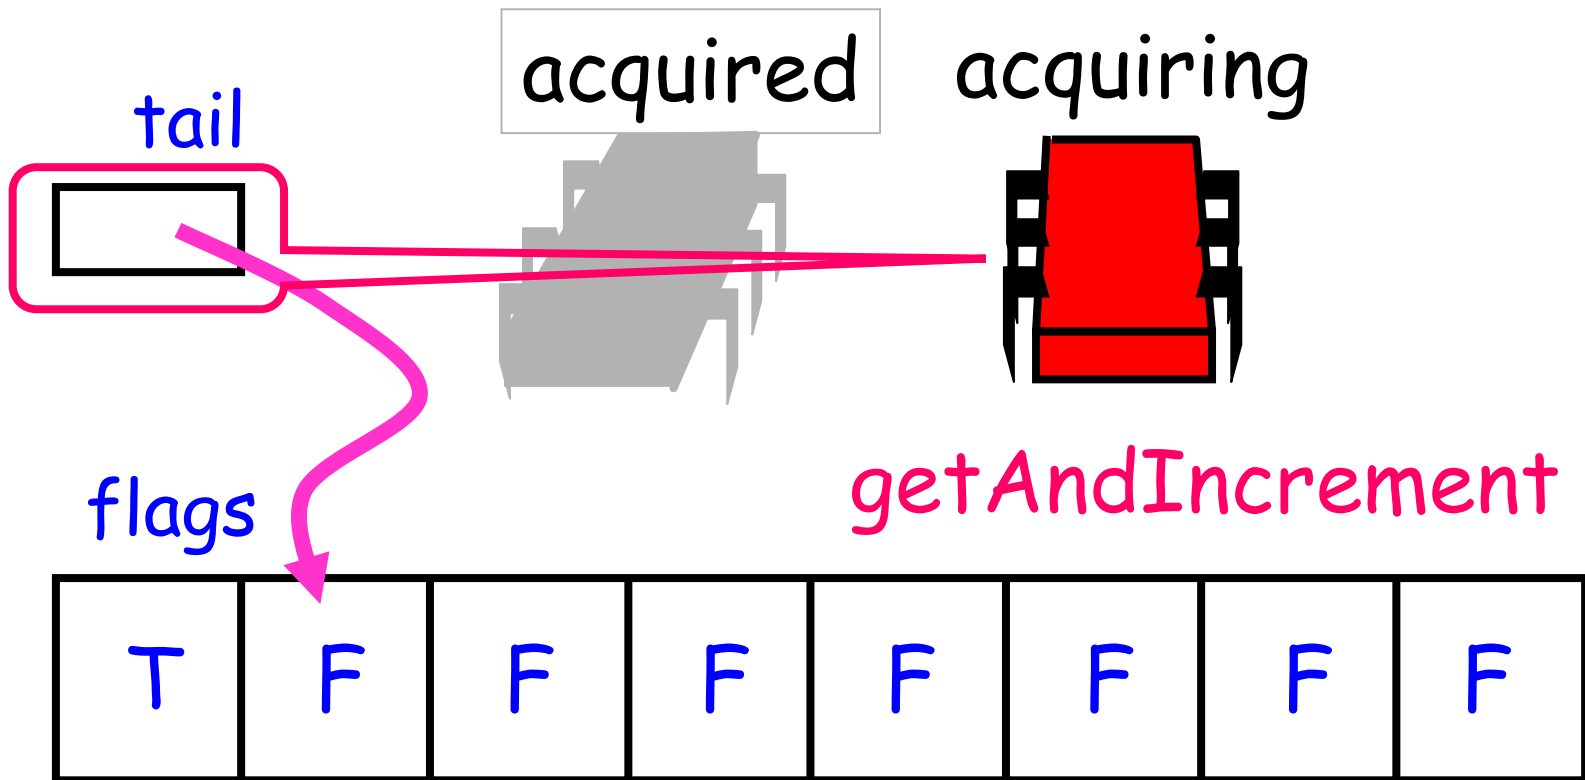
# Anderson Queue Lock




Age Group	Percentage
18-24	10%
25-34	15%
35-44	20%
45-54	25%
55-64	30%
65-74	35%
75-84	40%
85+	45%




# Anderson Queue Lock



# Anderson Queue Lock




# Anderson Queue Lock






# Anderson Queue Lock



# Anderson Queue Lock





# Anderson Queue Lock

```
class ALock implements Lock {  
    boolean[] flags={true,false,...,false};  
    AtomicInteger tail  
        = new AtomicInteger(0);  
    ThreadLocal<Integer> mySlot;
```

# Anderson Queue Lock

```
class Alock implements Lock {  
    boolean[] flags={true,false,...,false};  
    AtomicInteger tail  
        = new AtomicInteger(0);  
    ThreadLocal<Integer> mySlot;
```

One flag per thread

# Anderson Queue Lock

```
class ALock implements Lock {  
    boolean[] flags={true,false,...,false};  
    AtomicInteger tail  
        = new AtomicInteger(0);  
    ThreadLocal<Integer> mySlot;
```

Next flag to use

# Anderson Queue Lock

```
class ALock implements Lock {  
    boolean[] flags={true,false,...,false};  
    AtomicInteger tail  
        = new AtomicInteger(0);  
    ThreadLocal<Integer> mySlot;
```

Thread-local variable

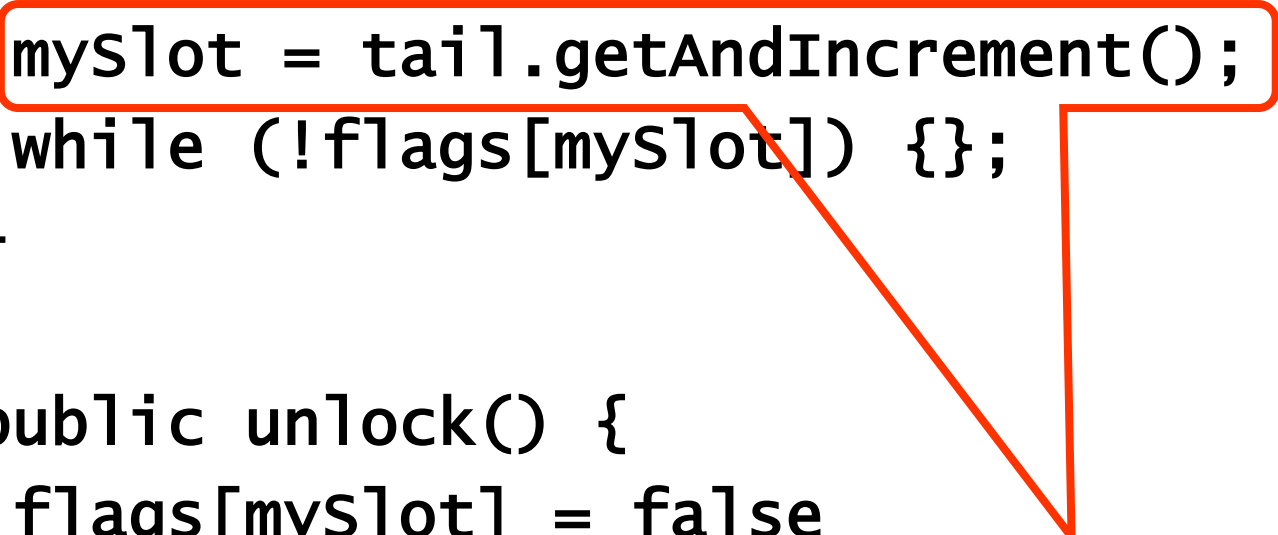


# Anderson Queue Lock

```
public lock() {  
    mySlot = tail.getAndIncrement();  
    while (!flags[mySlot]) {};  
}  
  
public unlock() {  
    flags[mySlot] = false;  
    flags[(mySlot+1)] = true;  
}
```

# Anderson Queue Lock

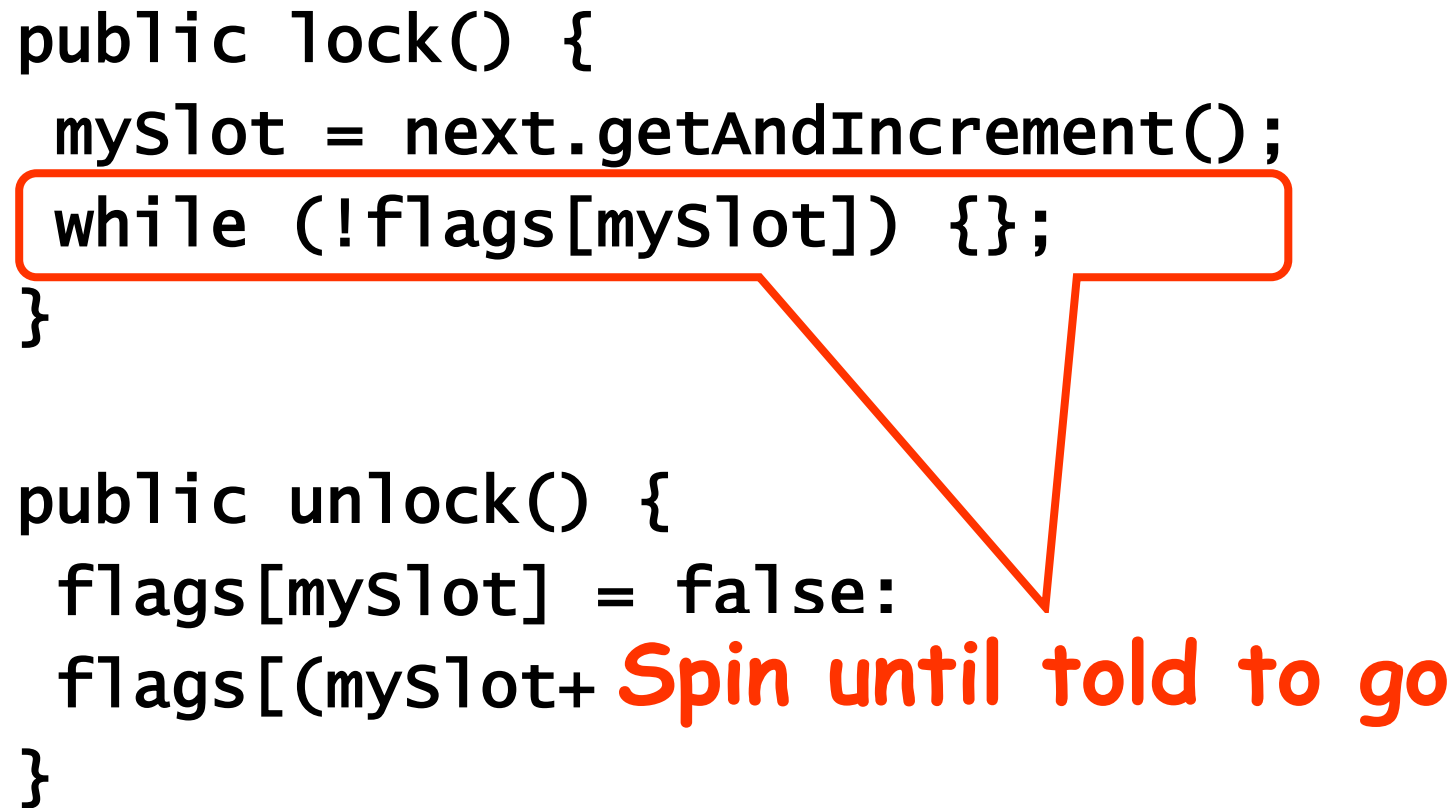
```
public lock() {  
    mySlot = tail.getAndIncrement();  
    while (!flags[mySlot]) {};  
}  
  
public unlock() {  
    flags[mySlot] = false  
    flags[(mySlot+1)] = Take next slot  
}
```

A red line originates from the `mySlot = tail.getAndIncrement();` line in the `lock()` method, extends horizontally to the right, and then diagonally down and to the left, pointing towards the `flags[(mySlot+1)] = Take next slot` line in the `unlock()` method. This indicates that the value of `mySlot` is used to determine the next slot to be taken when the lock is released.



# Anderson Queue Lock

```
public lock() {  
    mySlot = next.getAndIncrement();  
    while (!flags[mySlot]) {};  
}  
  
public unlock() {  
    flags[mySlot] = false;  
    flags[(mySlot+1) % flags.length] = true;  
    // Spin until told to go  
}
```

A red rectangular box highlights the while loop in the lock() method. A red line extends from the bottom of this box, pointing towards the comment 'Spin until told to go' in the unlock() method, indicating that the thread spins during this period.

# Anderson Queue Lock

```
public lock() {  
    myslot = next.getAndIncrement();  
    while (!flags[myslot]) {};  
}
```

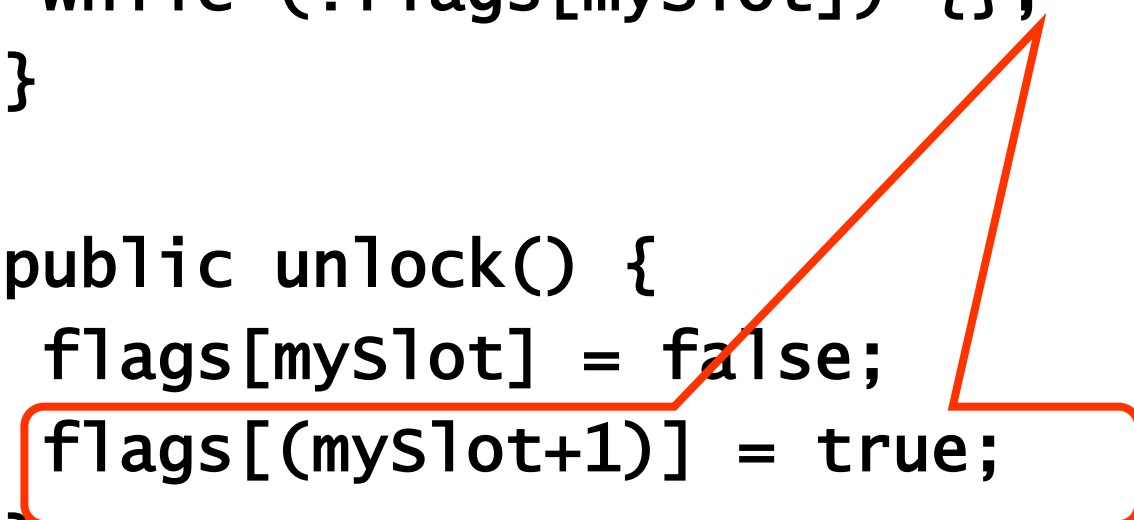
```
public unlock() {  
    flags[myslot] = false;  
    flags[(myslot+1)] = true;  
}
```

Prepare slot for re-use

# Anderson Queue Lock

```
public lock() {  
    mySlot = nextSlot;  
    while (!flags[mySlot]) {}  
}  
  
public unlock() {  
    flags[mySlot] = false;  
    flags[(mySlot+1)] = true;  
}
```

**Tell next thread to go**

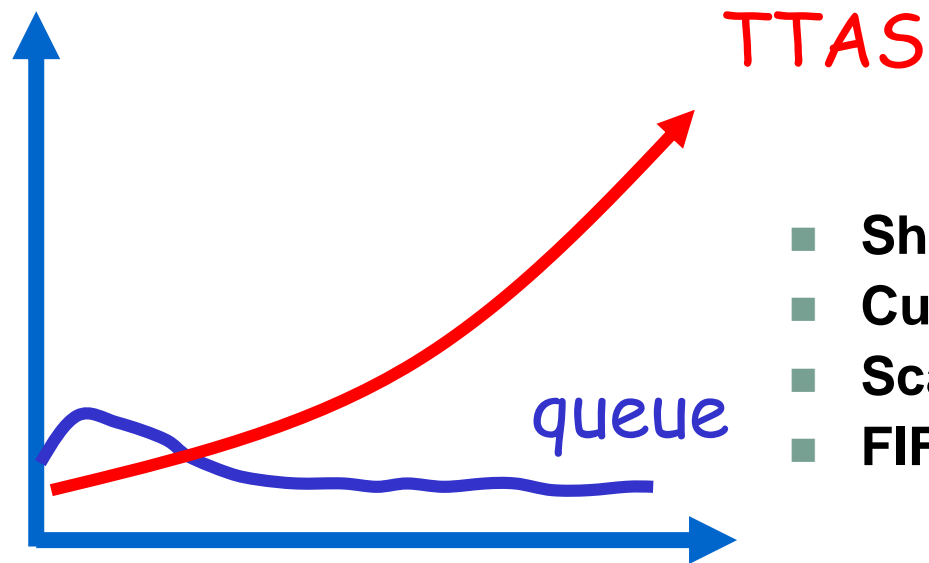




# Anderson Queue Lock

- Although the flags [] array is shared, contention on the array locations are minimised since each thread spins on its own locally cached copy of a single array location

# Performance



- Shorter handover than backoff
- Curve is practically flat
- Scalable performance
- FIFO fairness



# Anderson Queue Lock

- Good

- ☐ First truly scalable lock
- ☐ Simple, easy to implement

- Bad

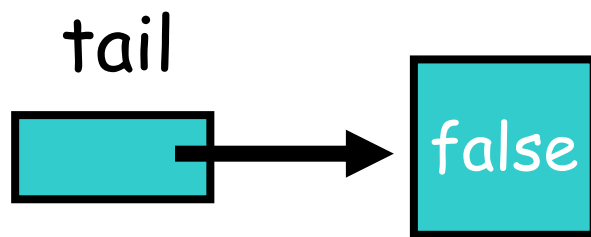
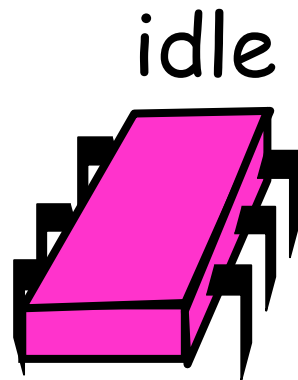
- ☐ Not space efficient
- ☐ One bit per thread
  - Unknown number of threads?
  - Small number of actual contenders?



# CLH Queue Lock



- Virtual Linked List keeps track of the queue
- Each thread's status is saved in its node:
  - True – has acquired the lock or wants to acquire the lock
  - False – is finished with the lock and has released it
- Each node keeps track of its predecessors status

# Initially



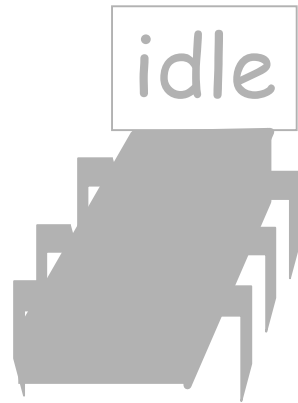


# Initially

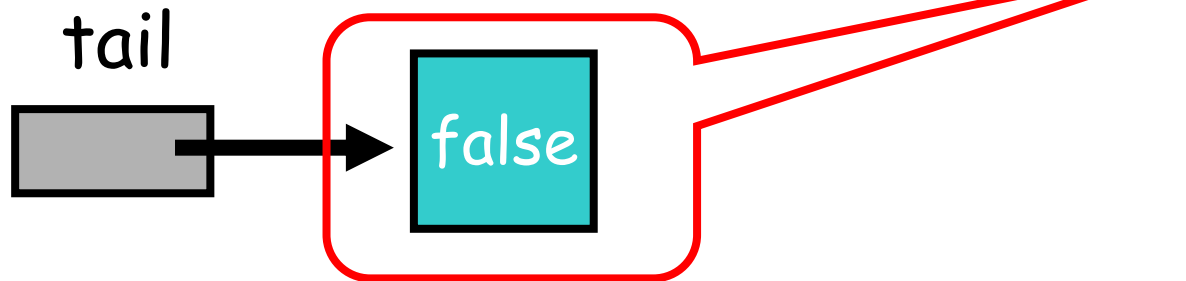


Queue tail

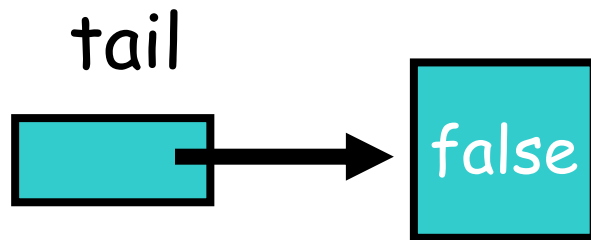
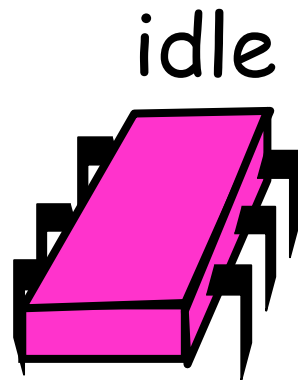
# Initially



Locked field: Lock is free




# Initially



# Purple Wants the Lock

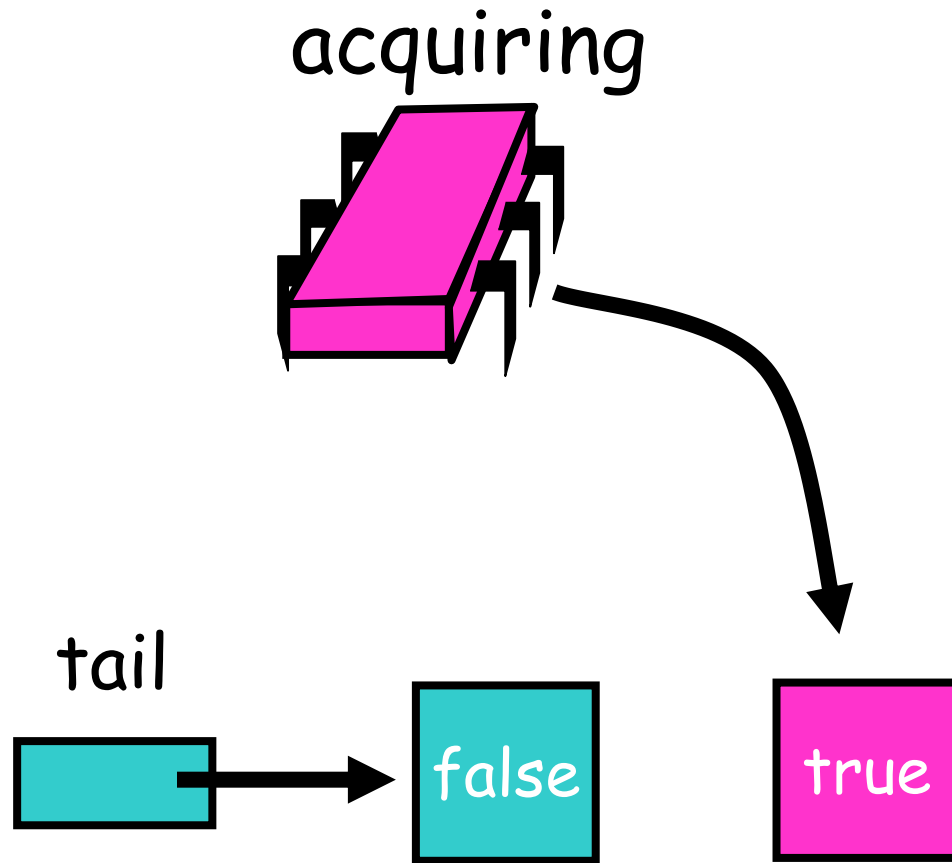
acquiring



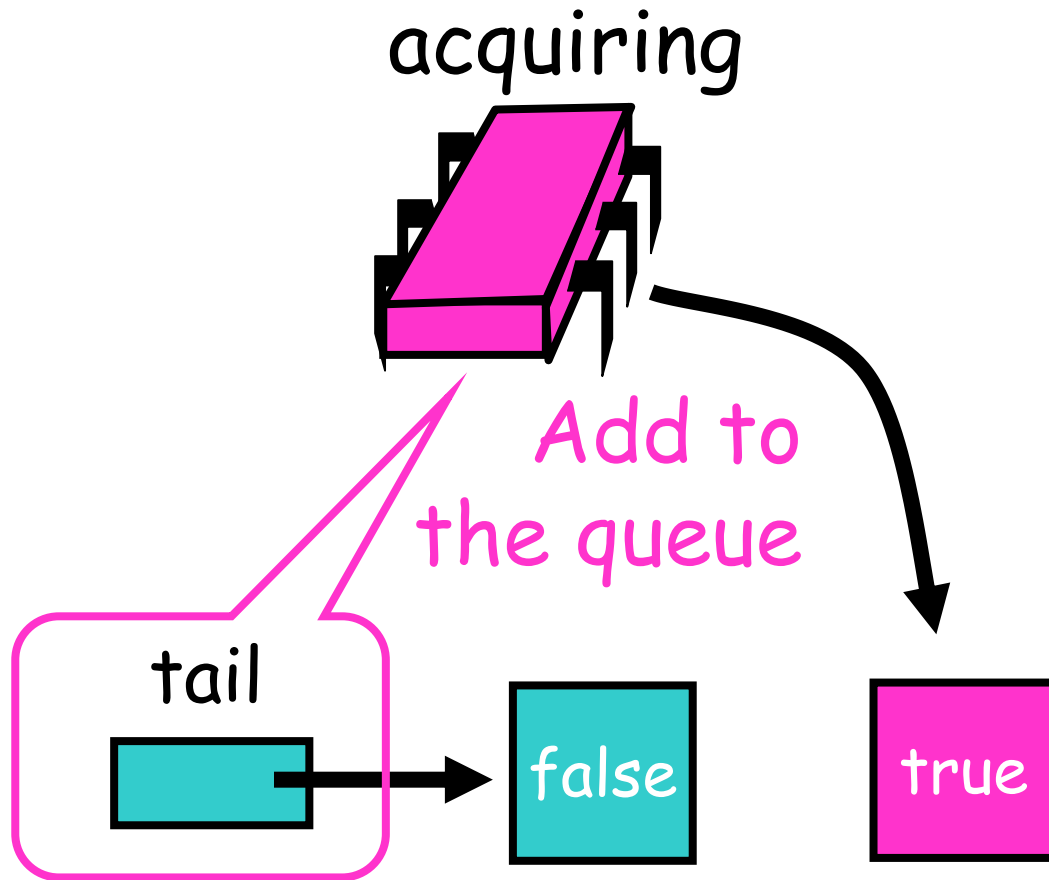
tail



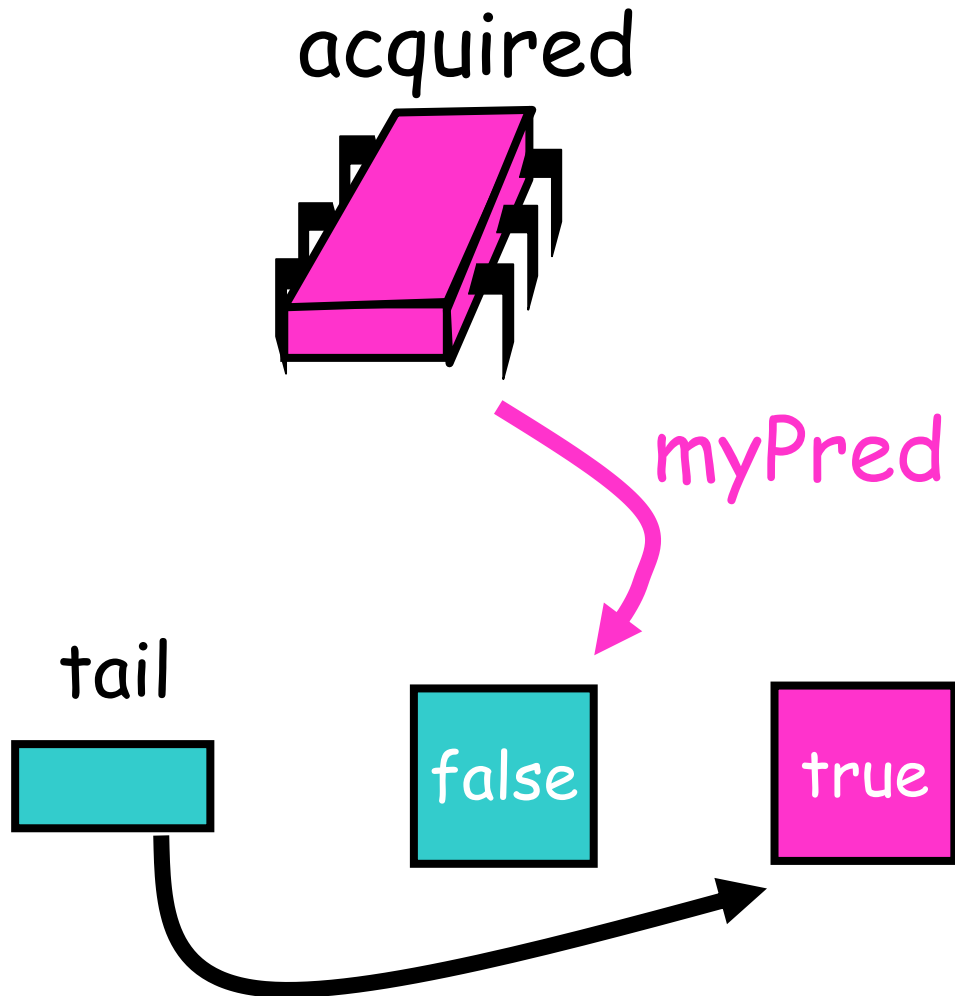
# Purple Wants the Lock




# Purple Wants the Lock



# Purple Has the Lock

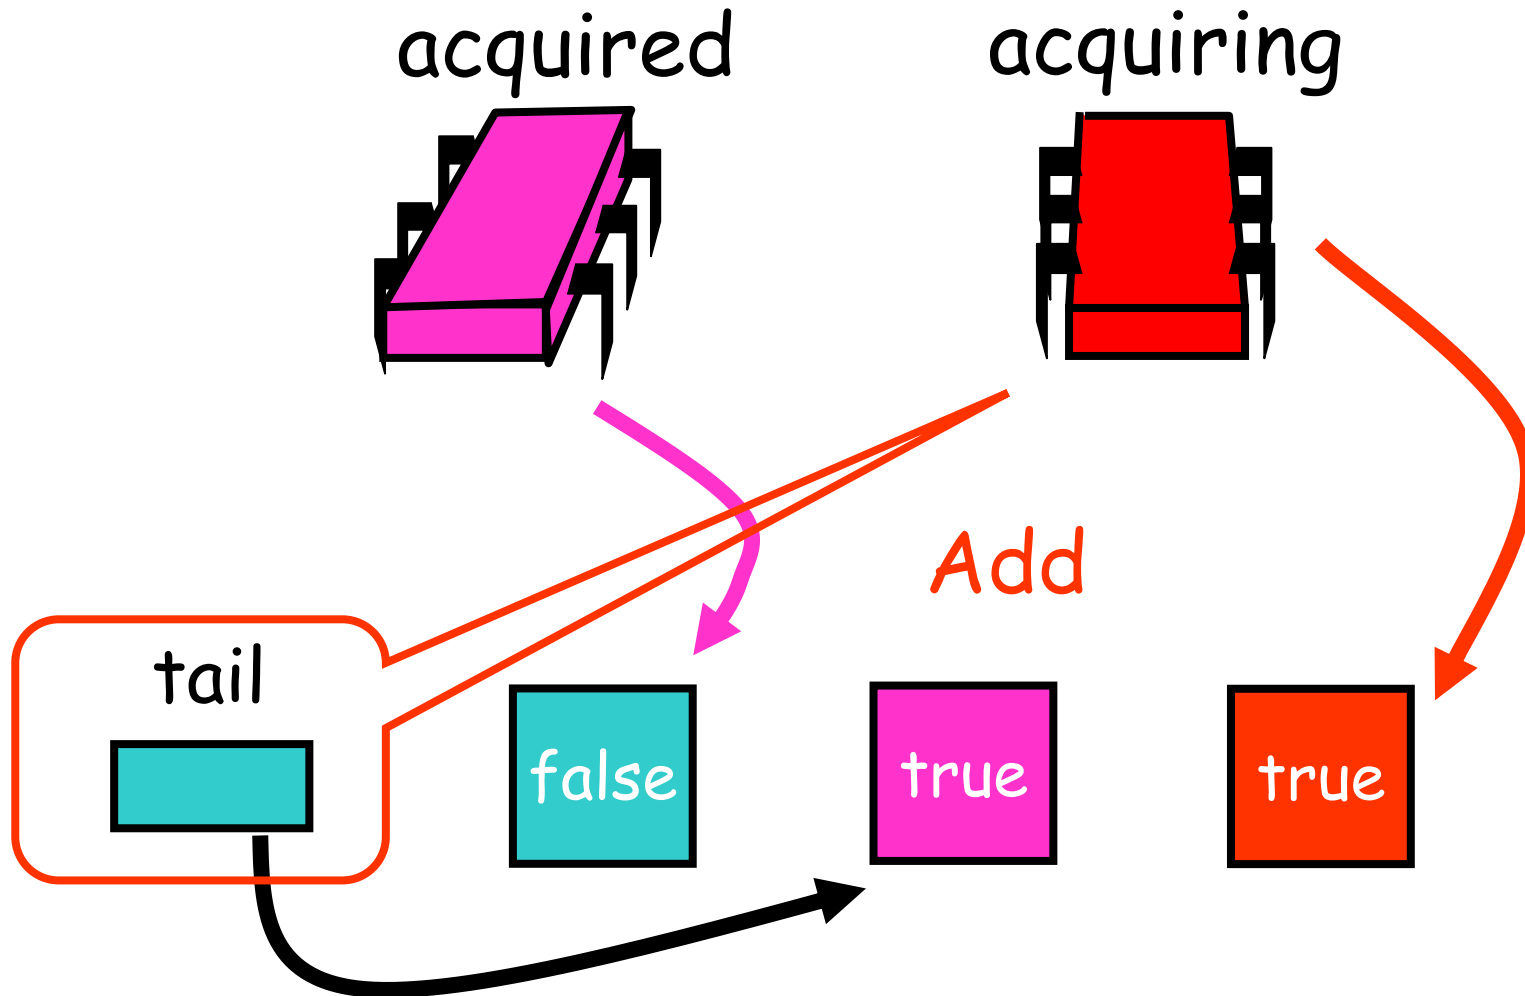


# Red Wants the Lock

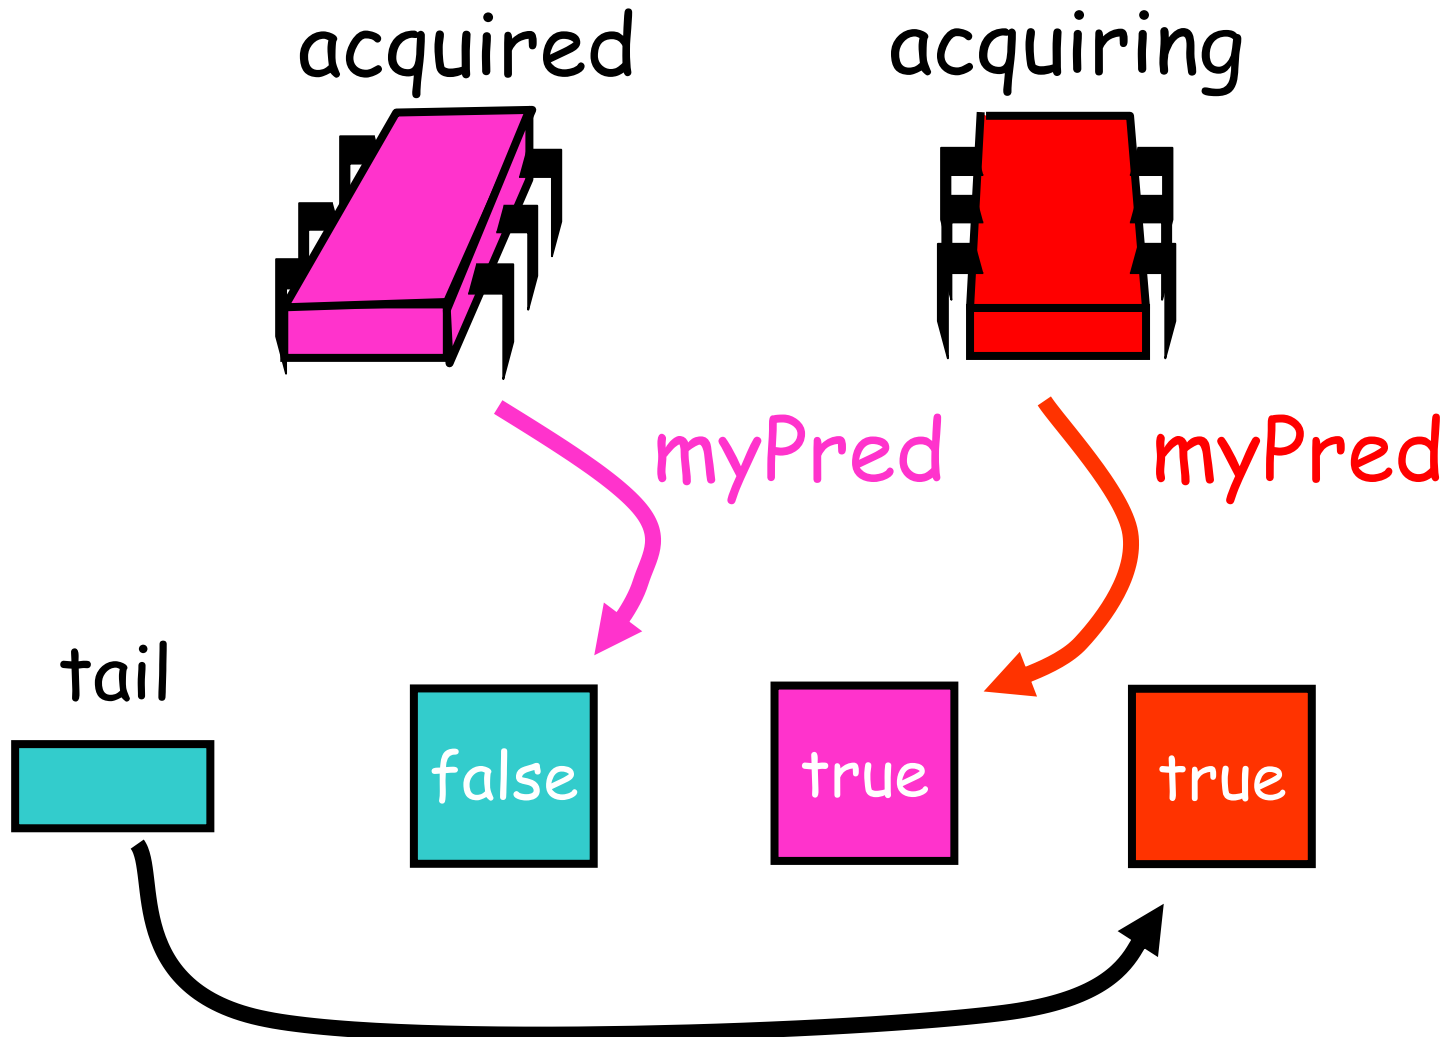





# Red Wants the Lock



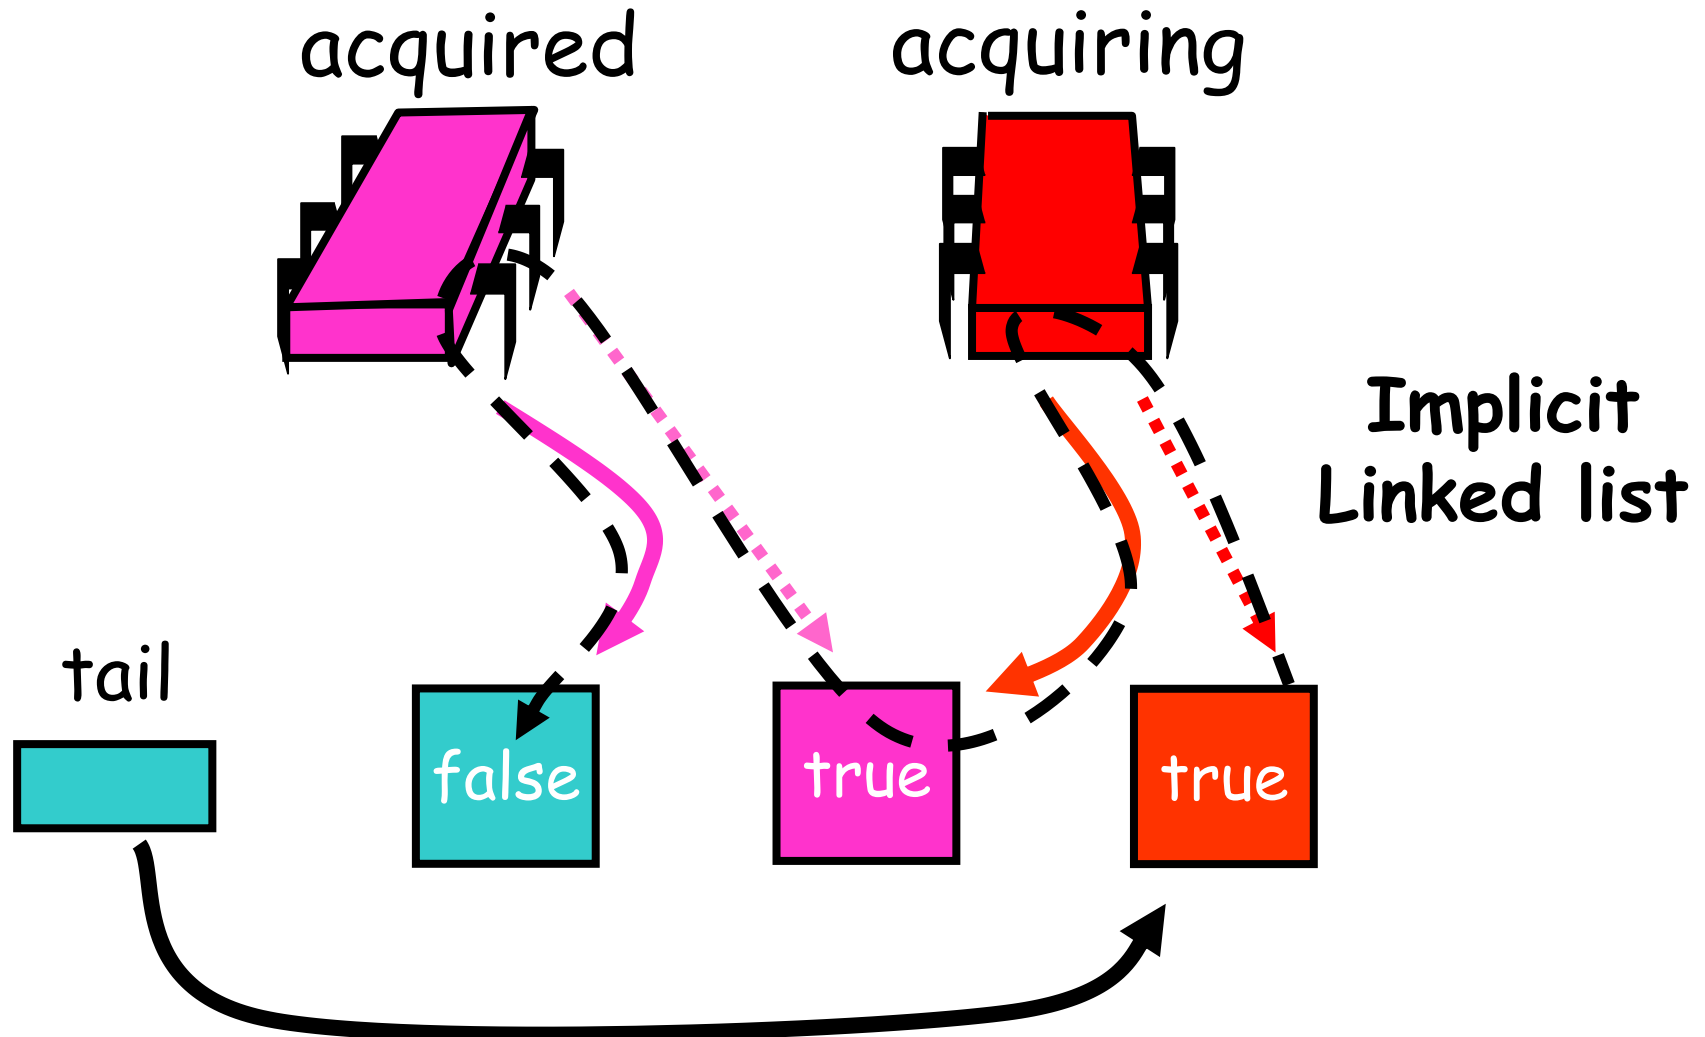
# Red Wants the Lock




# Red Wants the Lock



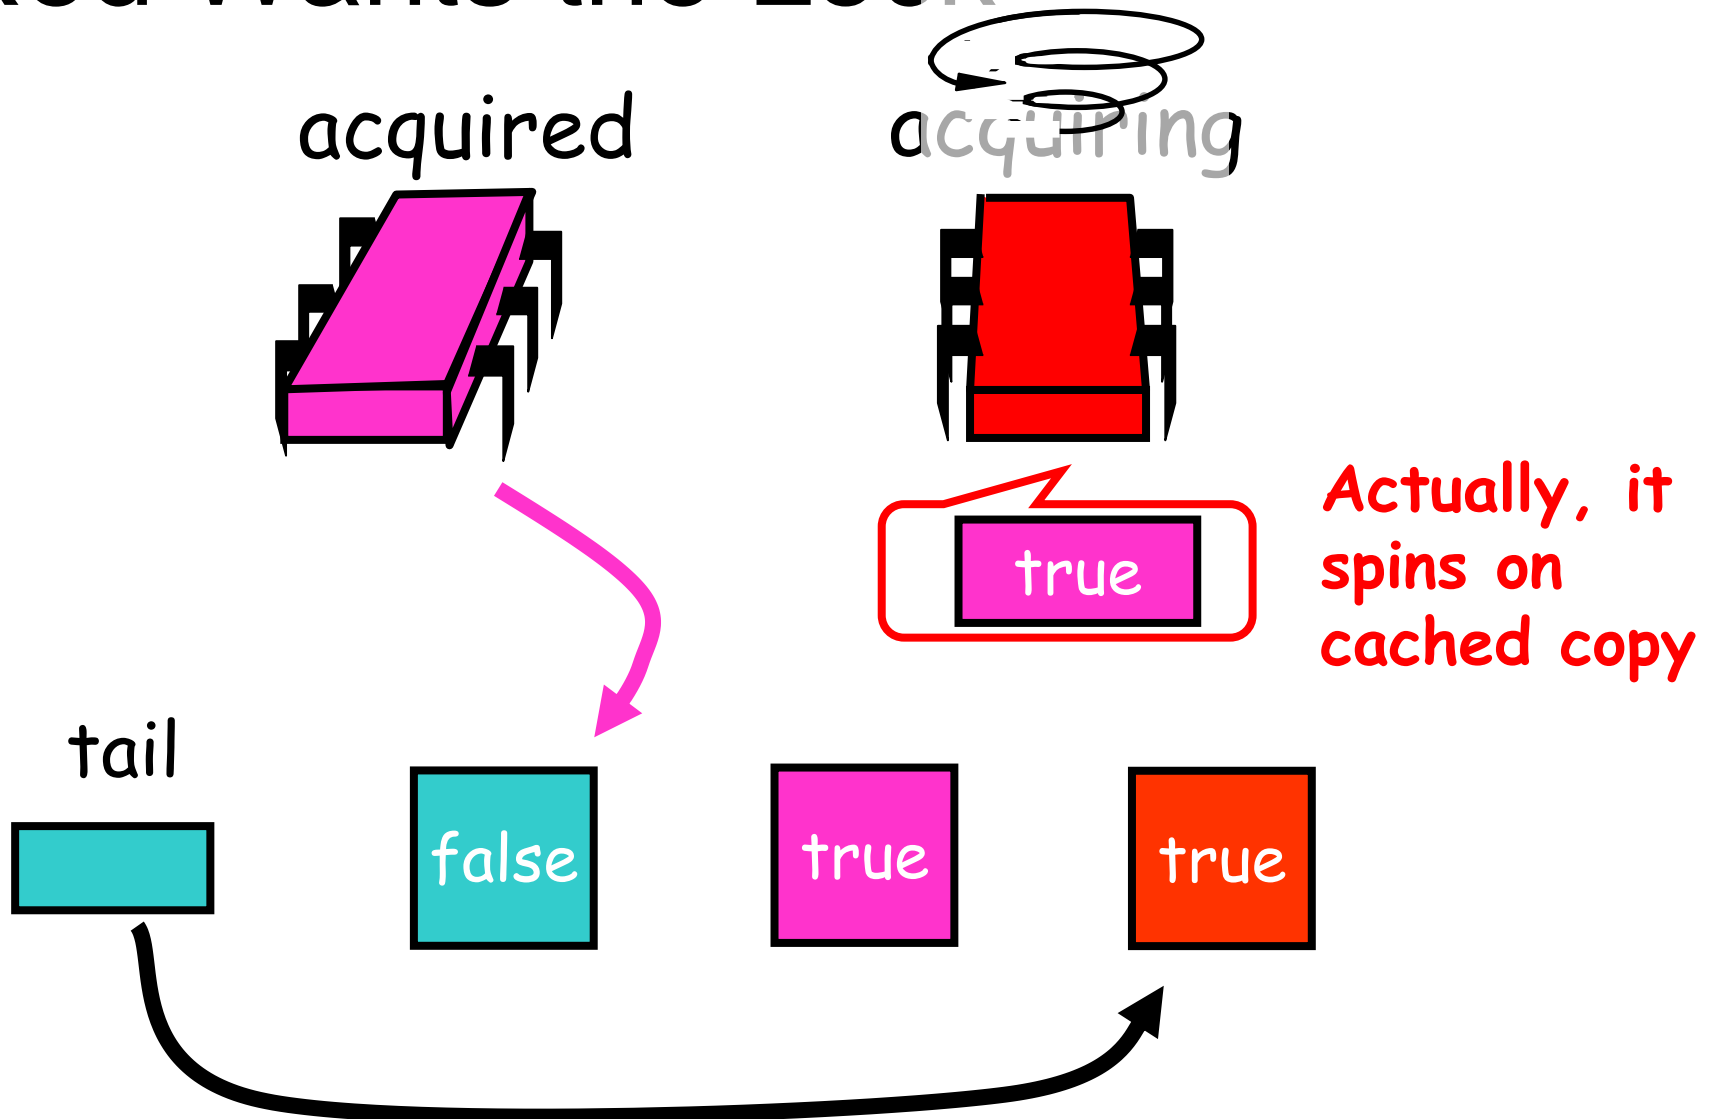
# Red Wants the Lock



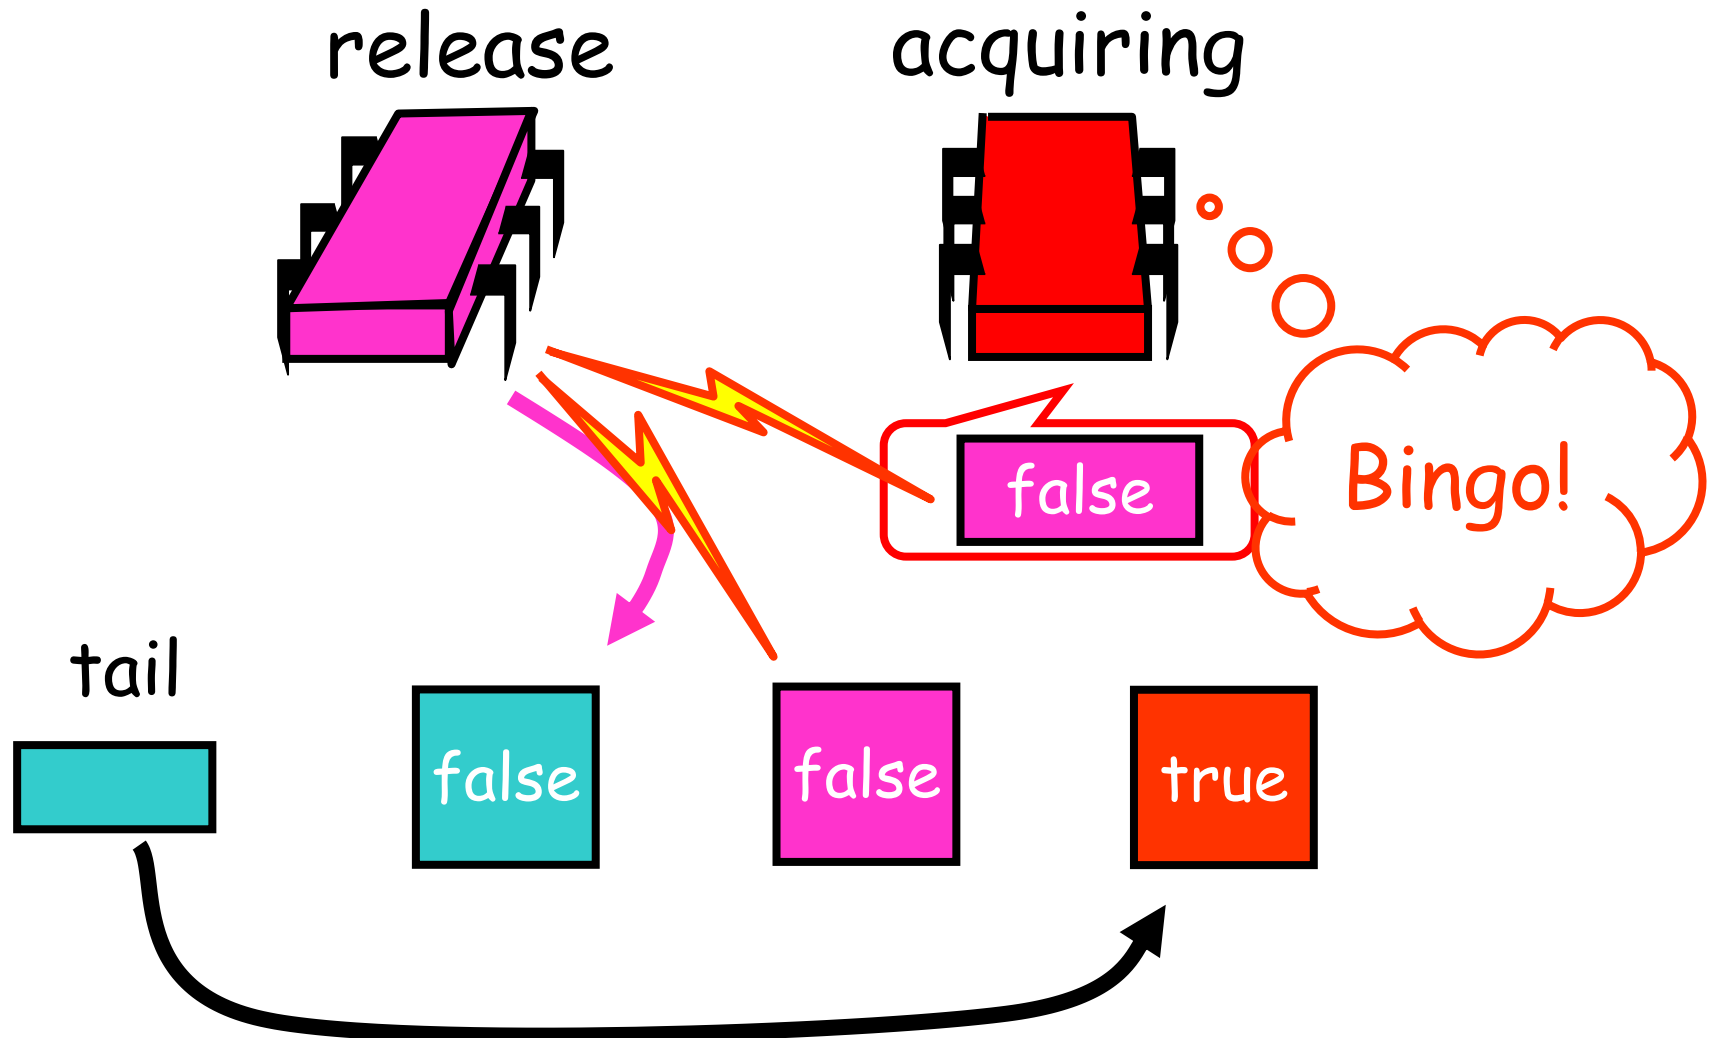
# Red Wants the Lock



# Red Wants the Lock

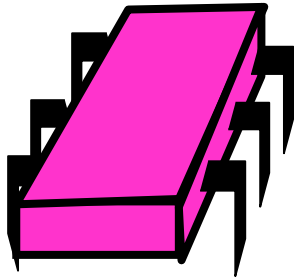


# Purple Releases

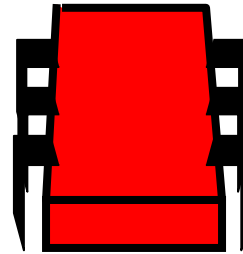


# Purple Releases

released



acquired



tail



true





# CLH Queue Lock

```
class CLHLock implements Lock {  
    AtomicReference<Qnode> tail;  
    ThreadLocal<Qnode> myNode  
        = new Qnode();  
    public void lock() {  
        myNode.locked = true;  
        Qnode pred  
            = tail.getAndSet(myNode);  
        while (pred.locked) {}  
    }  
}
```

# CLH Queue Lock

```
class CLHLock implements Lock {  
    AtomicReference<Qnode> tail;  
    ThreadLocal<Qnode> myNode  
        = new Qnode();  
    public void lock() {  
        myNode.locked = true;  
        Qnode pred  
            = tail.getAndSet(myNode);  
        while (pred.locked) {}  
    }  
}
```

Queue tail

# CLH Queue Lock

```
class CLHLock implements Lock {  
    AtomicReference<Qnode> tail;  
    ThreadLocal<Qnode> myNode  
        = new Qnode();  
    public void lock() {  
        myNode.locked = true;  
        Qnode pred  
            = tail.getAndSet(myNode);  
        while (pred.locked) {}  
    }  
}
```

Thread-local Qnode

# CLH Queue Lock

```
class CLHLock implements Lock {  
    AtomicReference<Qnode> tail;  
    ThreadLocal<Qnode> myNode  
        = new Qnode();  
    public void lock() {  
        myNode.locked = true;  
        Qnode pred  
            = tail.getAndSet(myNode);  
        while (pred.locked) {}  
    }  
}
```

Add my node

# CLH Queue Lock

```
class CLHLock implements Lock {  
    AtomicReference<Qnode> tail;  
    ThreadLocal<Qnode> myNode  
        = new Qnode();  
    public void lock() {  
        myNode.locked = true;  
        Qnode pred  
            = tail.getAndSet(myNode);  
        while (pred.locked) {}  
    }  
}
```

Spin until predecessor  
releases lock





# CLH Queue Lock

```
class CLHLock implements Lock {  
    ...  
    public void unlock() {  
        myNode.locked.set(false);  
        myNode = pred;  
    }  
}
```

# CLH Queue Lock

```
class CLHLock implements Lock {  
    ...  
    public void unlock() {  
        myNode.locked.set(false);  
        myNode = pred;  
    }  
}
```

Notify successor

# CLH Queue Lock

```
Class CLHLock implements Lock {  
    ...  
    public void unlock() {  
        myNode.locked.set(false);  
        myNode = pred;  
    }  
}
```

Recycle  
predecessor's node





# CLH Lock

- Good

- ☐ Lock release affects predecessor only
- ☐ Small, constant-sized space



# MCS Lock

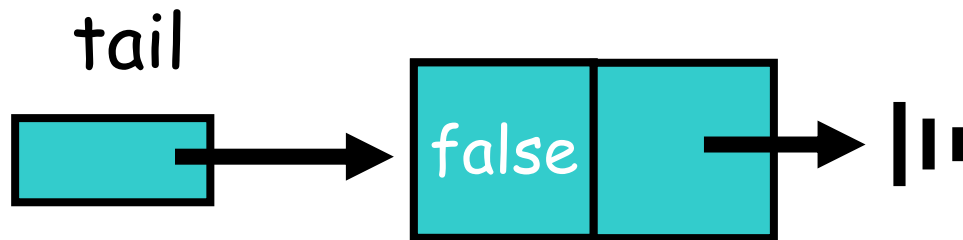

- FIFO order
- Spin on local memory only
- Small, Constant-size overhead



# MCS Queue Lock

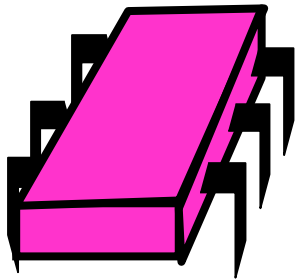
- Similar to CLHLock, but the linked list is explicit instead of implicit
- Each node in the Queue has a next field

# Initially

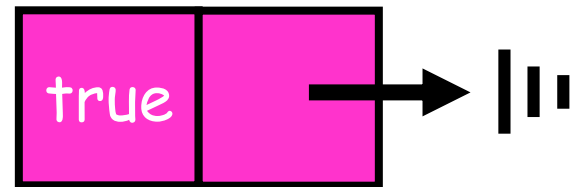


# Acquiring


acquiring



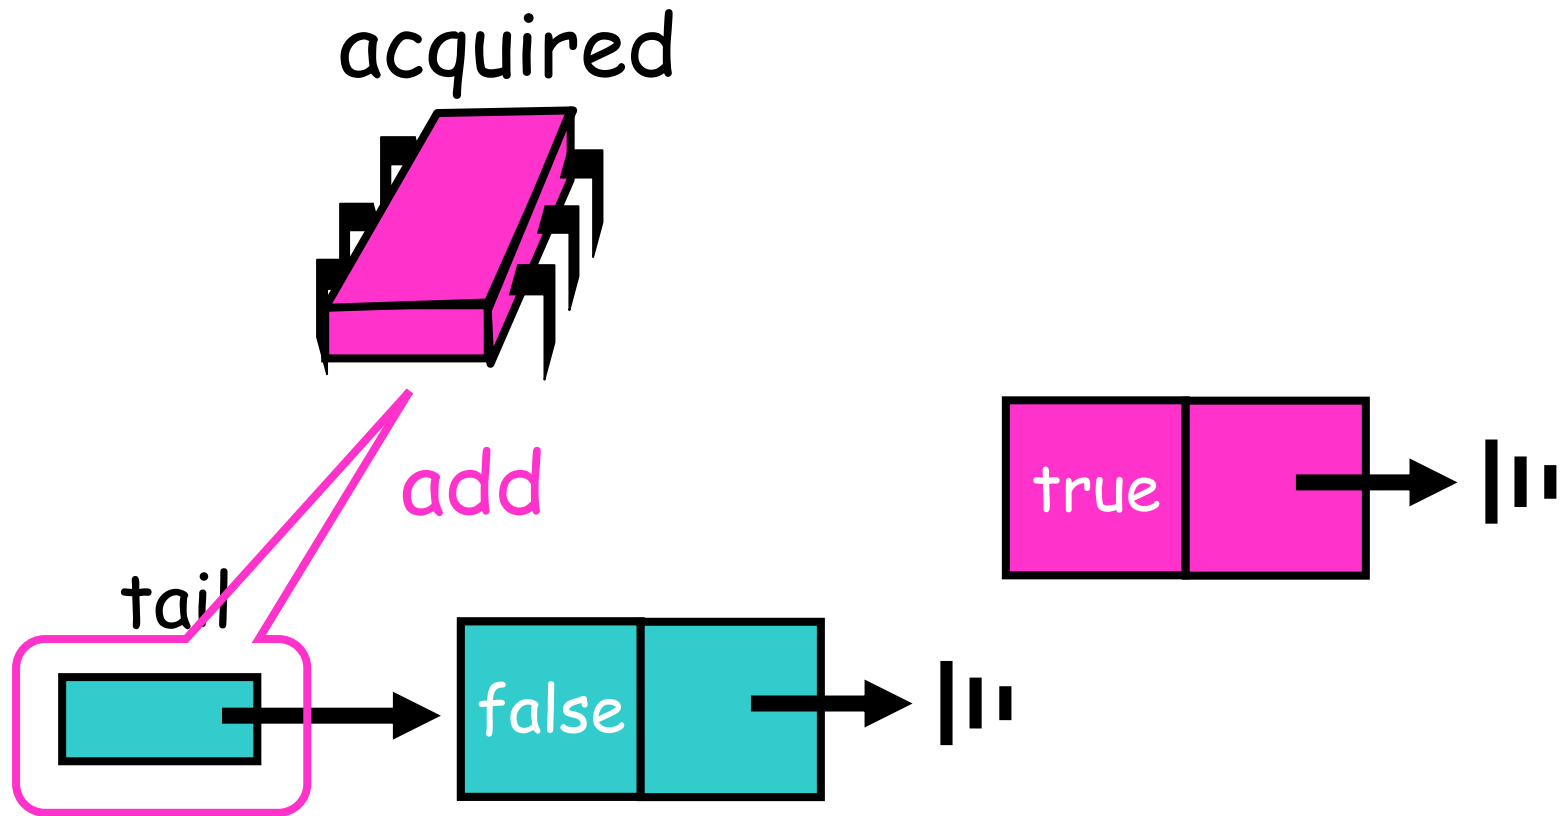
(add Qnode)



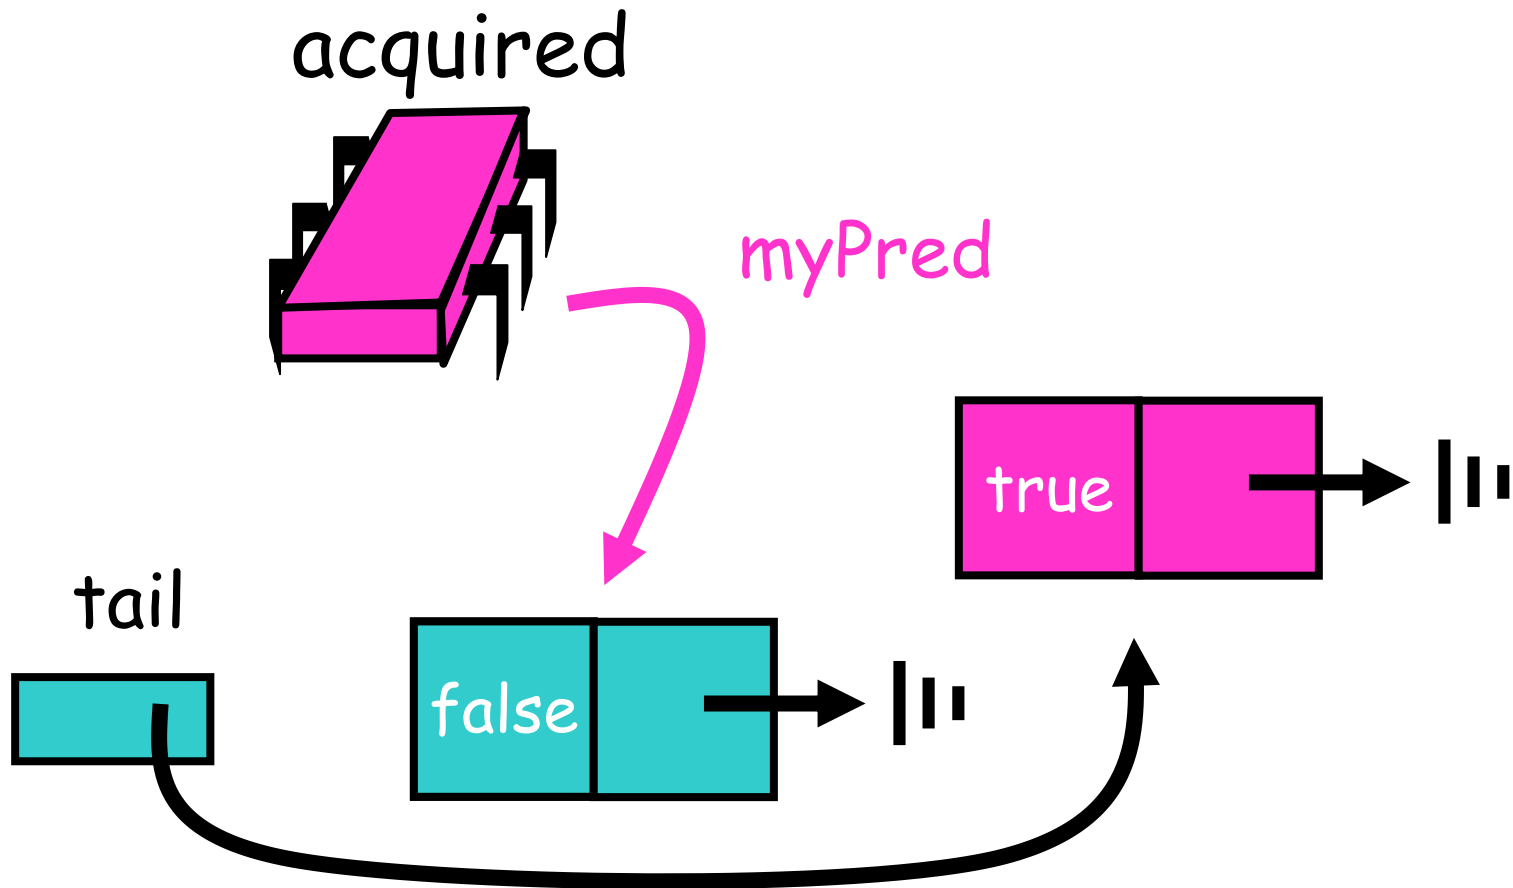
tail



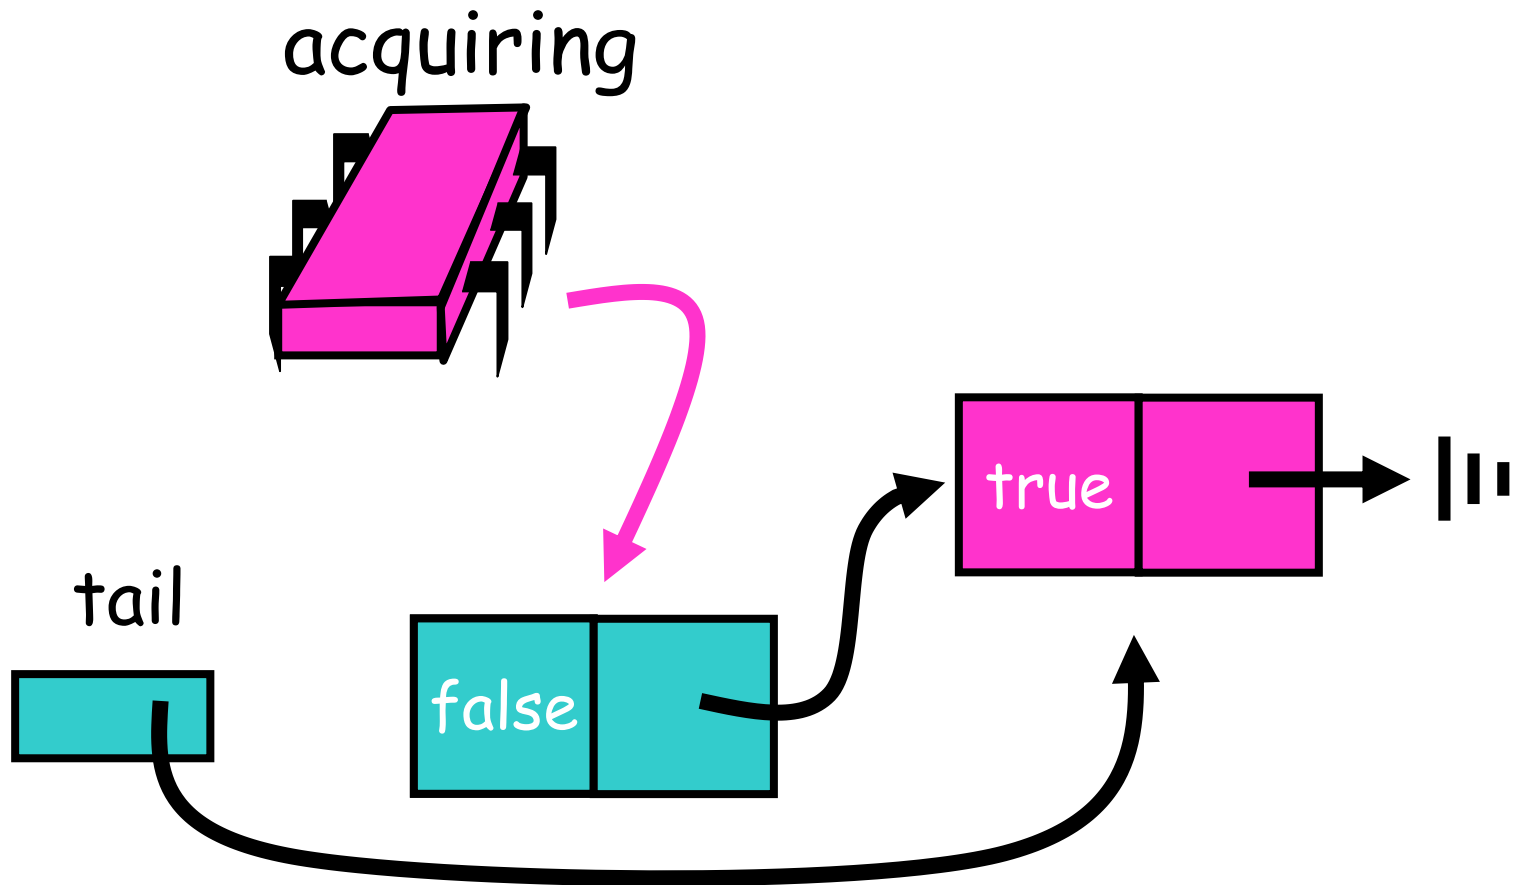
# Acquiring



# Acquiring




# Acquiring







# Acquiring



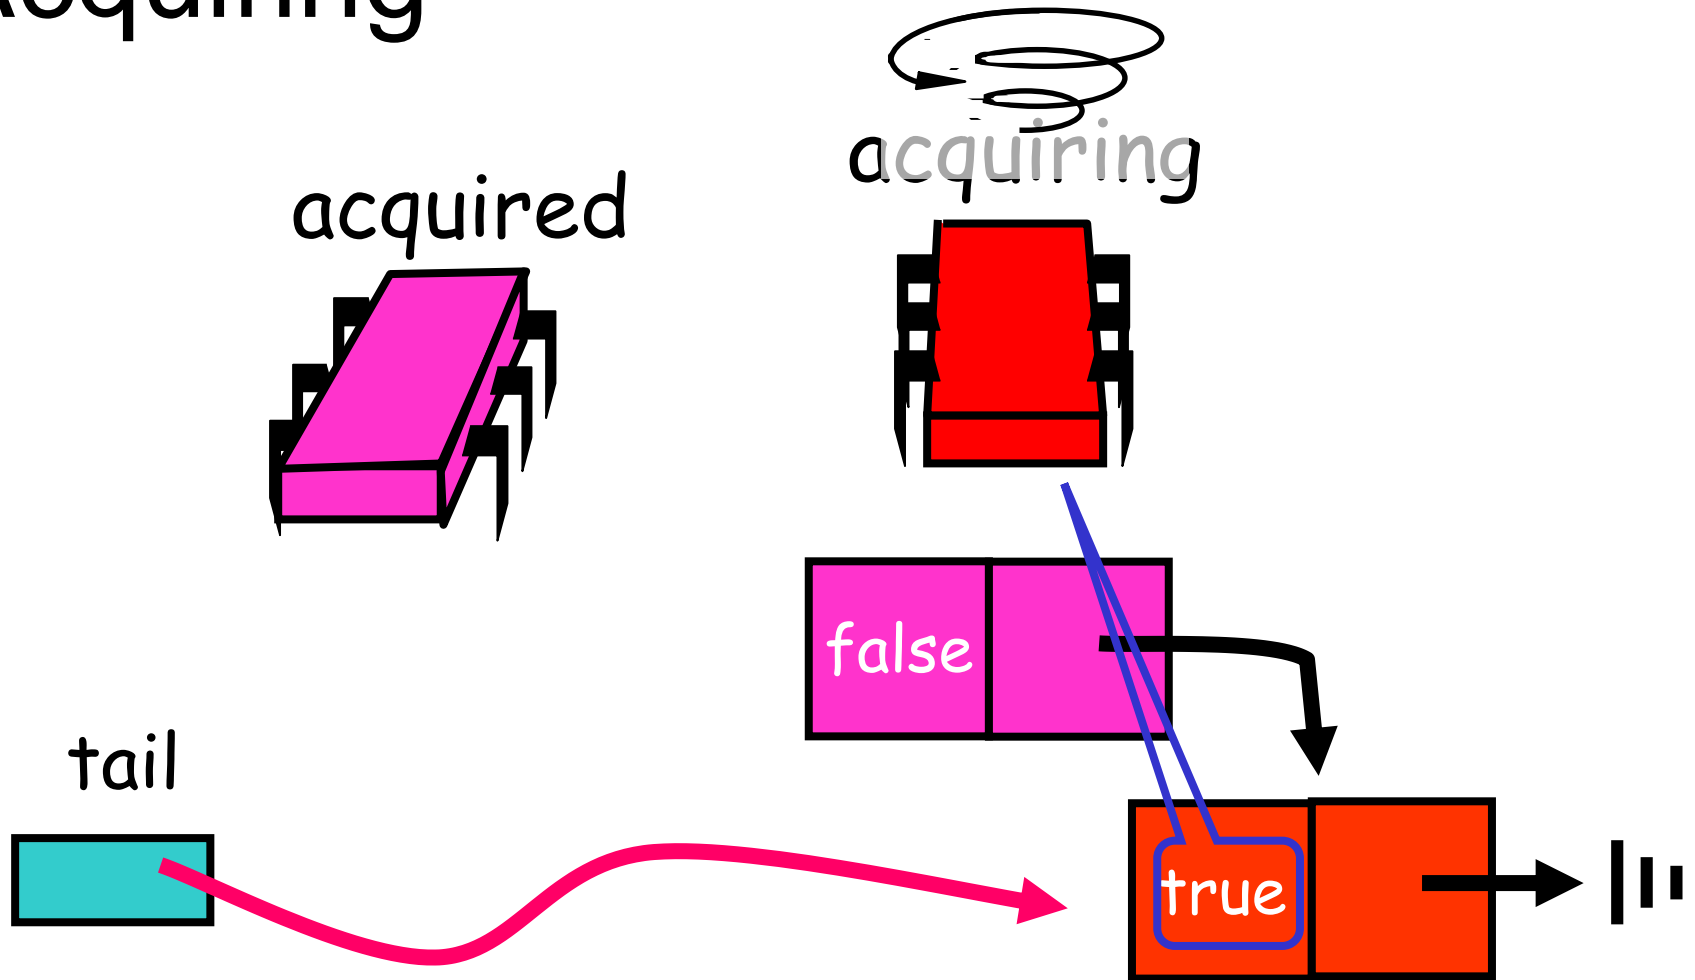
# Acquiring



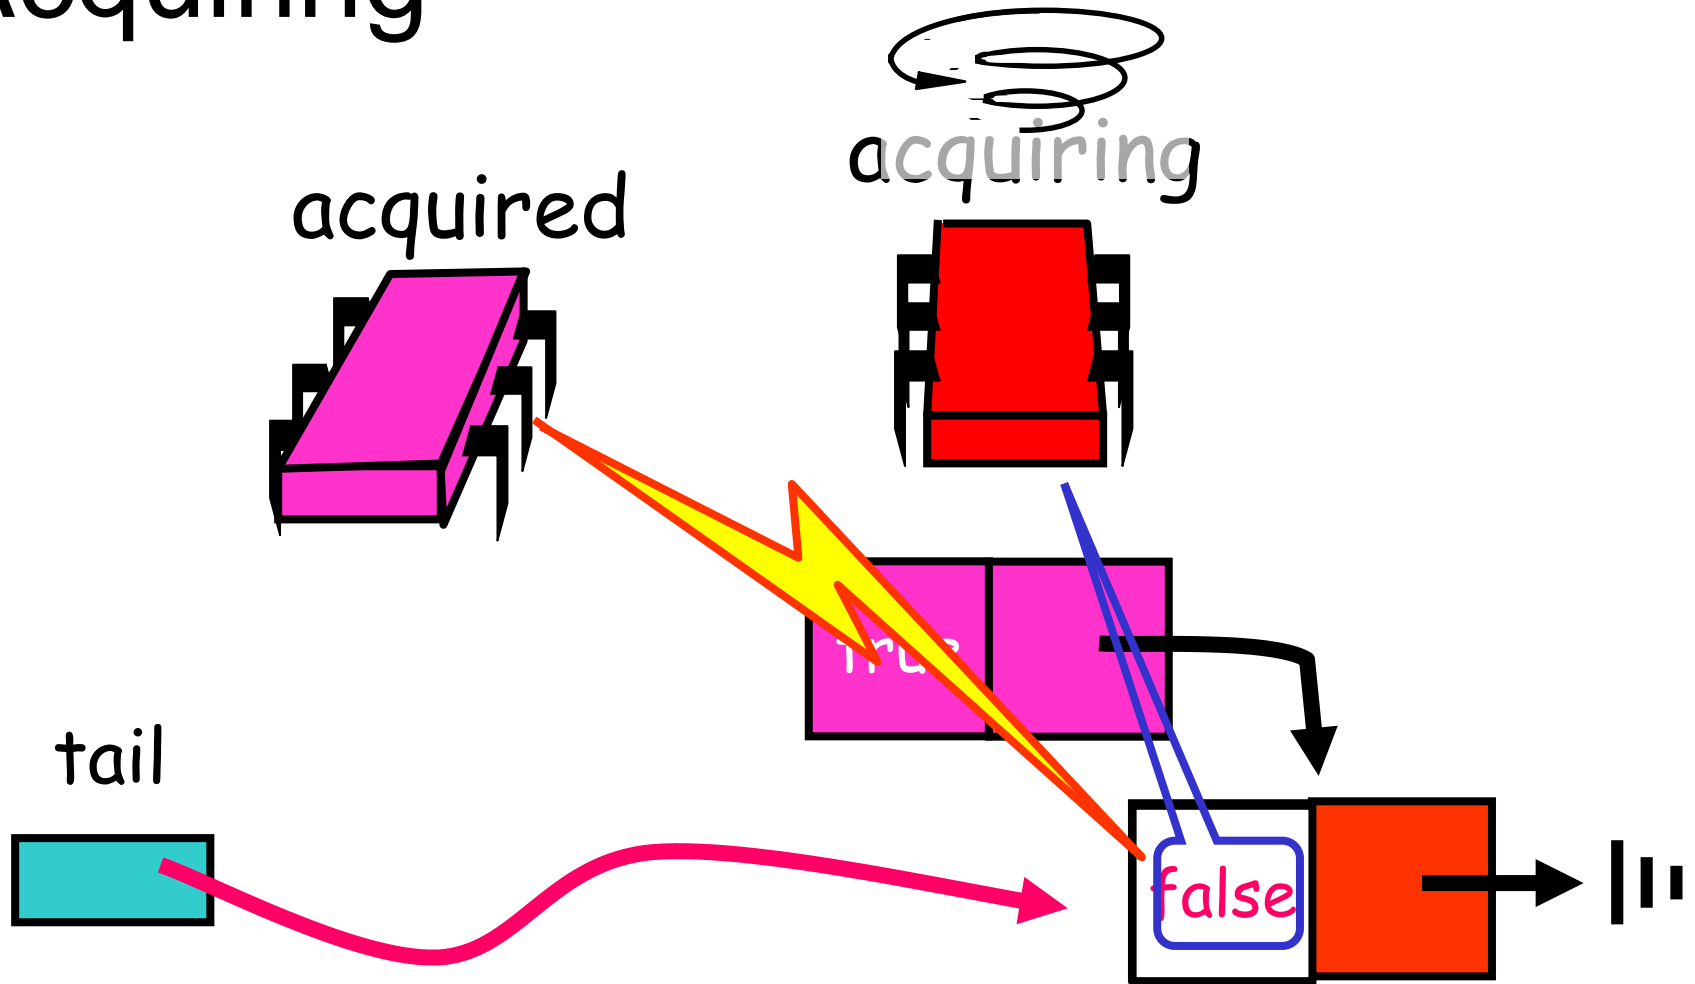
# Acquiring




# Acquiring



# Acquiring



# Acquiring





# MCS Queue Lock

```
class Qnode {  
    boolean locked = false;  
    qnode next = null;  
}
```

# MCS Queue Lock

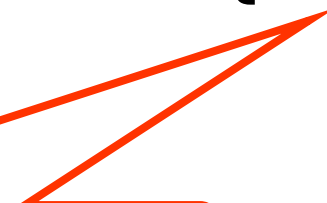
```
class MCSLock implements Lock {  
    AtomicReference tail;  
    public void lock() {  
        Qnode qnode = new Qnode();  
        Qnode pred = tail.getAndSet(qnode);  
        if (pred != null) {  
            qnode.locked = true;  
            pred.next = qnode;  
            while (qnode.locked) {}  
        }  
    }  
}
```



# MCS Queue Lock

```
class MCSLock implements Lock {  
    AtomicReference tail;  
    public void lock() {  
        Qnode qnode = new Qnode();  
        Qnode pred = tail.getAndSet(qnode);  
        if (pred != null) {  
            qnode.locked = true;  
            pred.next = qnode;  
            while (qnode.locked) {}  
        }  
    }  
}
```

Make a  
QNode



# MCS Queue Lock

```
class MCSLock implements Lock {  
    AtomicReference tail;  
    public void lock() {  
        Qnode qnode = new Qnode();  
        Qnode pred = tail.getAndSet(qnode);  
        if (pred != null) {  
            qnode.locked = true;  
            pred.next = qnode;  
            while (qnode.locked) {}  
        }  
    }  
}
```

add my Node to  
the tail of  
queue

# MCS Queue Lock

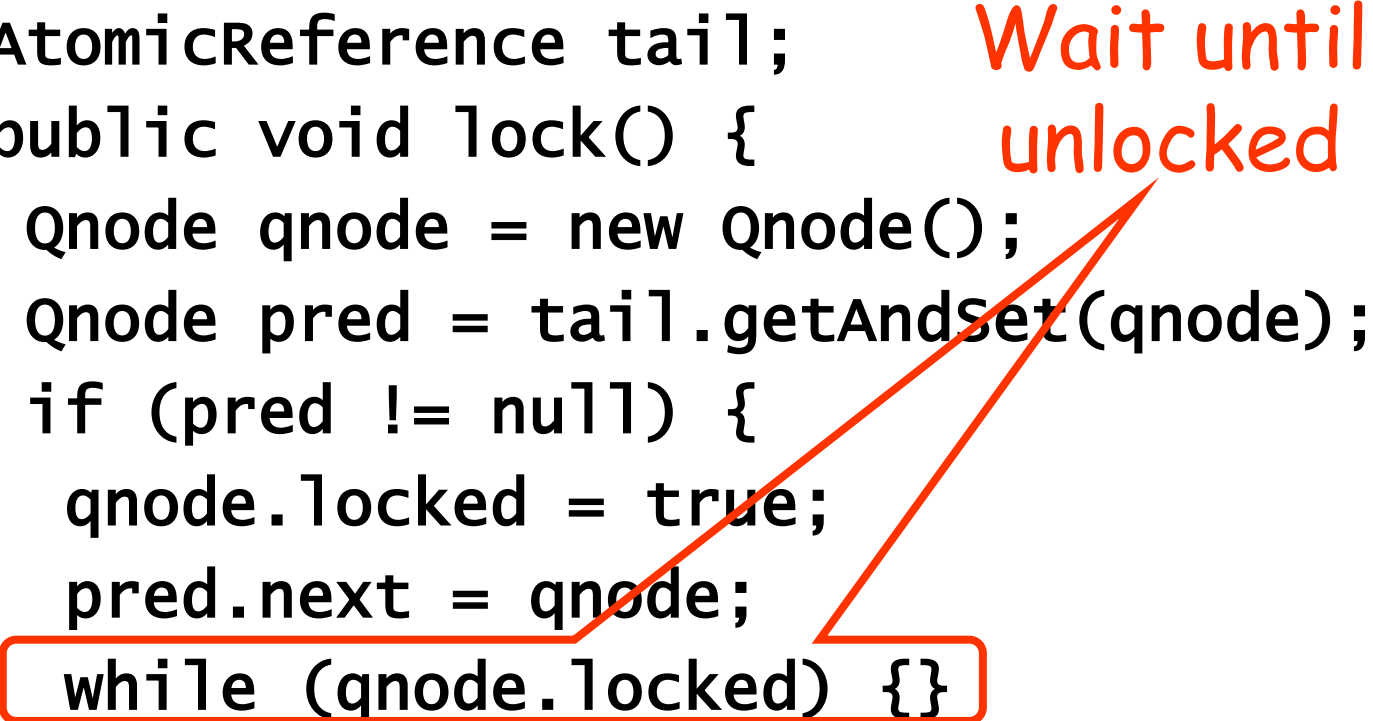
```
class MCSLock implements Lock {  
    AtomicReference tail;  
    public void lock() {  
        Qnode qnode = new Qnode();  
        Qnode pred = tail.getAndSet(qnode);  
        if (pred != null) {  
            qnode.locked = true;  
            pred.next = qnode;  
            while (qnode.locked) {}  
        }  
    }  
}
```

Fix if queue  
was non-empty

# MCS Queue Lock

```
class MCSLock implements Lock {  
    AtomicReference tail;  
    public void lock() {  
        Qnode qnode = new Qnode();  
        Qnode pred = tail.getAndSet(qnode);  
        if (pred != null) {  
            qnode.locked = true;  
            pred.next = qnode;  
            while (qnode.locked) {}  
        }  
    }  
}
```

Wait until  
unlocked



# MCS Queue Unlock

```
class MCSLock implements Lock {  
    AtomicReference tail;  
    public void unlock() {  
        if (qnode.next == null) {  
            if (tail.CAS(qnode, null)  
                return;  
            while (qnode.next == null) {}  
        }  
        qnode.next.locked = false;  
    }  
}
```

# MCS Queue Lock

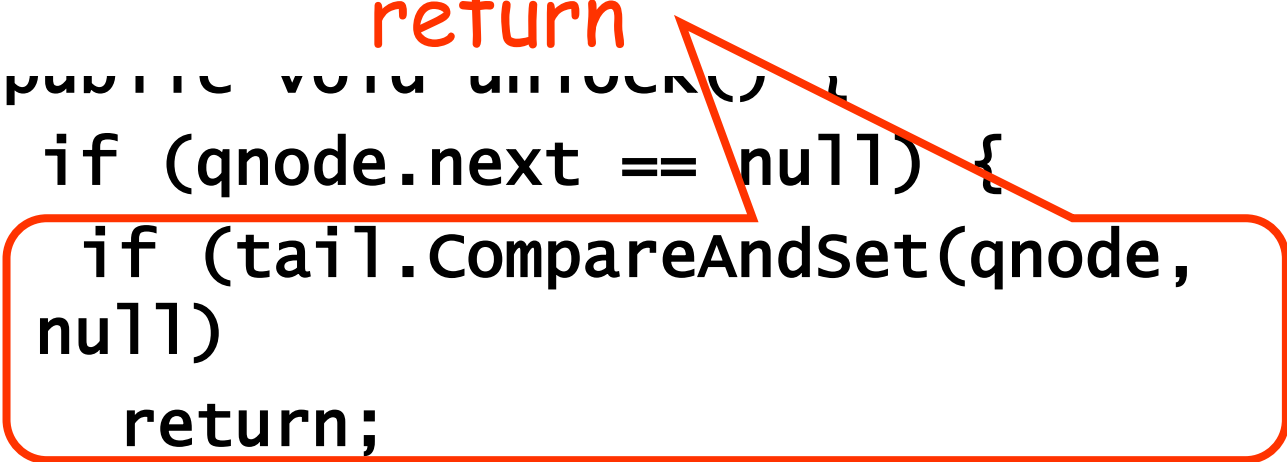
```
class MCSLock implements Lock {  
    AtomicReference tail;  
    public void unlock() {  
        if (qnode.next == null) {  
            if (tail.CAS(qnode, null)  
                return;  
            while (qnode.next == null) {}  
        }  
        qnode.next.locked = false;  
    }  
}
```

Missing  
successor?

# MCS Queue Lock

```
public void unlock() {  
    if (qnode.next == null) {  
        if (tail.CompareAndSet(qnode,  
                                null)  
            return;  
        while (qnode.next == null) {}  
    }  
    qnode.next.locked = false;  
}
```

*If really no successor,  
return*



# MCS Queue Lock

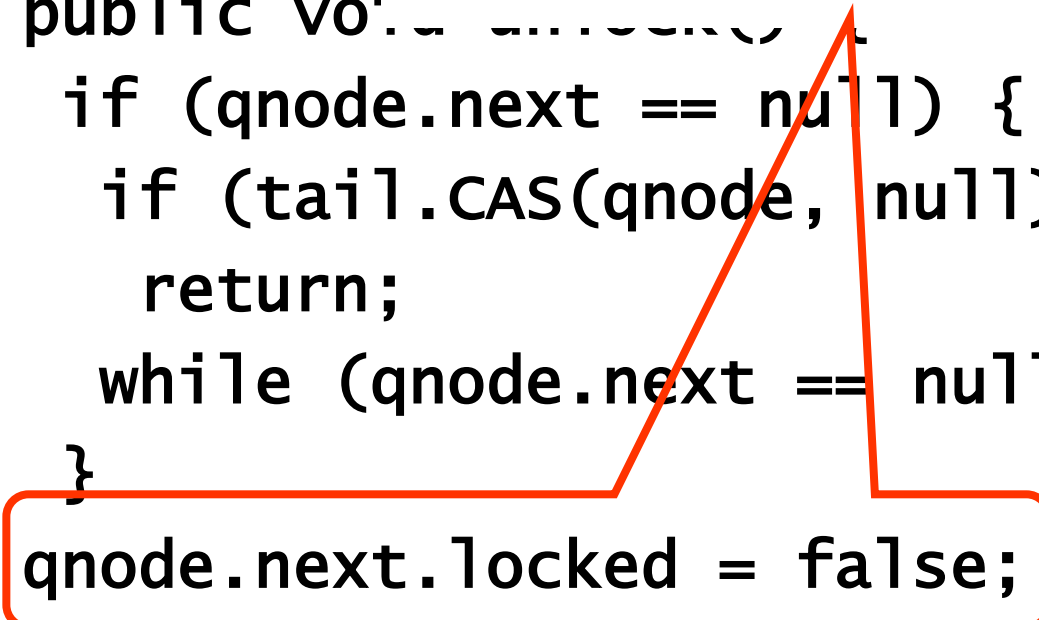
```
        }  
        Otherwise wait for  
        successor to catch up  
        public void unlock() {  
            if (qnode.next == null) {  
                if (tail.CAS(qnode, null)  
                    return;  
                while (qnode.next == null) {}  
            }  
            qnode.next.locked = false;  
        }  
    }  
}
```




# MCS Queue Lock

```
class MCSLock implements Lock {  
    AtomicRefNode tail;  
    public void lock() {  
        if (qnode.next == null) {  
            if (tail.CAS(qnode, null)  
                return;  
            while (qnode.next == null) {}  
        }  
        qnode.next.locked = false;  
    }  
}
```

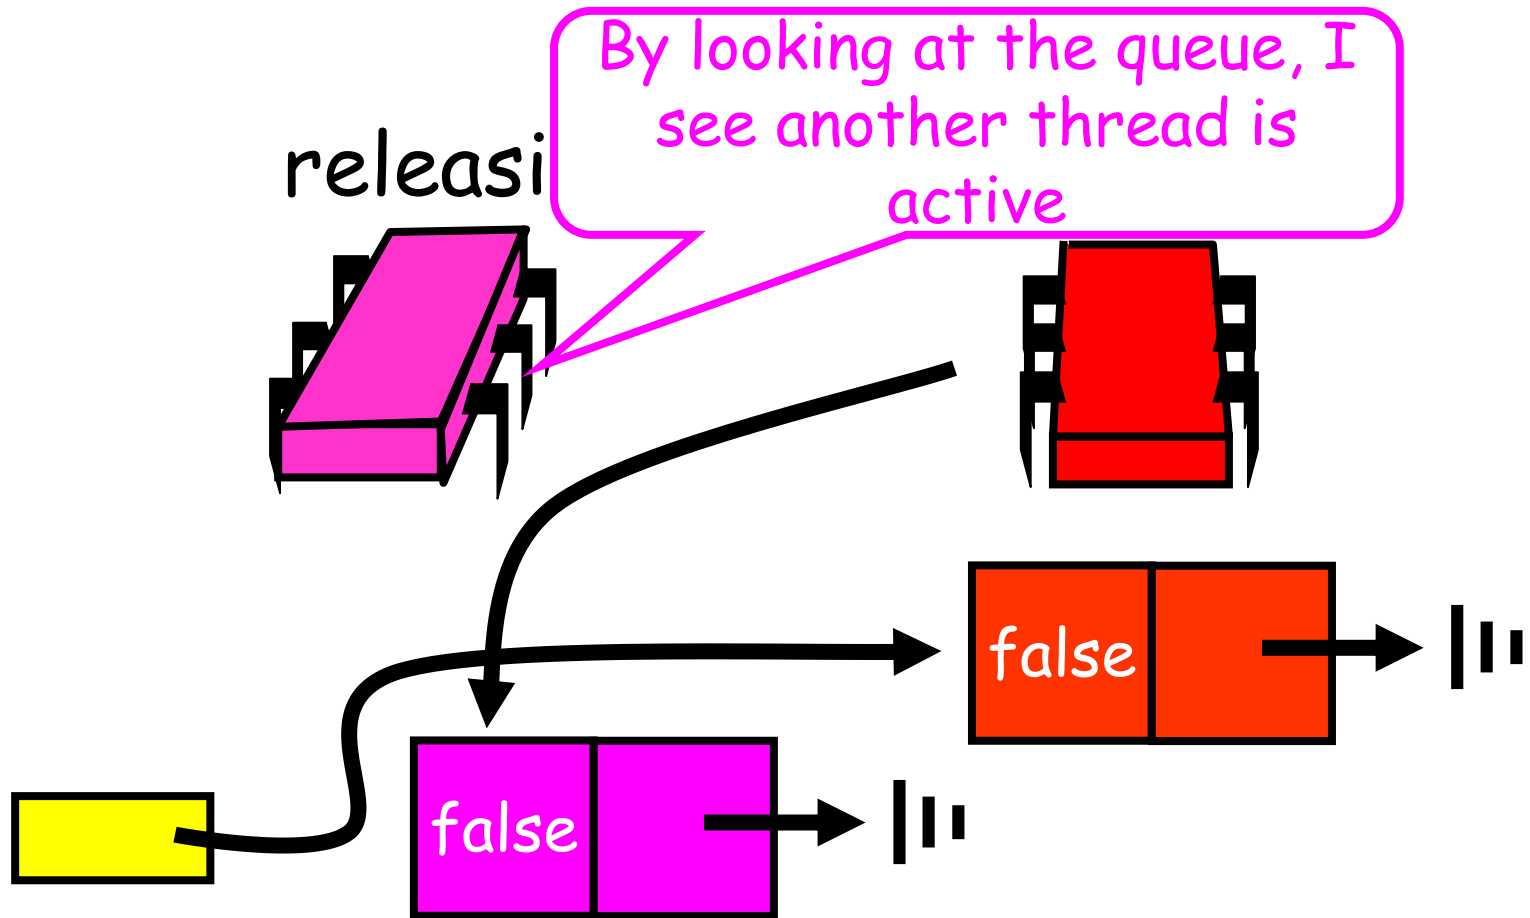
Pass lock to successor



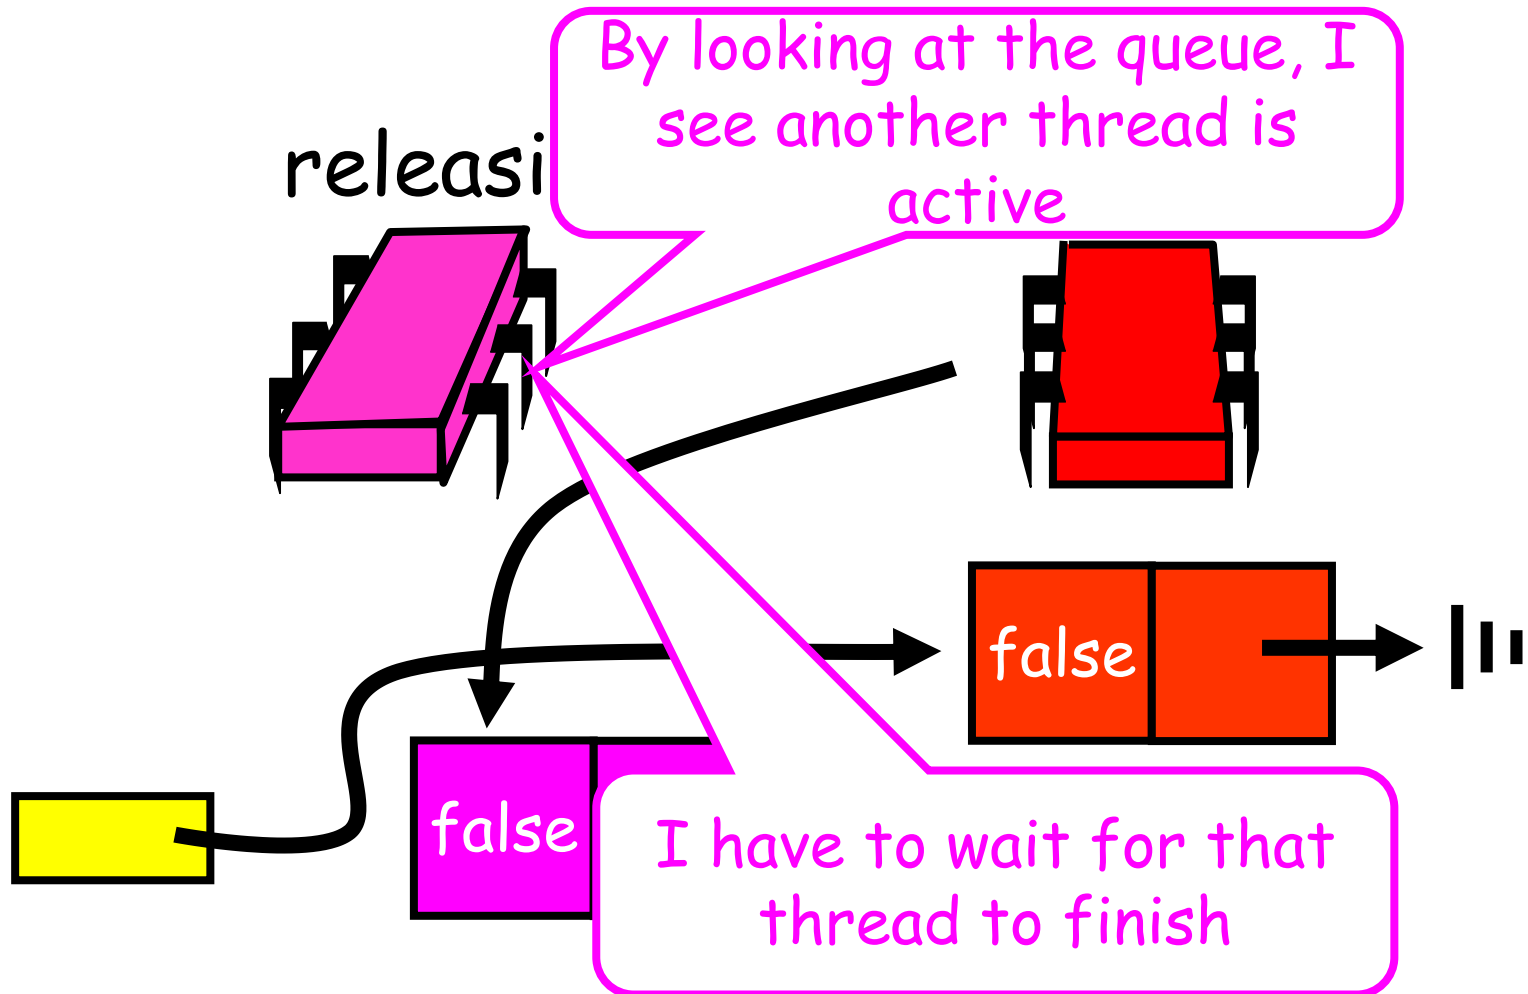
# Purple Release



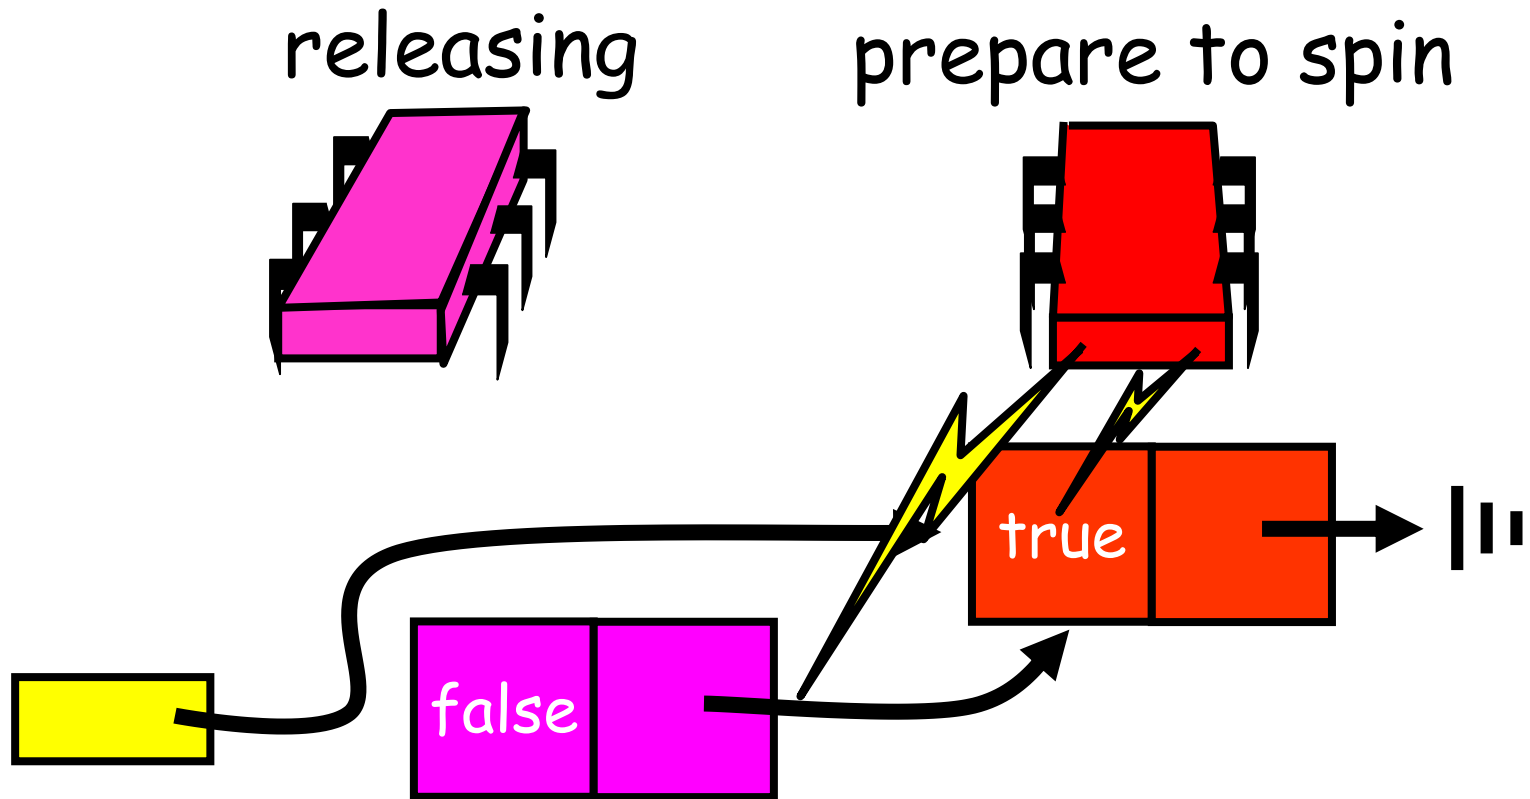
# Purple Release



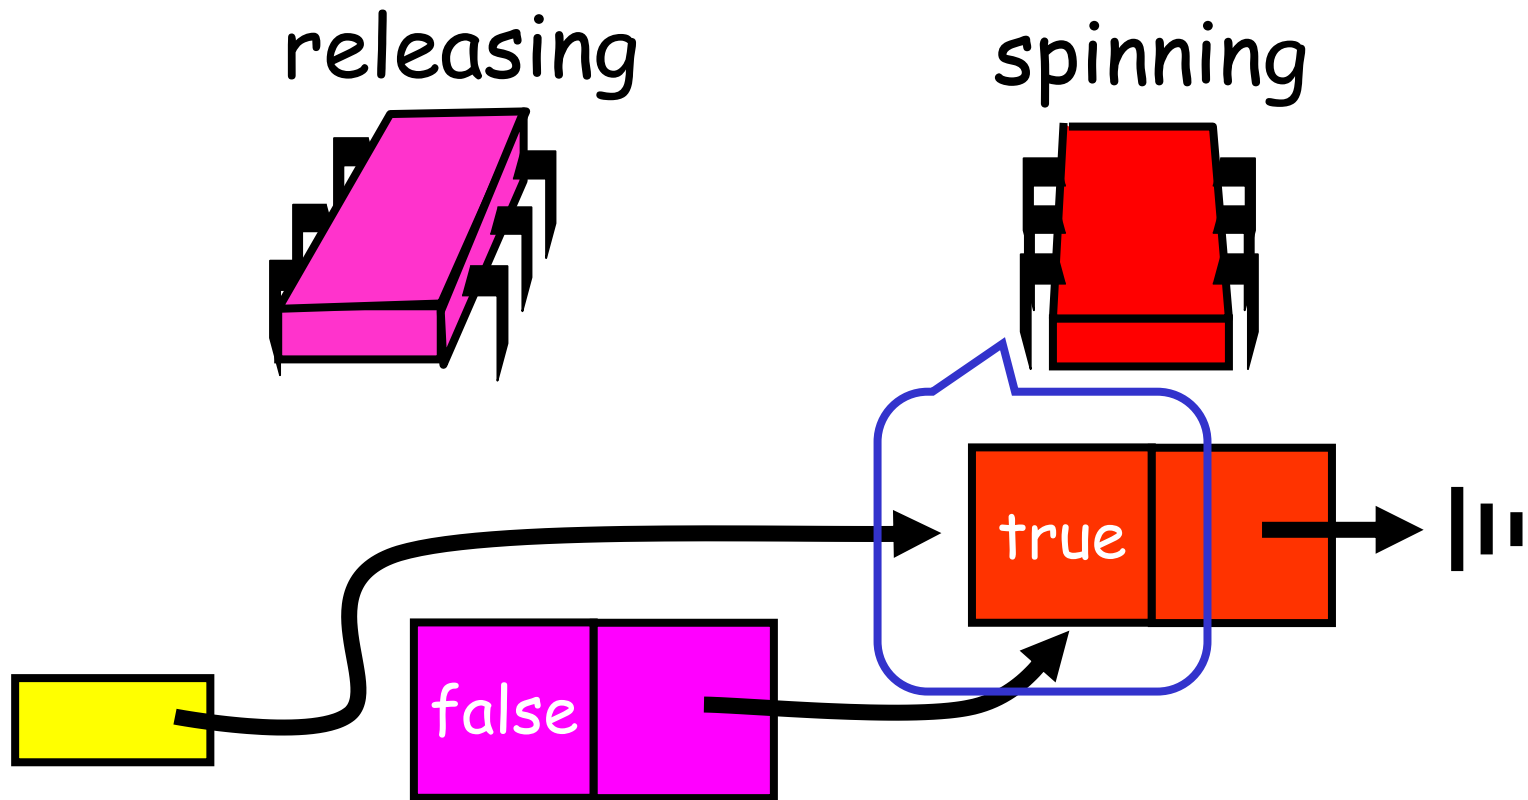
# Purple Release



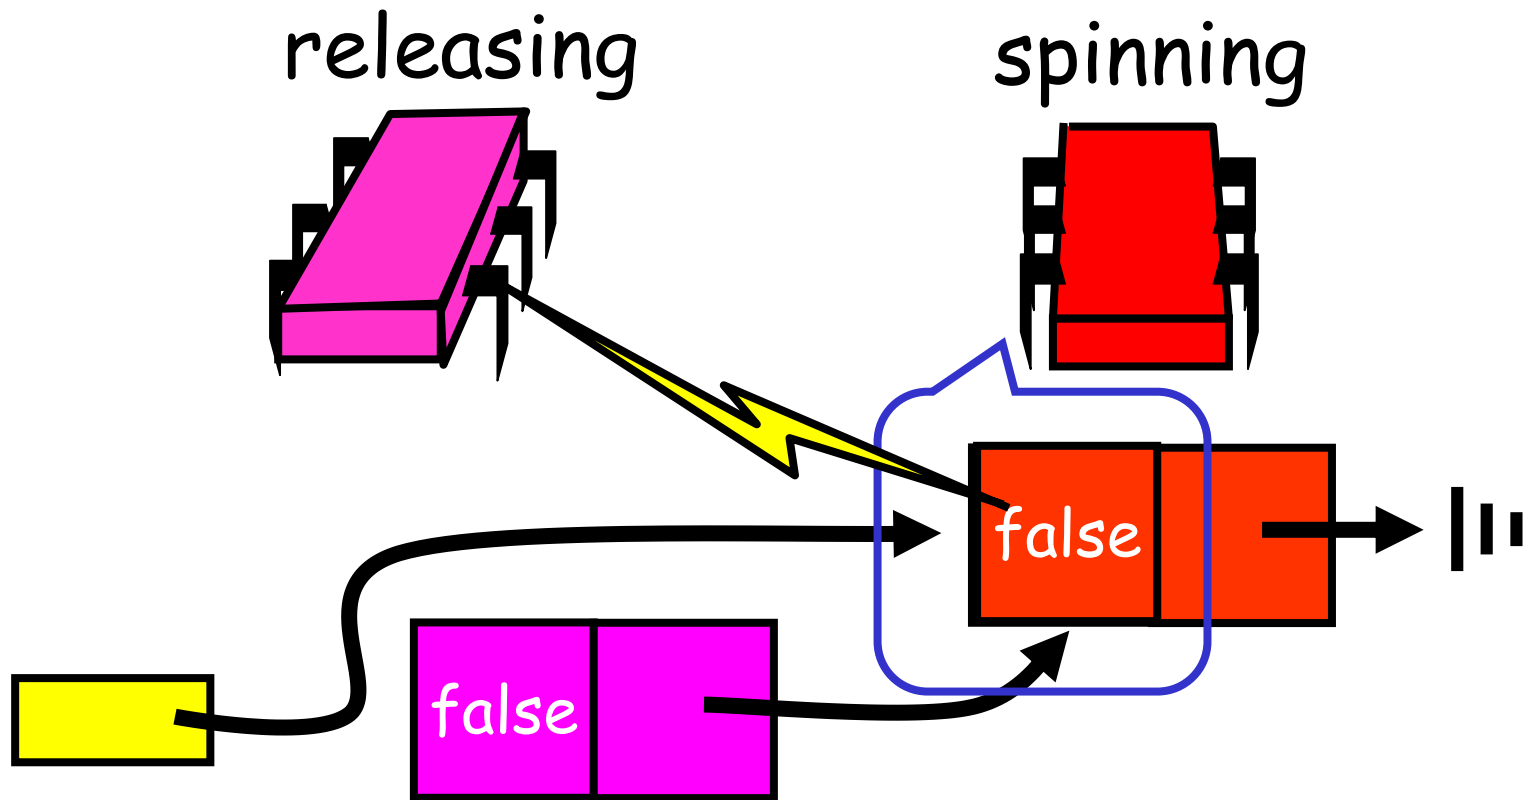
# Purple Release



# Purple Release

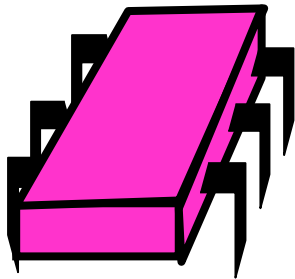


# Purple Release

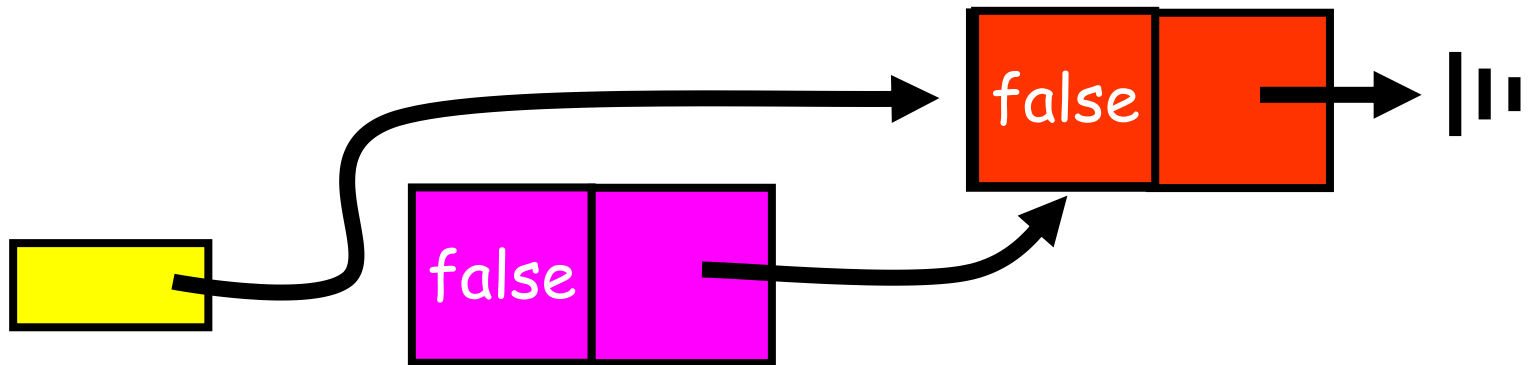
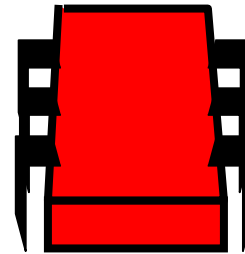


# Purple Release

releasing



Acquired lock







# Lock with timeout

- Java interface includes a `tryLock()` method to specify a maximum duration the thread is willing to wait to acquire the lock
- Should the thread not acquire the lock in the designated time, the thread will timeout



# Abortable Locks

- What if you want to give up waiting for a lock?
- For example
  - Timeout
  - Database transaction aborted by user



# Lock with timeout

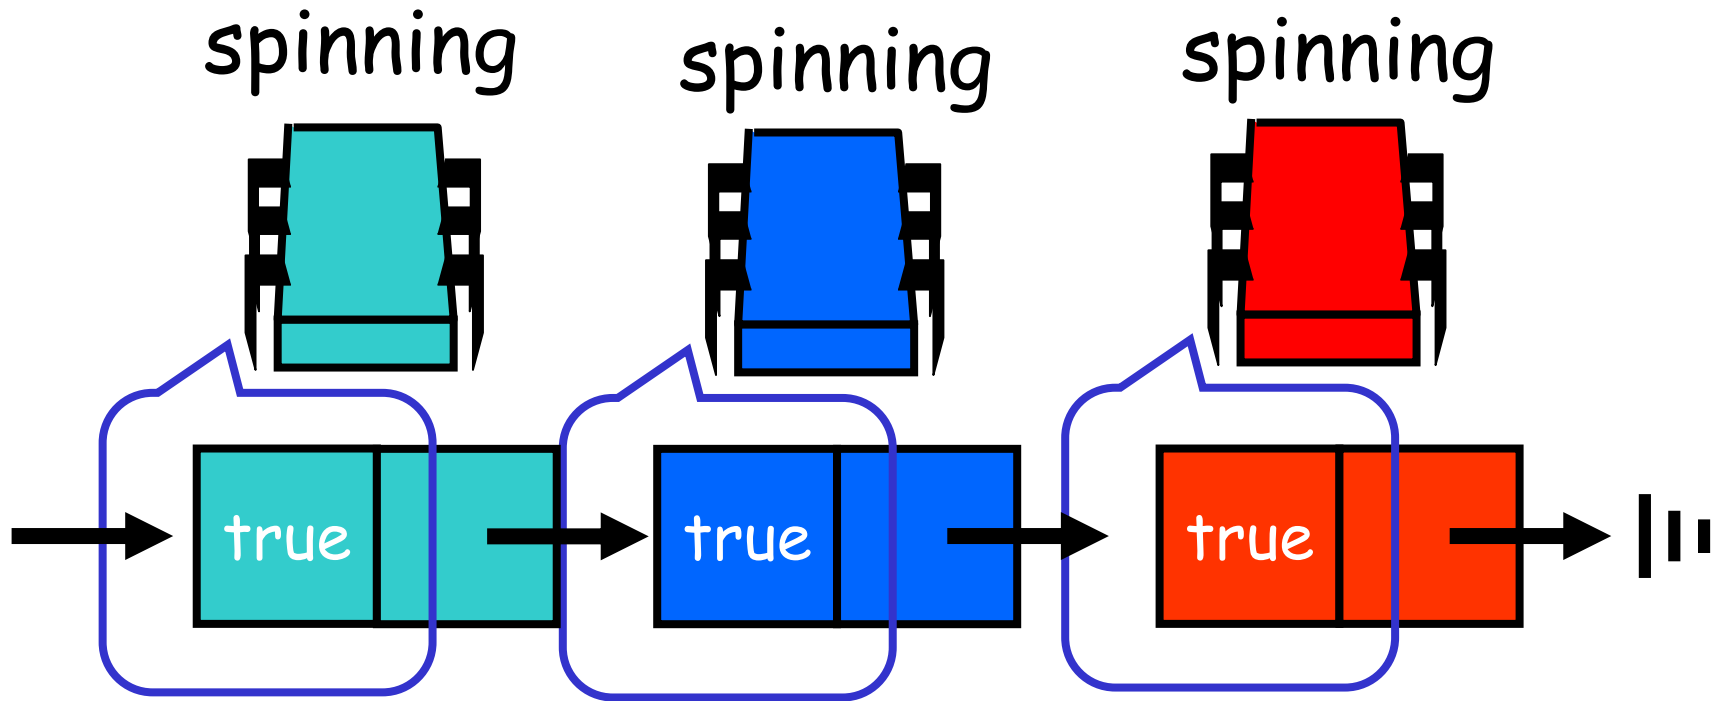
- What happens to other threads when you timeout?




# Back-off Lock

- Aborting is trivial
  - Just return from lock() call
- Extra benefit:
  - No cleaning up
  - Wait-free
  - Immediate return


# MCS Queue Locks



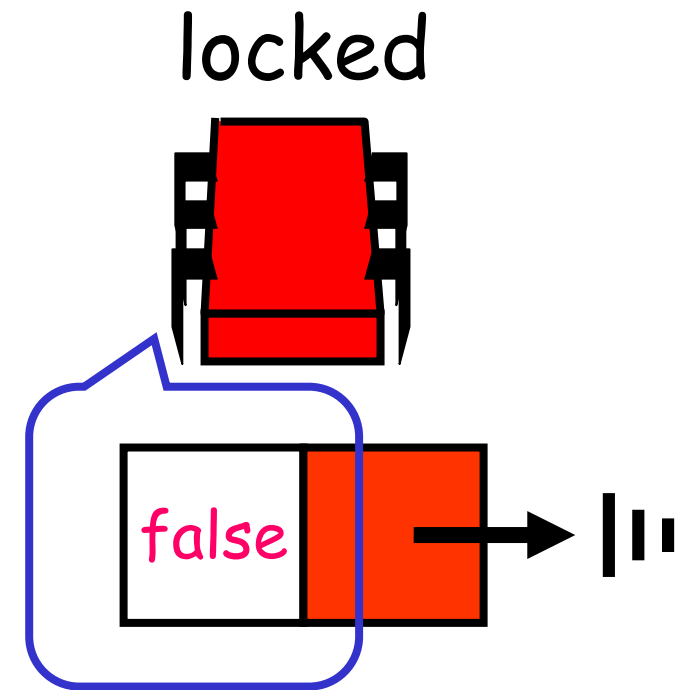
# MCS Queue Locks



# MCS Queue Locks




# MCS Queue Locks

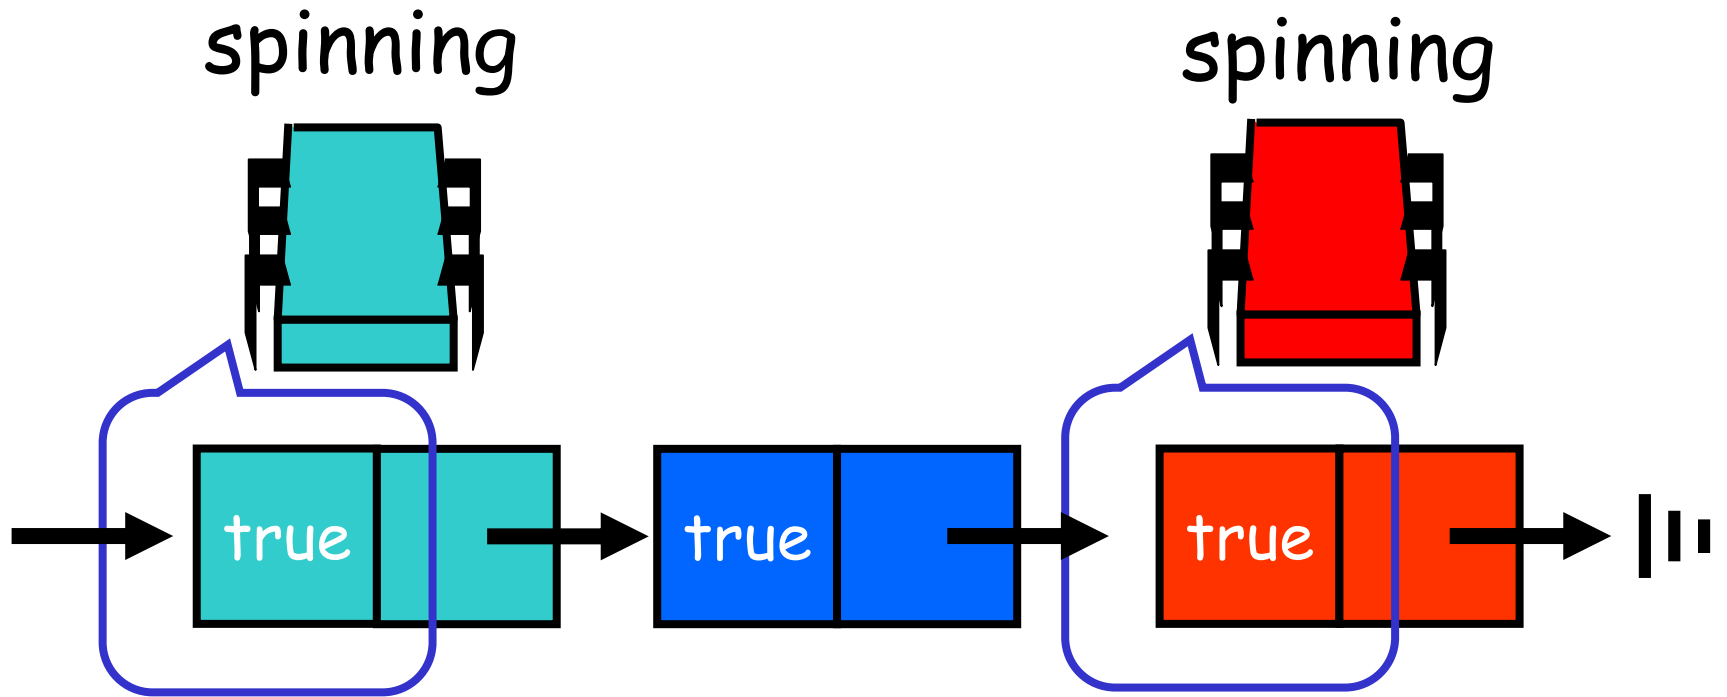




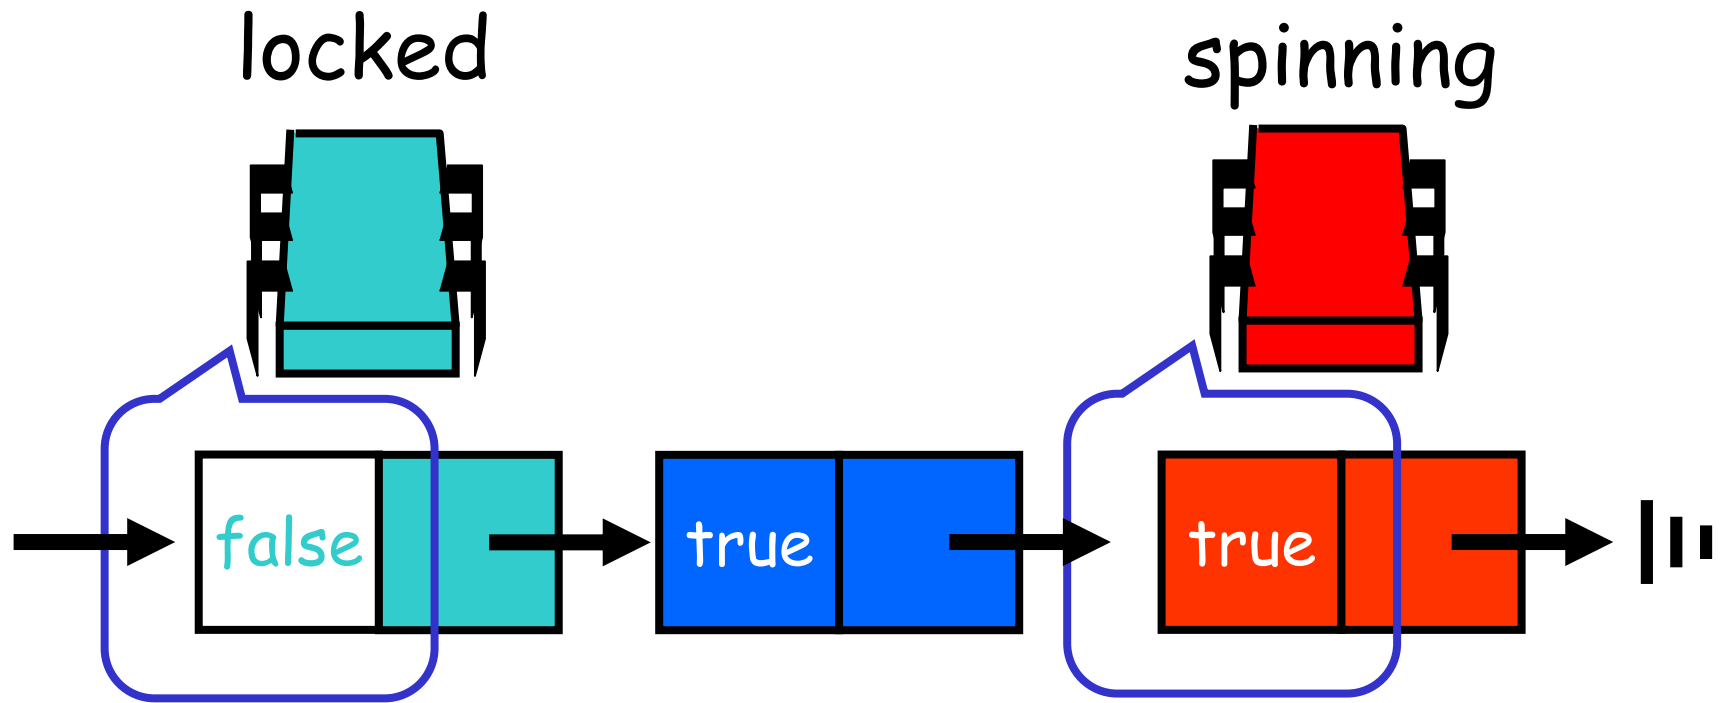
# MCS Queue Locks



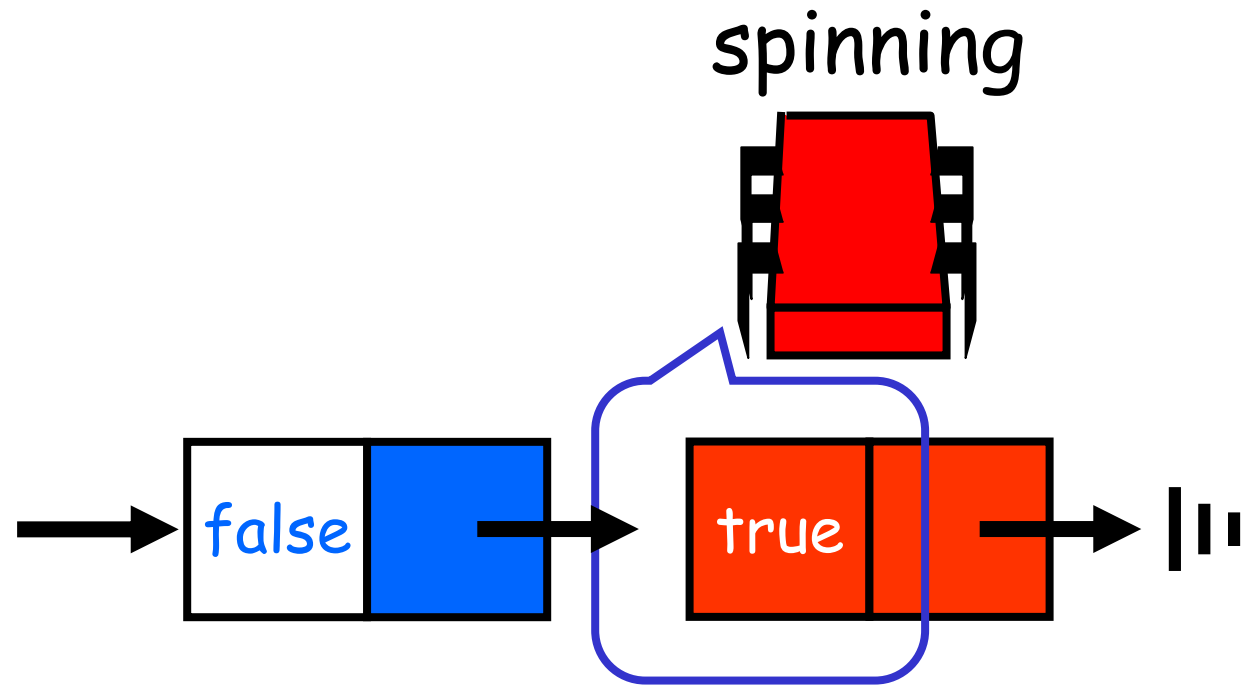
# MCS Queue Locks



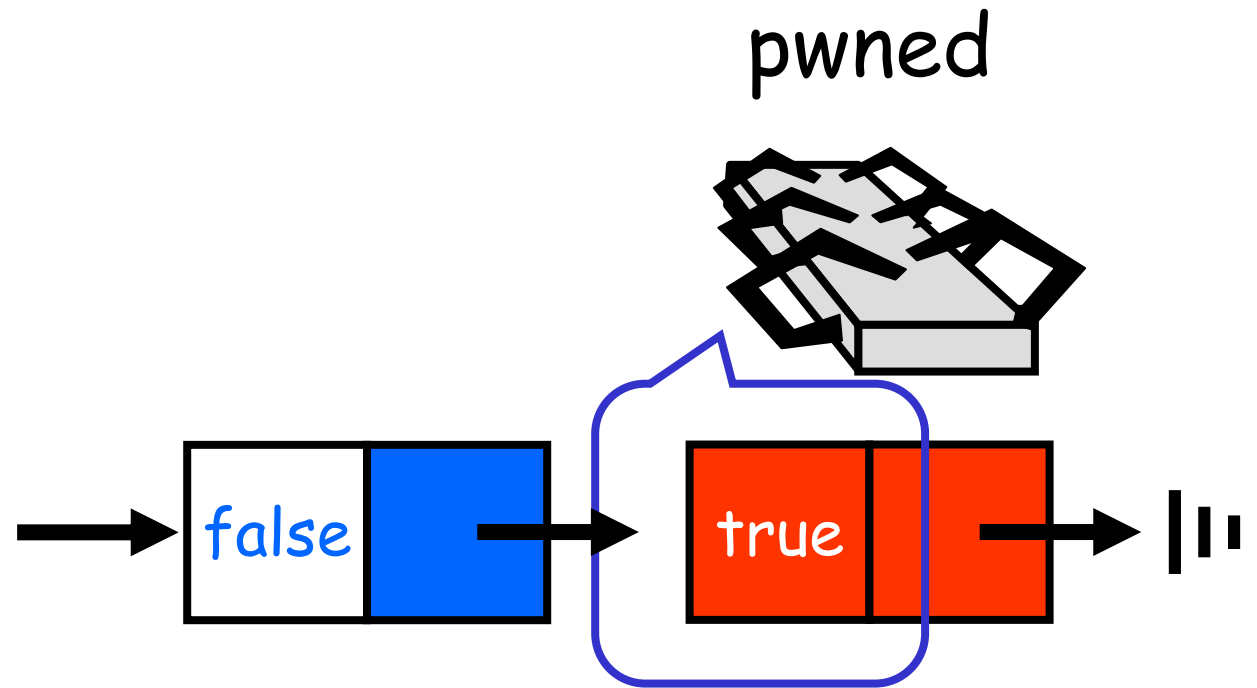
# MCS Queue Locks



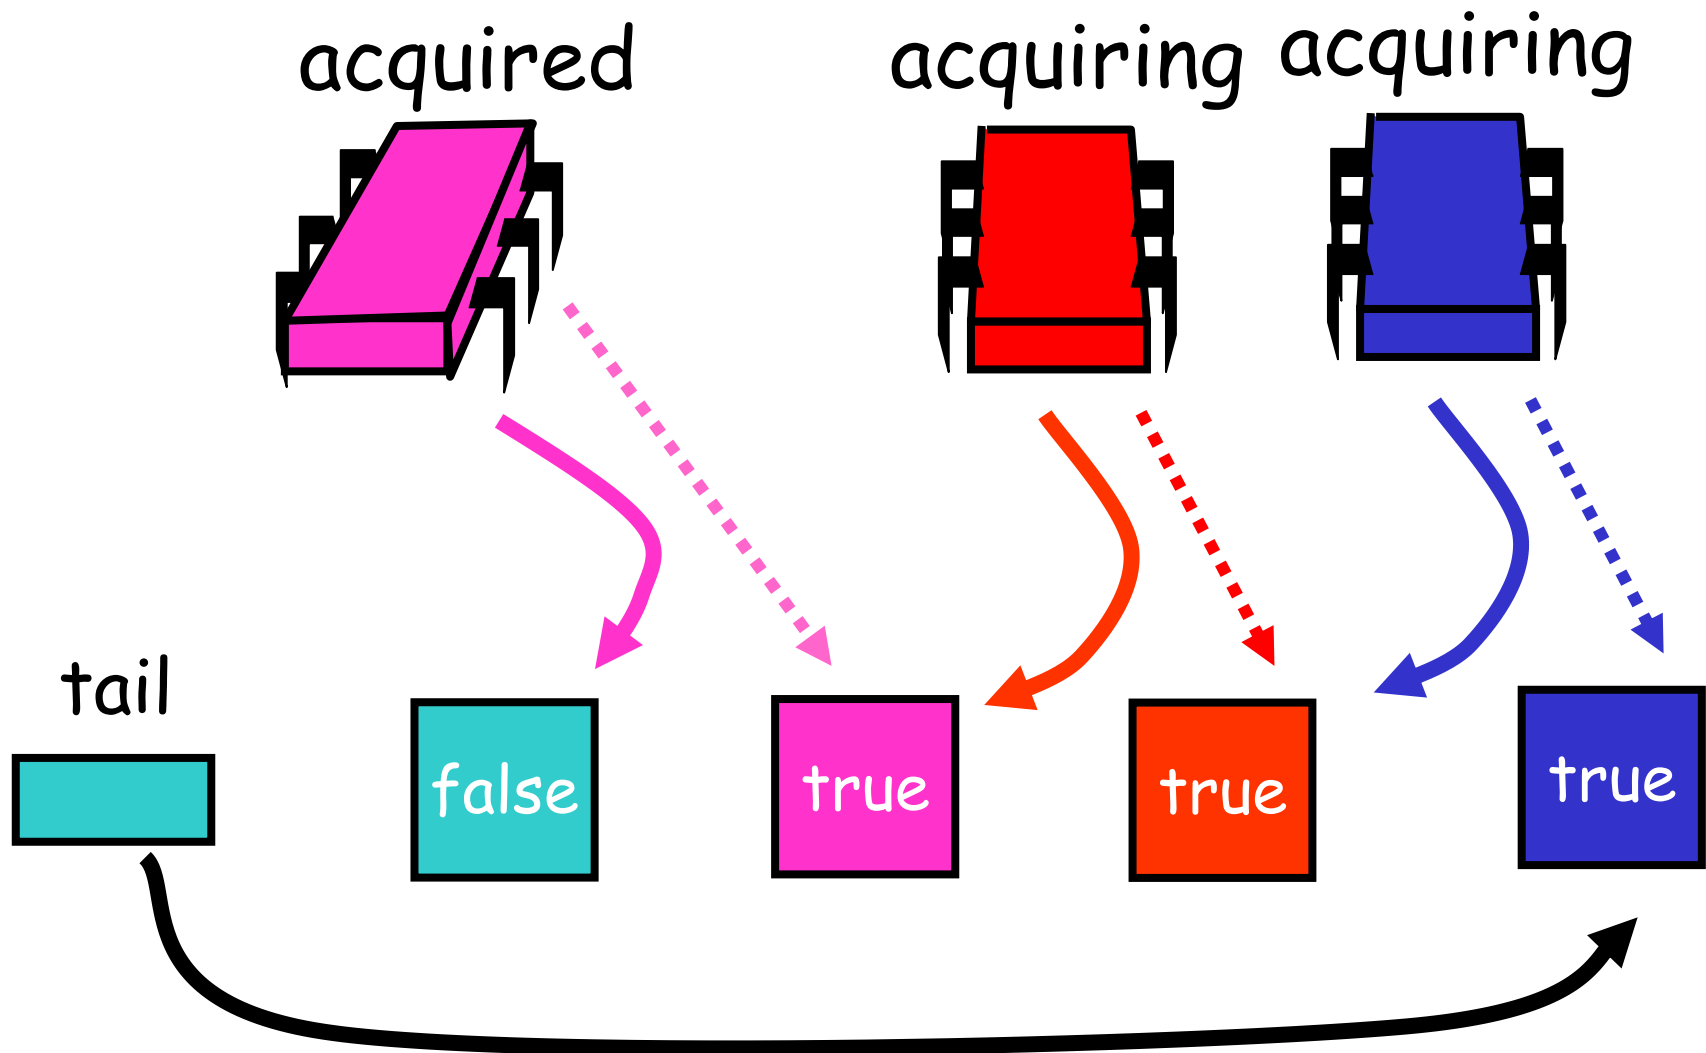
# MCS Queue Locks




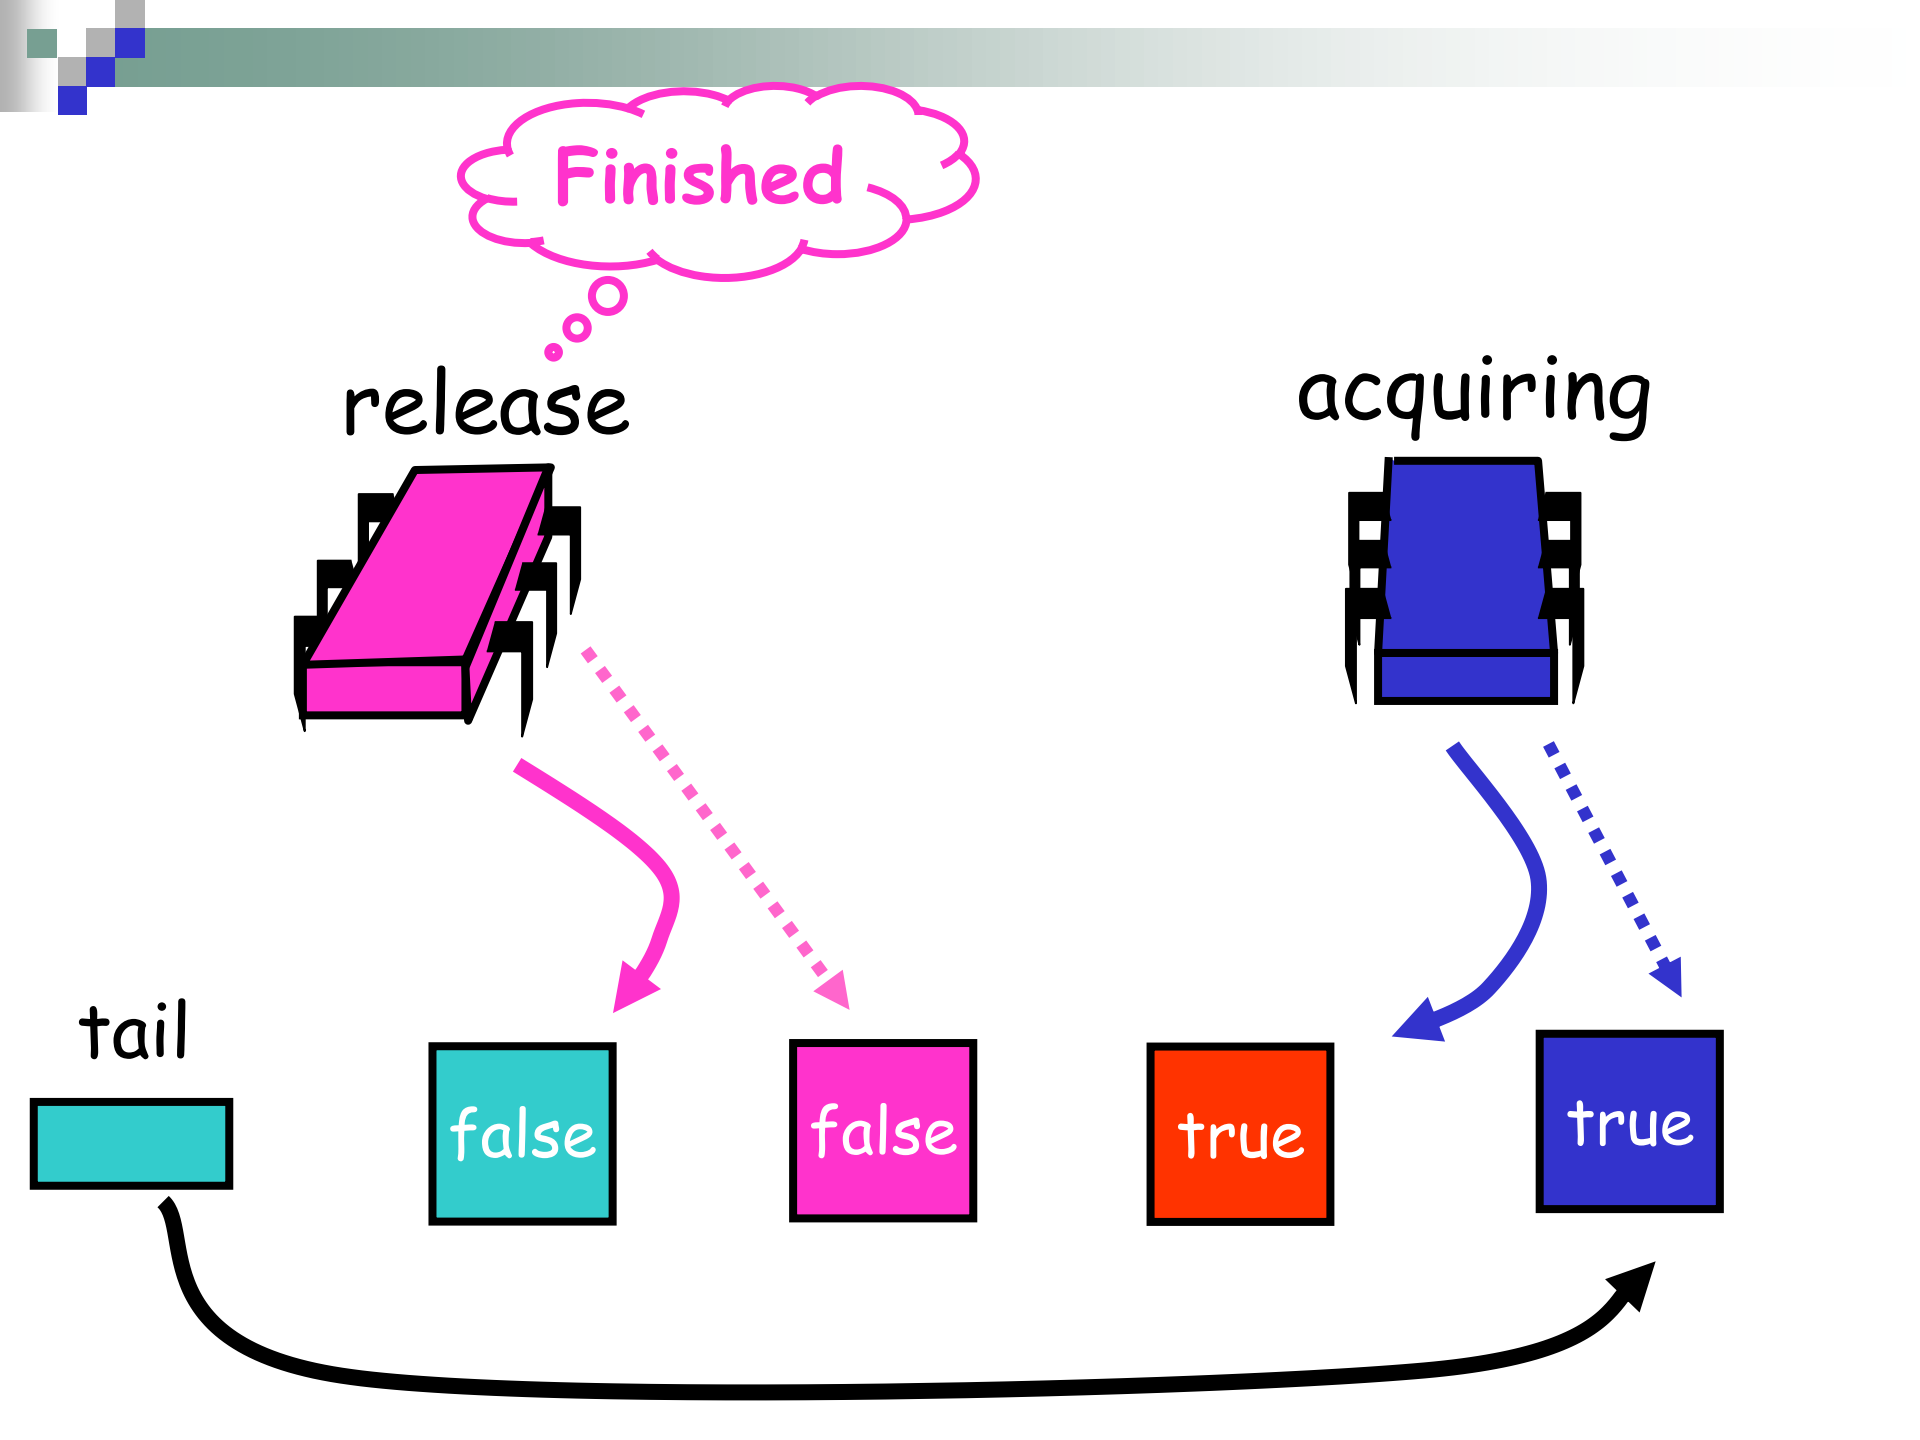
# MCS Queue Locks



# CLH Queue Lock











# Queue Locks

- Can't just quit
  - Thread in line behind will starve
- Need a graceful way out



# Abortable CLH Lock

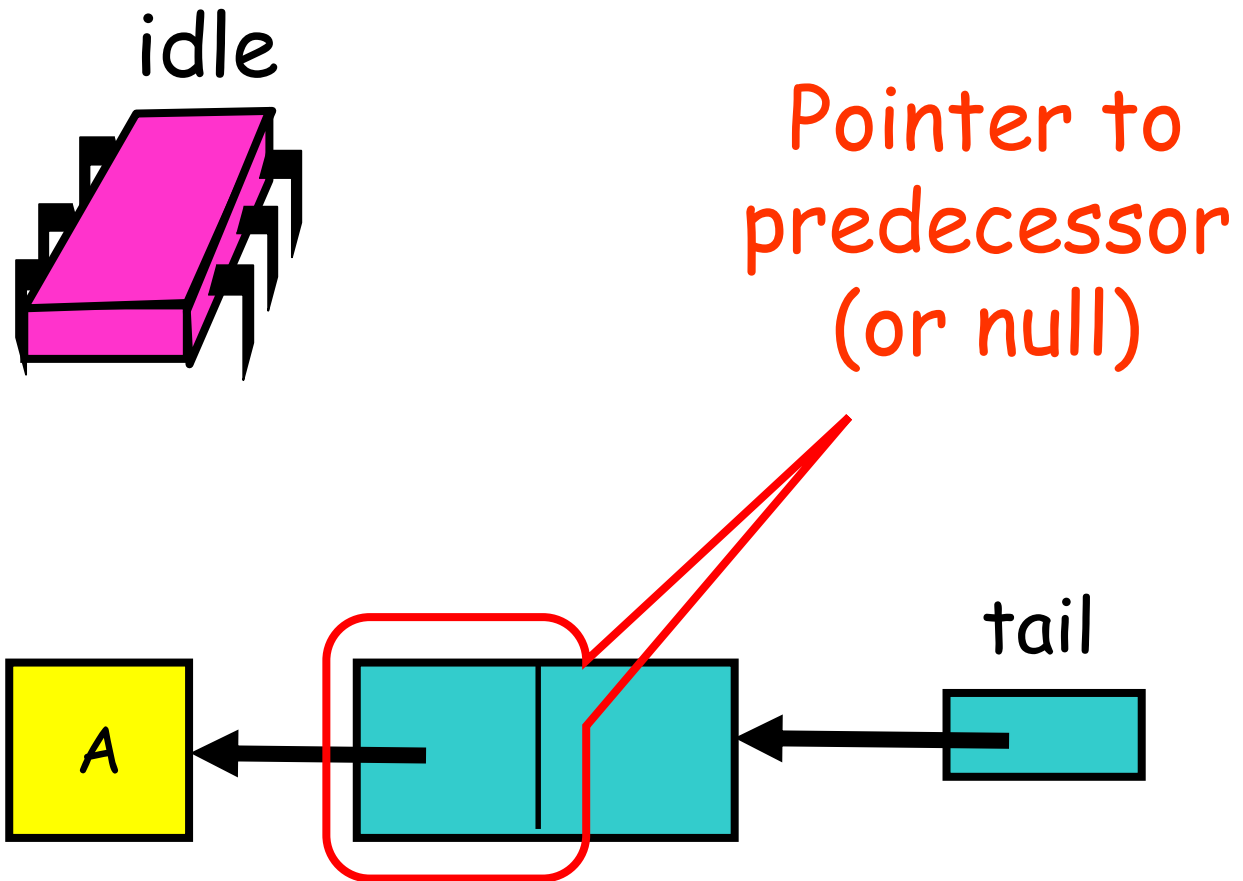
- When a thread gives up
  - Removing node from queue in a wait-free way is hard
- Idea for lazy approach:
  - let successor deal with it.



# A queue lock with timeouts

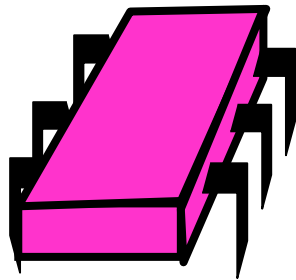
- When a thread times out, it marks its node as abandoned
- The successor in the queue notices that the node has been abandoned
- Successor starts spinning on the abandoned node's predecessors

# Initially

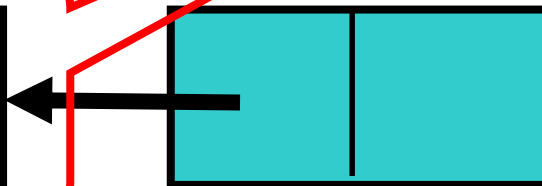
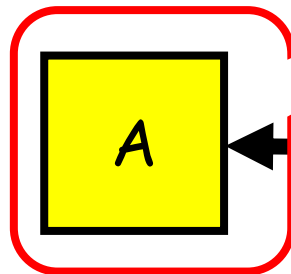


# Acquiring

idle



Distinguished  
available node  
means lock is  
free

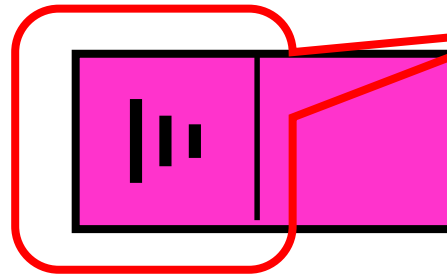
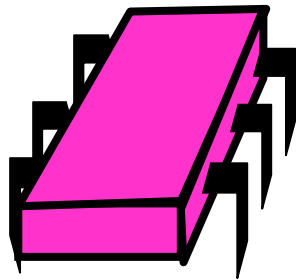


tail

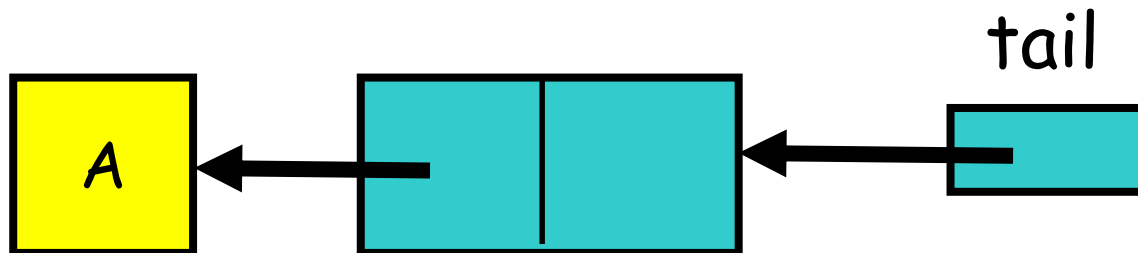


# Acquiring

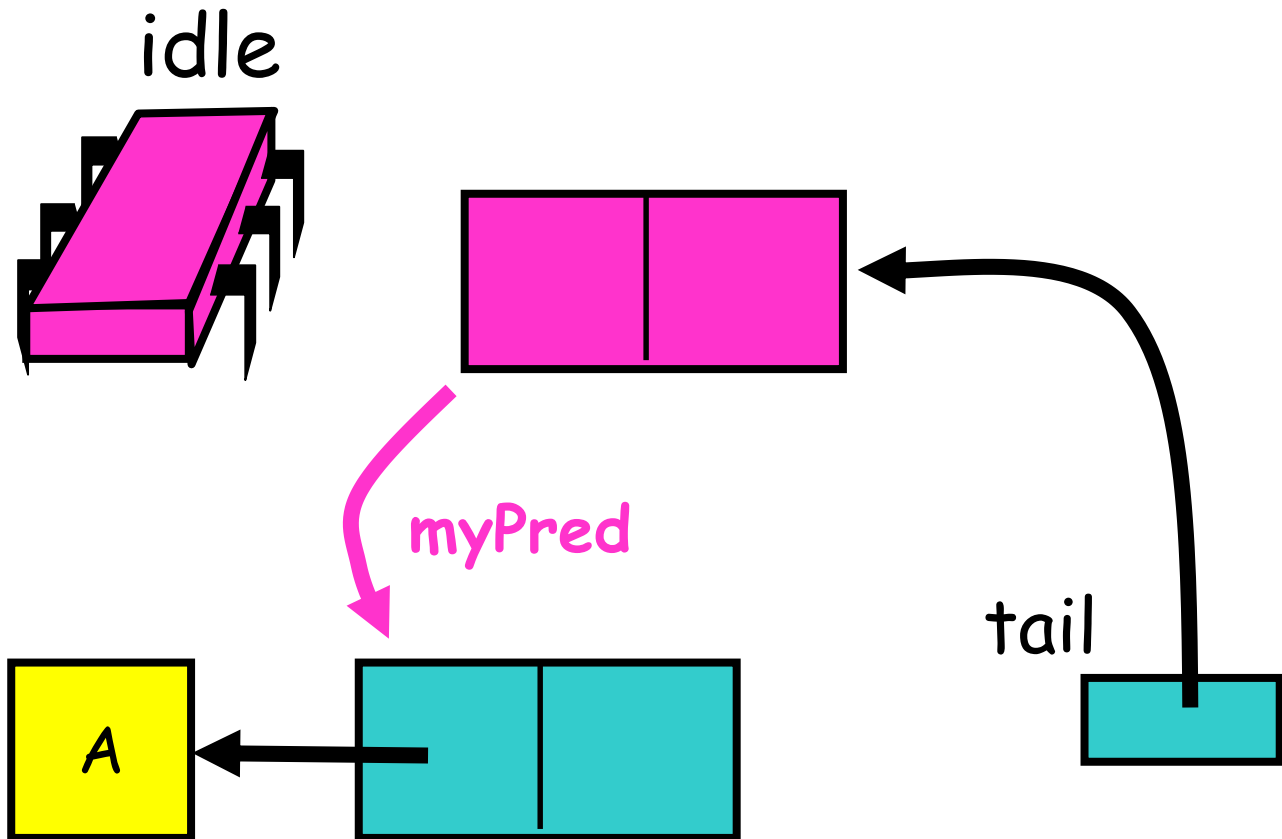
idle




Null predecessor  
means lock not  
released or  
aborted



# Acquiring



# Acquired







# Time-out Lock

```
public class TOLock implements Lock {  
    static Qnode AVAILABLE  
        = new Qnode();  
    AtomicReference<Qnode> tail;  
    ThreadLocal<Qnode> myNode;
```

# Time-out Lock

```
public class TOLock implements Lock {  
    static Qnode AVAILABLE  
        = new Qnode();  
    AtomicReference<Qnode> tail;  
    ThreadLocal<Qnode> myNode;
```

Distinguished node to  
signify free lock

# Time-out Lock

```
public class TOLock implements Lock {  
    static Qnode AVAILABLE  
        = new Qnode();  
    AtomicReference<Qnode> tail;  
    ThreadLocal<Qnode> myNode;
```

Tail of the queue

# Time-out Lock

```
public class TOLock implements Lock {  
    static Qnode AVAILABLE  
        = new Qnode();  
    AtomicReference<Qnode> tail;  
    ThreadLocal<Qnode> myNode;
```

Remember my node ...



# Time-out Lock

```
public boolean lock(long timeout) {  
    Qnode qnode = new Qnode();  
    myNode.set(qnode);  
    qnode.prev = null;  
    Qnode myPred =  
    tail.getAndSet(qnode);  
    if (myPred == null  
        || myPred.prev == AVAILABLE) {  
        return true;  
    }  
}
```

...

# Time-out Lock

```
public boolean lock(long timeout) {  
    Qnode qnode = new Qnode();  
    myNode.set(qnode);  
    qnode.prev = null;  
    Qnode myPred =  
    tail.getAndSet(qnode);  
    if (myPred == null  
        || myPred.prev == AVAILABLE) {  
        return true;  
    }  
}
```

Create & initialize node

# Time-out Lock

```
public boolean lock(long timeout) {  
    Qnode qnode = new Qnode();  
    myNode.set(qnode);  
    qnode.prev = null;  
    Qnode myPred =  
    tail.getAndSet(qnode);  
    if (myPred == null  
        || myPred.prev == AVAILABLE) {  
        return true;  
    }  
}
```

Swap with tail

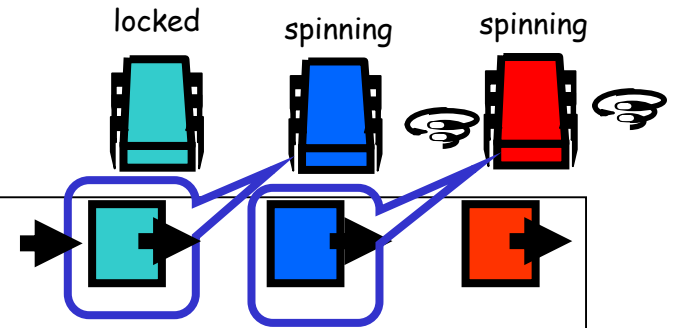
# Time-out Lock

```
public boolean lock(long timeout) {  
    Qnode qnode = new Qnode();  
    myNode.set(qnode);  
    qnode.prev = null;  
    Qnode myPred =  
    tail.getAndSet(qnode);  
    if (myPred == null  
        || myPred.prev == AVAILABLE) {  
        return true;  
    }  
    ...  
}
```

If predecessor absent or  
released, we are done



# Time-out Lock



...

```
long start = now();
while (now() - start < timeout) {
    Qnode predPred = myPred.prev;
    if (predPred == AVAILABLE) {
        return true;
    } else if (predPred != null) {
        myPred = predPred;
    }
}
```

...

# Time-out Lock

...

```
long start = now();  
while (now() - start < timeout) {  
    Qnode predPred = myPred.prev;  
    if (predPred == AVAILABLE) {  
        return true;  
    } else if (predPred != null) {  
        myPred = predPred;  
    }  
}
```

Keep trying for a while ...

...

# Time-out Lock

...

```
long start = now();  
while (now() - start < timeout) {  
    Qnode predPred = myPred.prev;  
    if (predPred == AVAILABLE) {  
        return true;  
    } else if (predPred != null) {  
        myPred = predPred;  
    }  
}
```

...

Spin on predecessor's  
prev field

# Time-out Lock

...

```
long start = now();  
while (now() - start < timeout) {  
    Qnode predPred = myPred.prev;  
    if (predPred == AVAILABLE) {  
        return true;  
    } else if (predPred != null) {  
        myPred = predPred;  
    }  
}
```

Predecessor released lock

...

# Time-out Lock

...

```
long start = now();  
while (now() - start < timeout) {  
    Qnode predPred = myPred.prev;  
    if (predPred == AVAILABLE) {  
        return true;  
    } else if (predPred != null) {  
        myPred = predPred;  
    }  
}
```

...

Predecessor aborted,  
advance one



# Time-out Lock

```
...  
if (!tail.compareAndSet(qnode,  
    myPred))  
    qnode.prev = myPred;  
return false;  
}  
}
```

What do I do when I time out?

# Time-out Lock

```
...  
if (!tail.compareAndSet(qnode,  
    myPred))  
    qnode.prev = myPred;  
return false;  
}  
}
```

Do I have a successor? If CAS  
fails: I do have a successor,  
tell it about myPred

# Time-out Lock

```
...  
if (!tail.compareAndSet(qnode,  
    myPred))  
    qnode.prev = myPred;  
return false;  
}  
}
```

If CAS succeeds: no successor,  
simply return false





# Time-Out Unlock

```
public void unlock() {  
    Qnode qnode = myNode.get();  
    if (!tail.compareAndSet(qnode,  
        null))  
        qnode.prev = AVAILABLE;  
}
```

# Time-out Unlock

```
public void unlock() {  
    Qnode qnode = myNode.get();  
    if (!tail.compareAndSet(qnode,  
        null))  
        qnode.prev = AVAILABLE;  
}
```

If CAS failed: exists  
successor, notify successor  
it can enter

# Timing-out Lock

```
public void unlock() {  
    Qnode qnode = myNode.get();  
    if (!tail.compareAndSet(qnode,  
        null))  
        qnode.prev = AVAILABLE;  
}
```

*CAS successful: set tail to null, no clean up since no successor waiting*



# Composite Locks

- Spin lock algorithms impose trade-offs
- Most spin locks have advantages and disadvantages



# Composite Locks

## ■ Backoff Lock:

- + provide trivial timeout protocols
- are not scalable
- may have critical section underutilization if timeout parameters are not well-tuned.



# Composite Locks

- Queue locks:

- + provide first-come-first-served fairness
- + fast lock release
- nontrivial protocols for abandoned nodes



# Composite Locks

- Composite locks combine the best of both approaches



# Composite Locks

- Idea:

- In a queue only the threads at the front of the queue require lock handoffs

- Solution:

- Keep a small number of waiting threads in a queue and have the rest use exponential backoff





# Composite Locks

- Keep a short, fixed-sized array of lock nodes
- Each thread that tries to acquire the lock selects a node in the array at random
- If the node is in use, the thread backs off and tries again



# Composite Locks

- Once the thread acquires a node, it enqueues that node in a TOLock-style queue
- The thread spins on the preceding node and when that node's owner signals it is done, the thread enters the critical section



# Composite Locks

- Every node in the array has a state:
  - FREE, WAITING, RELEASED, ABORTED
- Actions depend on the state of the randomly selected node



# Composite Lock

- A FREE node is available for other threads to acquire
- A WAITING node is linked into the queue and either in the critical section or waiting to enter
- A node becomes RELEASED when the owner leaves the critical section
- If a node is enqueued but the owner has quit, the node is ABORTED



# Composite Lock

- A thread acquires the lock in three steps:
  - Acquires a node in the waiting array
  - Enqueues the node in the queue
  - Waits until the node reaches the head of the queue

Qnode:

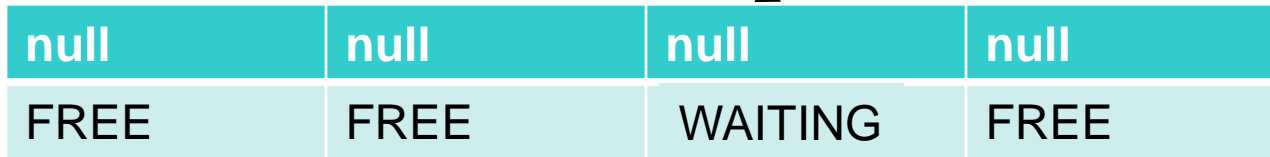
pred

state



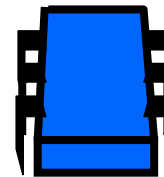
myPred

Tail:



null	null	null	null
FREE	FREE	WAITING	FREE

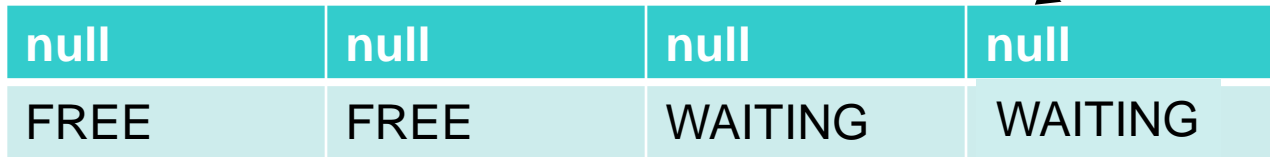
myPred = null



In CS

Randomly  
select array  
location

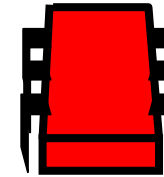
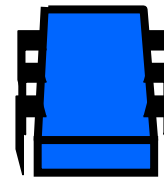
Tail:



null	null	null	null
FREE	FREE	WAITING	WAITING

myPred = null

myPred



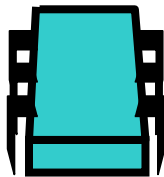
In CS



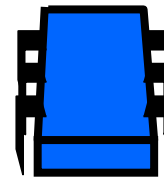
Tail:

null	null	null	null
WAITING	FREE	WAITING	WAITING

myPred

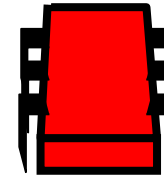


myPred = null



In CS

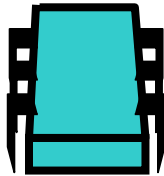
myPred



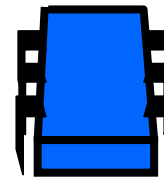
Tail:

null	null	null	
WAITING	FREE	WAITING	ABORTED

myPred





myPred = null



In CS


myPred



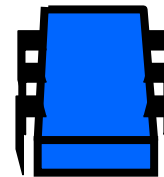
Tail:

null	null	null	
WAITING	FREE	WAITING	FREE

myPred




myPred = null



In CS

myPred



Better  
clean this  
up

# Composite Lock

```
public class CompositeLock implements  
    Lock {  
    AtomicReference<Qnode> tail;  
    Qnode [] waiting;  
    ...
```

Tail field is either null or  
points to last value entered  
into queue (waiting array)

# Composite Lock

```
public class CompositeLock implements  
    Lock {  
    AtomicReference<Qnode> tail;  
    Qnode [] waiting;  
    ...
```

Short queue of waiting  
threads - implemented as  
array



# Composite Lock

```
enum State {FREE, WAITING, RELEASED,
            ABORTED}
class Qnode {
    AtomicReference<State> state;
    Qnode pred;

    public Qnode() {
        state = new
        AtomicReference<State>
        (State.FREE);
    }
}
```

# Composite Lock

```
enum State {FREE, WAITING, RELEASED,
            ABORTED}
class Qnode {
    AtomicReference<State> state;
    Qnode pred;
    public Qnode() {
        state = new
        AtomicReference<State>
        (State.FREE);
    }
}
```

Starting states for all threads is FREE



# Composite Lock

```
public boolean tryLock() {  
    Qnode node = acquireNode();  
    Qnode pred = spliceNode(node);  
    waitForPredecessor(pred, node);  
}
```



# Composite Lock

```
public boolean tryLock() {  
    Qnode node = acquireNode();  
    Qnode pred = spliceNode(node);  
    waitForPredecessor(pred, node);  
}
```

Randomly select node from  
waiting array and return  
available node

# Composite Lock

```
public boolean tryLock() {  
    Qnode node = acquireNode();  
    Qnode pred = spliceNode(node);  
    waitForPredecessor(pred, node);  
}
```

*Add selected node to queue*

# Composite Lock

```
public boolean tryLock() {  
    Qnode node = acquireNode();  
    Qnode pred = spliceNode(node);  
    waitForPredecessor(pred, node);  
}
```

*Add selected node to queue*



# acquireNode()

- Thread selects a node at random
- Tries to acquire the node by setting its state from FREE to WAITING
- If it fails it examines the state
  - If ABORTED or RELEASED it recycles the node



# Composite Lock

```
private Qnode acquireNode() {  
    Qnode node = waiting[random.nextInt()];  
    while (true) {  
        if (state.compareAndSet(FREE, WAITING))  
            return node;  
    }  
}
```

# Composite Lock

Choose random location  
in waiting array

```
private Qnode acquireNode() {  
    Qnode node = waiting[random.nextInt()];  
    while (true) {  
        if (state.compareAndSet(FREE, WAITING))  
            return node;  
    }  
}
```

# Composite Lock

If state was FREE then  
no problem

```
private Qnode acquireNode() {  
    Qnode node = waiting[random.nextInt()];  
    while (true) {  
        if (state.compareAndSet(FREE, WAITING))  
            return node;  
    }  
}
```

If state **ABORTED** or  
**RELEASED**

# Composite Lock

```
currTail = tail;
if (state == ABORTED || state ==
RELEASED) {
    if (node == currTail) {
        Qnode myPred = null;
        if (state == ABORTED)
            myPred = node.pred;
        if (tail.compareAndSet(
            currTail, myPred)
            node.state.set(WAITING);
        return node;
    }
}
```



# Composite Lock

Can only be recycled if  
last node in queue

```
currTail = tail;
if (state == ABORTED || state ==
RELEASED) {
    if (node == currTail) {
        Qnode myPred = null;
        if (state == ABORTED)
            myPred = node.pred;
        if (tail.compareAndSet(
            currTail, myPred)
            node.state.set(WAITING);
        return node;
    }
}
```

# Composite Lock

If last node is  
**ABORTED** tail is set to  
predecessor otherwise

```
currTail = tail;
if (state == ABORTED || state ==  
RELEASED) {  
    if (node == currTail) {  
        Qnode myPred = null;  
        if (state == ABORTED)  
            myPred = node.pred;  
        if (tail.compareAndSet(  
            currTail, myPred)  
            node.state.set(WAITING);  
        return node;  
    }  
}
```

# Composite Lock

Successfully recycle  
node and set to  
**WAITING**

```
currTail = tail;
if (state == ABORTED || state ==
RELEASED) {
    if (node == currTail) {
        Qnode myPred = null;
        if (state == ABORTED)
            myPred = node.pred;
        if (tail.compareAndSet(
            currTail, myPred)
            node.state.set(WAITING);
        return node;
    }
}
```



# spliceNode()

- Once node is acquired, the node is spliced into the queue
- Thread tries to set tail to node
- If timeout, set node to FREE and throw TimeoutException
- If succeeds return prior value of tail (node's predecessor)



# Composite Lock

```
private Qnode spliceNode(Qnode node) {
    Qnode currTail;
    do {
        currTail = tail;
        if (timeout()) {
            node.state.set(FREE);
            throw new TimeoutException();
        }
    } while (!tail.compareAndSet(currTail,
        node));
    return currTail;
}
```

If waited too long...

# Composite Lock

```
private Qnode spliceNode(Qnode node) {  
    Qnode currTail;  
    do {  
        currTail = tail;  
        if (timeout()) {  
            node.state.set(FREE);  
            throw new TimeoutException();  
        }  
    } while (!tail.compareAndSet(currTail,  
        node));  
    return currTail;  
}
```

# Composite Lock

```
private Qnode spliceNode(Qnode node) {  
    Qnode currTail;  
    do {  
        currTail = tail;  
        if (timeout()) {  
            node.state.set(FREE);  
            throw new TimeoutException();  
        }  
    } while (!tail.compareAndSet(currTail,  
        node);  
    return currTail;  
}
```

Try to set tail to new  
node



# Composite Lock

- Finally thread has to wait its turn
- If predecessor is null thread can enter CS
- If predecessor is not RELEASED, check if it is ABORTED
  - If so, mark as FREE and get ABORTED node's predecessor
- If predecessor is RELEASED enter CS





# Composite Lock

```
private void waitforPredecessor {  
    if (pred == null)  
        return;  
    State predState = pred.state.get();  
    while (predState != RELEASED) {  
        if (predState == ABORTED) {  
            Qnode temp = pred;  
            pred = pred.pred;  
            temp.state.set(FREE);  
        }  
    }  
}
```

# Composite Lock

```
private void waitforPredecessor {  
    if (pred == null)  
        return;  
    State predState = pred.state.get();  
    while (predState != RELEASED) {  
        if (predState == ABORTED) {  
            Qnode temp = pred;  
            pred = pred.pred;  
            temp.state.set(FREE);  
        }  
    }
```

If I am first in line

# Composite Lock

```
private void waitforPredecessor {  
    if (pred == null)  
        return;
```

```
    State predState = pred.state.get();  
    while (predState != RELEASED) {  
        if (predState == ABORTED) {  
            Qnode temp = pred;  
            pred = pred.pred;  
            temp.state.set(FREE);  
        }  
    }
```

Spin on predecessor's  
state

# Composite Lock

```
private void waitforPredecessor {  
    if (pred == null)  
        return;  
    State predState = pred.state.get();  
    while (predState != RELEASED) {  
        if (predState == ABORTED) {  
            Qnode temp = pred;  
            pred = pred.pred;  
            temp.state.set(FREE);  
        }  
    }
```

If predecessor **ABORTED** set to  
**FREE** and spin on predecessors  
predecessor

# Composite Lock

```
    if (timeout()) {  
        node.pred = pred;  
        node.state.set(ABORTED);  
    } throw new TimeoutException();  
    predState = pred.state.get();  
}  
pred.state.set(FREE);  
return;  
}
```

If I have waited long enough...

# Composite Lock

```
    if (timeout()) {  
        node.pred = pred;  
        node.state.set(ABORTED);  
    } throw new TimeoutException();  
    predState = pred.state.get();  
}  
pred.state.set(FREE);  
return;  
}
```

When predecessor RELEASED



# Composite Lock

```
public void unlock() {  
    myNode.state.set(RELEASED);  
}
```